



**Hewlett Packard**  
Enterprise

# OCT 2023 CHAPEL TUTORIAL UPDATED FOR USE IN SPRING 2025 CS372 CLASS

---

Chapel Team, edited by Michelle Strout

March 27 and April 1, 2025

# PLAN

- **Announcements**

- MT2 and LA2 grades have been posted
- SA6 is due Monday, survey for final project preferences
- Some reading assignment for Chapel have been posted

- **Last time**

- Midterm 2

- **Today**

- Chapel introduction
- Please start the docker pull for Chapel (see next slide)

# HANDS ON: HOW TO DO THE HANDS ON



01-hello.chpl

## Example codes for Chapel tutorial slides

- <https://github.com/UofA-CSc-372-Spring-2025/CSc372Spring2025-CourseMaterials/tree/main/Sandboxes/ChapelTutorialExamples>

## Using a container on your laptop

- First, install docker for your machine and start it up (see the README.md for more info)
- Then, use the chapel-gasnet docker container

```
docker pull docker.io/chapel/chapel-gasnet      # takes about 5 minutes
cd CSc372Spring2025-CourseMaterials/Sandboxes/ChapelTutorialExamples/
docker run --rm -it -v "$PWD":/workspace chapel/chapel-gasnet
root@589405d07f6a:/opt/chapel# cd /workspace
root@xxxxxxxx:/myapp# chpl 01-hello.chpl
root@xxxxxxxx:/myapp# ./01-hello -nl 1
```



## OUTLINE: OVERVIEW OF PROGRAMMING IN CHAPEL

---

- Chapel Goals, Usage, and Comparison with other Tools
- Hello World (Hands On)
- Chapel Execution Model and Parallel Hello World (Hands On)
- kmer counting using file IO, config consts, strings, maps (Hands On)
- Parallelizing a program that processes files (Hands On)
- GPU programming support
- Learning goals for rest of tutorial



CHAPEL GOALS, USAGE, AND COMPARISON WITH OTHER TOOLS

# CHAPEL PROGRAMMING LANGUAGE

---

Chapel is a general-purpose programming language that provides **ease of parallel programming, high performance, and portability.**

And is being used in applications in various ways:


**refactoring** existing codes,

**developing** new codes,

serving high performance to Python codes (**Chapel server with Python client**), and **providing distributed and shared memory parallelism** for existing codes.




# APPLICATIONS OF CHAPEL: LINKS TO USERS' TALKS (SLIDES + VIDEO)




CHAMPS: 3D Unstructured CFD

[CHIUW 2021](#) [CHIUW 2022](#)




Arkouda: Interactive Data Science at Massive Scale

[CHIUW 2020](#) [CHIUW 2023](#)




ChOp: Chapel-based Optimization

[CHIUW 2021](#) [CHIUW 2023](#)




ChpUltra: Simulating Ultralight Dark Matter

[CHIUW 2020](#) [CHIUW 2022](#)



Lattice-Symmetries: a Quantum Many-Body Toolbox Desk dot chpl: Utilities for Environmental Eng.

[CHIUW 2022](#)




[CHIUW 2022](#)




RapidQ: Mapping Coral Biodiversity

[CHIUW 2023](#)




ChapQG: Layered Quasigeostrophic CFD




Chapel-based Hydrological Model Calibration

[CHIUW 2023](#)




CrayAI HyperParameter Optimization (HPO)

[CHIUW 2021](#)



CHGL: Chapel Hypergraph Library

[CHIUW 2020](#)




Your Application Here?

# HIGHLIGHTS OF CHAPEL USAGE

## **CHAMPS:** Computational Fluid Dynamics framework for airplane simulation

- Professor Eric Laurendeau's team at Polytechnique Montreal
- Performance: achieves competitive results w.r.t. established, world-class frameworks from Stanford, MIT, etc.
- Programmability: *"We ask students at the master's degree to do stuff that would take 2 years and they do it in 3 months."*



## **Arkouda:** data analytics framework (<https://github.com/Bears-R-Us/arkouda>)

- Mike Merrill, Bill Reus, et al., US DOD
- Python front end client, Chapel server that processes dozens of terabytes in seconds
- April 2023: 1200 GiB/s for argsort on an HPE EX system



## **Other recent users**

- Marjan Asgari et al, "Development of a knowledge-sharing parallel computing approach for calibrating distributed watershed hydrologic models", Environmental Modeling and Software.
- Scott Bachman has written some coral reef image analysis applications in Chapel.





# CHAPEL IS HIGHLY PERFORMANT AND SCALABLE

## HPE Apollo (May 2021)



- HDR-100 Infiniband network (100 Gb/s)
- 576 compute nodes
- 72 TiB of 8-byte values
- ~480 GiB/s (~150 seconds)

## HPE Cray EX (April 2023)




- Slingshot-11 network (200 Gb/s)
- 896 compute nodes
- 28 TiB of 8-byte values
- ~1200 GiB/s (~24 seconds)

## HPE Cray EX (May 2023)



- Slingshot-11 network (200 Gb/s)
- 8192 compute nodes
- 256 TiB of 8-byte values
- ~8500 GiB/s (~31 seconds)

Arkouda Argosort Performance



A notable performance achievement in ~100 lines of Chapel



# COMPARE WITH OTHER PARALLEL PROGRAMMING MODELS

- Shared-memory parallelism

- Pthreads: low-level library for creating and managing threads
- OpenMP: pragmas added before loops and other statements
- Rust, Julia: programming languages with some threaded parallelism
- RAJA, Kokkos: C++ libraries that use template metaprogramming

- Distributed-memory parallelism and shared-memory parallelism

- MPI+X:

- MPI stands for message passing interface
- MPI is a library for sending and receiving messages between processes
- All processes allocate their own memory and run the same program
- There are many options for X: OpenMP, Pthreads, Python, Julia, R, etc.

- OpenSHMEM: library for implementing a partitioned global address space

- Spark: Python, Scala, and Java accessible library for especially iterative data processing

- Regent and Legion: programming language and runtime that implements implicit task parallelism

- Kokkos Remote Spaces: extends Kokkos C++ template views to distributed views

## Chapel:

- shared memory parallelism,
- distributed-memory parallelism,
- data parallelism,
- task parallelism,
- map-reduce parallelism,
- vector parallelism,
- GPU parallelism, ...

All can be expressed in the same programming language.

# COMPARE WITH OTHER PARALLEL PROGRAMMING MODELS (W/OUT CHAPEL BOX)

- Shared-memory parallelism
  - Pthreads: low-level library for creating and managing threads of execution that share memory
  - OpenMP: pragmas added before loops and other statements in C/C++/Fortran programs
  - Rust, Julia: programming languages with some threaded parallelism constructs
  - RAJA, Kokkos: C++ libraries that use template metaprogramming to express parallel policies
- Distributed-memory parallelism and shared-memory parallelism
  - MPI+X:
    - MPI stands for message passing interface
    - MPI is a library for sending and receiving messages between processes
    - All processes allocate their own memory and run the same program, SPMD: Single Program Multiple Data
    - There are many options for X: OpenMP, Pthreads, Python, Julia, RAJA, Kokkos, Chapel, ...
  - OpenSHMEM: library for implementing a partitioned global address space
  - Spark: Python, Scala, and Java accessible library for especially the map-reduce parallelism
  - Regent and Legion: programming language and runtime that implements implicit task parallelism
  - Kokkos Remote Spaces: extends Kokkos C++ template views to distributed views




## LET'S COMPARE WITH STREAM TRIAD: A PARALLEL COMPUTATION

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures:**




## LET'S COMPARE WITH STREAM TRIAD: A PARALLEL COMPUTATION

**Given:**  $n$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..n, A_i = B_i + \alpha \cdot C_i$

**In pictures, in parallel** (distributed memory multicore, global-view):



# STREAM TRIAD: IN MPI+OPENMP

```
#include <hpcc.h>
#ifdef _OPENMP
#include <omp.h>
#endif
```

```
static int VectorSize;
static double *a, *b, *c;
```

```
int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
```

```
rv = HPCC_Stream(params, 0, myRank);
MPI_Reduce(&rv, &errCount, 1, MPI_INT, MPI_SUM,
0, comm);


return errCount;
}
```

```
int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;
```

```
VectorSize = HPCC_LocalVectorSize(params, 3,
    sizeof(double), 0);
```

```
a = HPCC_XMALLOC(double, VectorSize);
b = HPCC_XMALLOC(double, VectorSize);
c = HPCC_XMALLOC(double, VectorSize);
```

**MPI + OpenMP**



```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf(outFile, "Failed to
        allocate memory (%d).\n",
        VectorSize);
        fclose(outFile);
    }
}
```

```
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 1.0;
}
scalar = 3.0;
```


```
#ifdef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);

return 0; }
```

```
#define N 2000000
```

**CUDA**



```
int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
```

```
scalar = 3.0;
STREAM_Triad<<<dimGrid,dimBlock>>>>(d_b, d_c, d_a, scalar, N);
cudaThreadSynchronize();

cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad(float *a, float *b, float *c,
    float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```

*HPC suffers from too many distinct notations for expressing parallelism and locality.  
This tends to be a result of bottom-up language design.*

# WHY SO MANY PROGRAMMING MODELS?

HPC tends to approach programming models bottom-up:

Given a system and its core capabilities...

...provide features that permit users to access the available performance.

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	MPI	executable
Intra-node/multicore	OpenMP / pthreads	iteration/task
Instruction-level vectors/threads	pragmas	iteration
GPU/accelerator	CUDA / Open[MP CL ACC]	SIMD function/task

**benefits:** lots of control; decent generality; easy to implement


**downsides:** lots of user-managed detail; brittle to changes



# STREAM TRIAD: IN CHAPEL

**The special sauce:**  
How should this index set—  
and any arrays and  
computations over it—be  
mapped to the system?

```
use BlockDist;  
  
config const m = 1000,  
           alpha = 3.0;  
  
const ProblemSpace = blockDist.createDomain({1..m});  
  
var A, B, C: [ProblemSpace] real;  
  
B = 2.0;  
C = 1.0;  
  
A = B + alpha * C;
```



**Philosophy:** *Top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.



HELLO WORLD (HANDS ON)



## Example codes for Chapel tutorial slides

- <https://github.com/UofA-CSc-372-Spring-2025/CSc372Spring2025-CourseMaterials/tree/main/Sandboxes/ChapelTutorialExamples>

## Using a container on your laptop

- First, install docker for your machine and start it up (see the README.md for more info)
- Then, use the chapel-gasnet docker container

```
docker pull docker.io/chapel/chapel-gasnet      # takes about 5 minutes
cd CSc372Spring2025-CourseMaterials/Sandboxes/ChapelTutorialExamples/
docker run --rm -it -v "$PWD":/workspace chapel/chapel-gasnet
root@589405d07f6a:/opt/chapel# cd /workspace
root@xxxxxxxx:/myapp# chpl 01-hello.chpl
root@xxxxxxxx:/myapp# ./01-hello -nl 1
```



# "HELLO WORLD" IN CHAPEL: TWO VERSIONS

- Fast prototyping

```
writeln("Hello, world!");
```



01-hello.chpl

- “Production-grade”

```
module Hello {  
  
    proc main() {  
        writeln("Hello, world!");  
    }  
  
}
```



01-hello-production.chpl



# "HELLO WORLD" IN CHAPEL: TWO VERSIONS

- Fast prototyping (configurable)

```
config const audience = "world";  
writeln("Hello, ", audience, "!");
```

 01-hello-configurable.chpl

- “Production-grade” (configurable)

```
module Hello {  
    config const audience = "world";  
  
    proc main() {  
        writeln("Hello, ", audience, "!");  
    }  
}
```

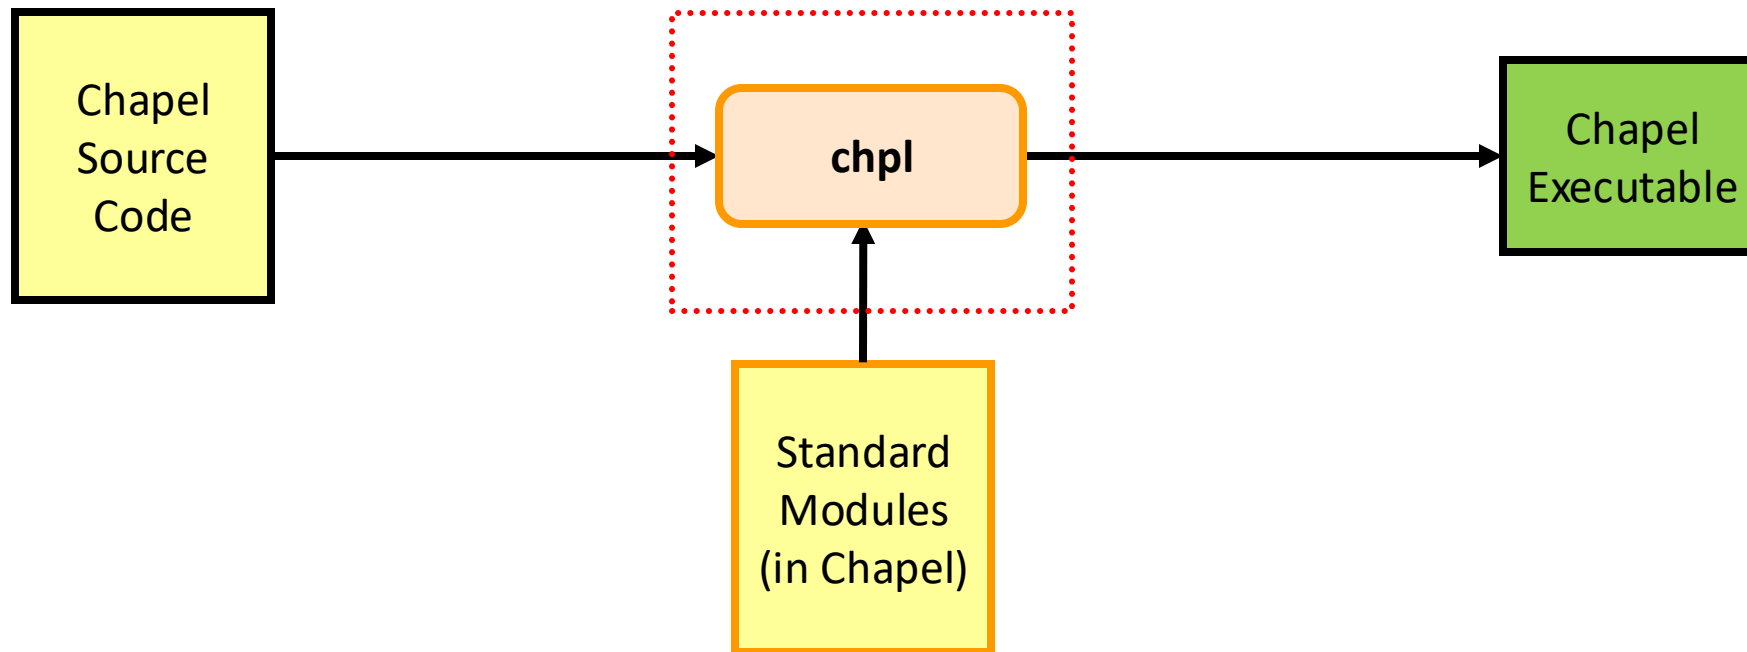
 01-hello-production-configurable.chpl

- To change ‘audience’ for a given run:

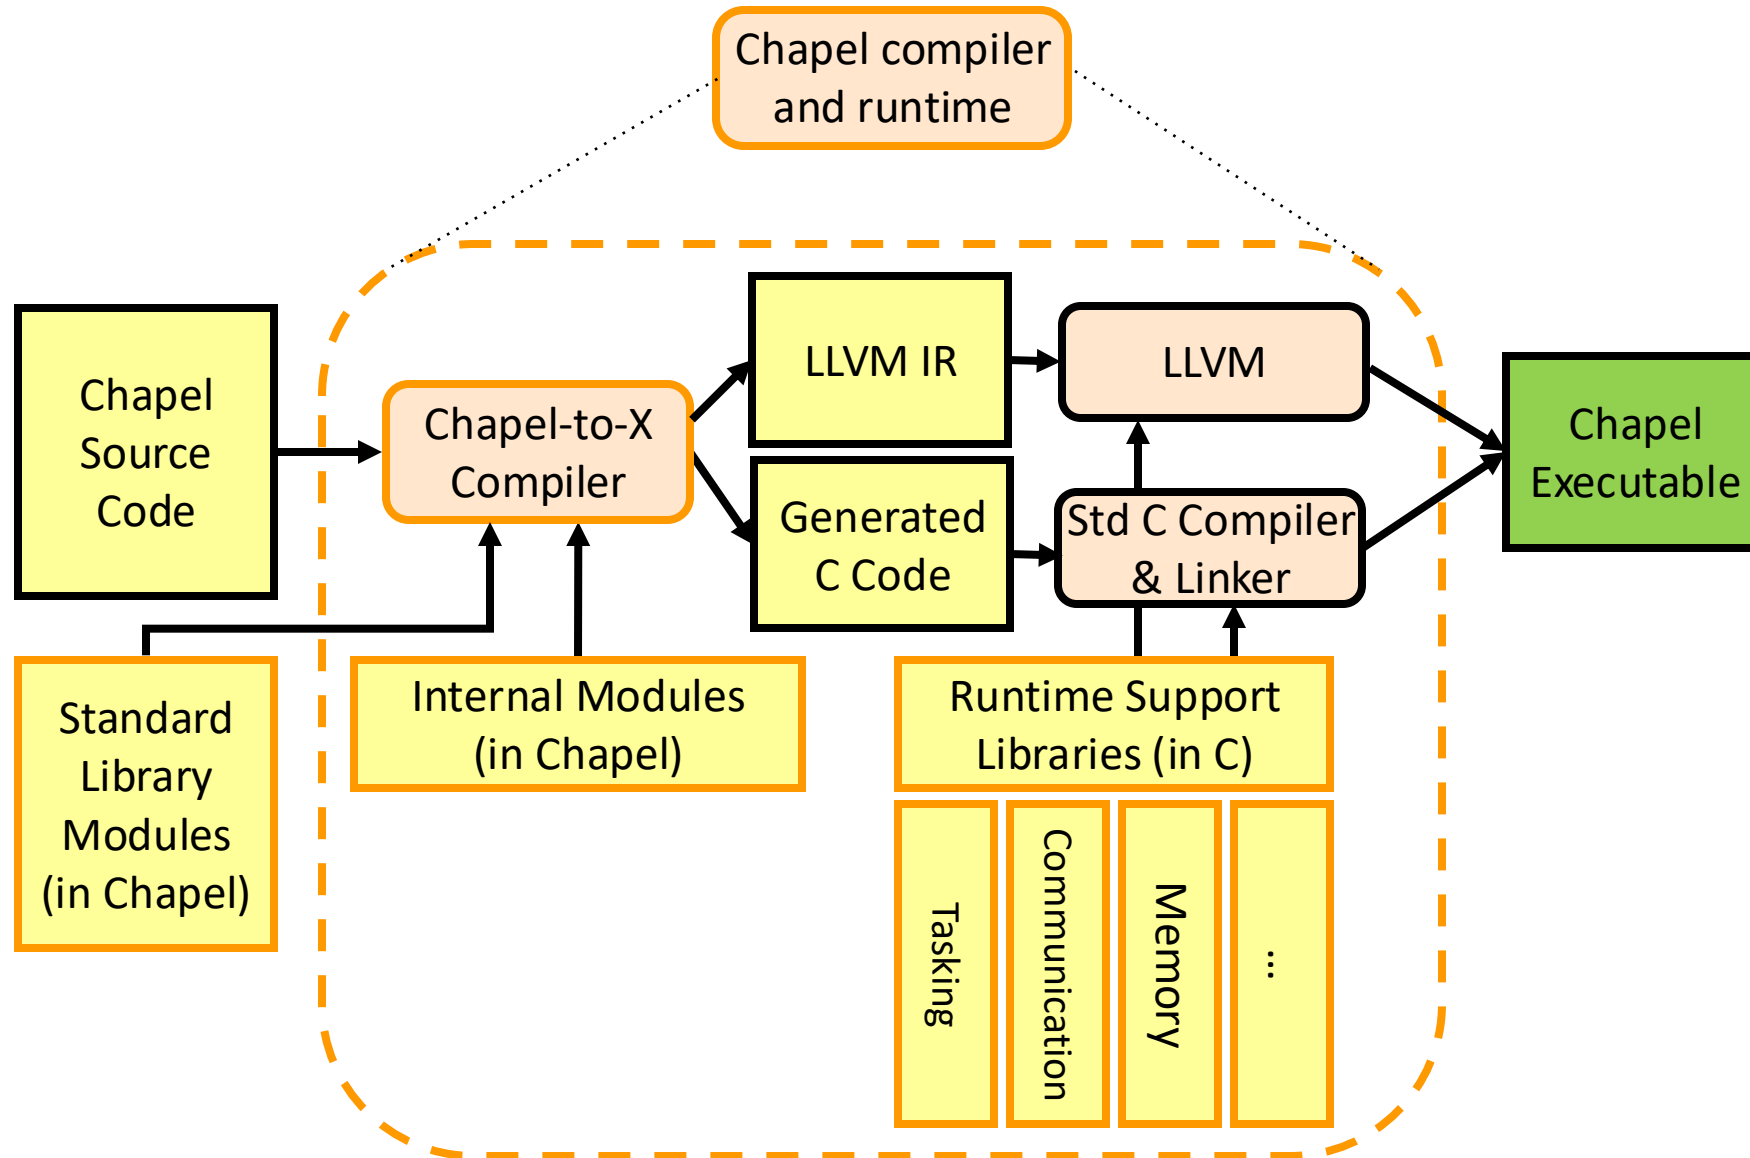
```
./01-hello-configurable -nl 1 --audience="y'all"
```



# COMPILING CHAPEL



# CHAPEL COMPILER ARCHITECTURE



# CHAPEL EXECUTION MODEL AND PARALLEL HELLO WORLD (HANDS ON)

# CHAPEL EXECUTION MODEL AND TERMINOLOGY: LOCALES


## Locales can run tasks and store variables

- Each locale executes on a “compute node” on a parallel system
- User specifies number of locales on executable’s command-line

```
prompt> ./myChapelProgram --numLocales=4 # or '-nl 4'
```

System has many nodes

**Locales array:**



User's code starts running as a single task on locale 0



# TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
    writef("Hello from task %n of %n on %s\n",  
          tid, numTasks, here.name);
```



# TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
    writef("Hello from task %n of %n on %s\n",  
          tid, numTasks, here.name);
```

‘here’ refers to the locale on which we’re currently running

how many concurrent tasks does this node support (typically the number of processor cores)?

what’s my locale’s name?



# TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
    writef("Hello from task %n of %n on %s\n",  
          tid, numTasks, here.name);
```

a 'coforall' loop executes each iteration as an independent task

```
> chpl 01-hello-dist-node-names.chpl  
> ./01-hello-dist-node-names -nl 1  
Hello from task 1 of 4 on n1032  
Hello from task 4 of 4 on n1032  
Hello from task 3 of 4 on n1032  
Hello from task 2 of 4 on n1032
```



# TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
    writef("Hello from task %n of %n on %s\n",  
          tid, numTasks, here.name);
```

```
> chpl 01-hello-dist-node-names.chpl  
> ./01-hello-dist-node-names -nl 1  
Hello from task 1 of 4 on n1032  
Hello from task 4 of 4 on n1032  
Hello from task 3 of 4 on n1032  
Hello from task 2 of 4 on n1032
```

**So far, this is a shared-memory program**

Nothing refers to remote locales,  
explicitly or implicitly

# TASK-PARALLEL “HELLO WORLD” (DISTRIBUTED VERSION)

01-hello-dist-node-names.chpl

```
coforall loc in Locales {  
  on loc {  
    const numTasks = here.maxTaskPar;  
    coforall tid in 1..numTasks do  
      writef("Hello from task %n of %n on %s\n",  
            tid, numTasks, here.name);  
  }  
}
```

the array of locales we're running  
on

**Locales array:**



# TASK-PARALLEL “HELLO WORLD” (DISTRIBUTED VERSION)

01-hello-dist-node-names.chpl

```
coforall loc in Locales {  
  on loc {  
    const numTasks = here.maxTaskPar;  
    coforall tid in 1..numTasks do  
      writef("Hello from task %n of %n on %s\n",  
            tid, numTasks, here.name);  
  }  
}
```

create a task per locale  
on which the program is running

have each task run 'on' its  
locale

then print a message per  
core,  
as before


```
> chpl 01-hello-dist-node-names.chpl  
> ./01-hello-dist-node-names -nl=4  
Hello from task 1 of 4 on n1032  
Hello from task 4 of 4 on n1032  
Hello from task 1 of 4 on n1034  
Hello from task 2 of 4 on n1032  
Hello from task 1 of 4 on n1033  
Hello from task 3 of 4 on n1034  
Hello from task 1 of 4 on n1035  
...
```

# LOCALES AND EXECUTION MODEL IN CHAPEL

In Chapel, a locale refers to a compute resource with...

- processors, so it can run tasks
- memory, so it can store variables


For now, think of each compute node as having one locale run on it



# LOCALES AND EXECUTION MODEL IN CHAPEL

Two key built-in variables for referring to locales in Chapel programs:

- **Locales:** An array of locale values representing the system resources on which the program is running
- **here:** The locale on which the current task is executing





# GETTING STARTED WITH LOCALES

- Users specify # of locales when running Chapel programs

```
% a.out --numLocales=8
```

```
% a.out -nl 8
```

- Chapel provides built-in locale variables

```
config const numLocales: int = ...;  
const Locales: [0..#numLocales] locale = ...;
```

- User's `main()` begins executing on locale #0, i.e. 'Locales[0]'



# LOCALE OPERATIONS

- Locale methods support queries about the target system:

```
proc locale.physicalMemory(...) { ... }  
proc locale.maxTaskPar { ... }  
proc locale.id { ... }  
proc locale.name { ... }
```

- *On-clauses* support placement of computations:


```
writeln("on locale 0");  
  
on Locales[1] do  
    writeln("now on locale 1");  
  
writeln("on locale 0 again");
```

```
on A[i,j] do  
    bigComputation(A);  
  
on node.left do  
    search(node.left);
```



# KEY CONCERNS FOR SCALABLE PARALLEL COMPUTING

1. **parallelism:** Which tasks should run simultaneously?
2. **locality:** Where should tasks run? Where should data be allocated?



# BASIC FEATURES FOR LOCALITY



01-basics-on.chpl

01-basics-on.chpl

```
writeln("Hello from locale ", here.id);  
  
var A: [1..2, 1..2] real;  
  
on Locales[1] {  
    var B: [1..2, 1..2] real;  
  
    B = 2 * A;  
}
```


All Chapel programs begin running as a single task on locale 0

Variables are stored using the memory local to the current task

on-clauses move tasks to other locales

remote variables can be accessed directly

This is a serial, but distributed computation



# BASIC FEATURES FOR LOCALITY




01-basics-for.chpl

01-basics-for.chpl

```
writeln("Hello from locale ", here.id);  
  
var A: [1..2, 1..2] real;  
  
for loc in Locales {  
  on loc {  
    var B = A;  
  }  
}
```

This loop will serially iterate over the program's locales

This is also a serial, but distributed computation




# MIXING LOCALITY WITH TASK PARALLELISM

01-basics-coforall.chpl

```
writeln("Hello from locale ", here.id);  
  
var A: [1..2, 1..2] real;  
  
coforall loc in Llocales {  
  on loc {  
    var B = A;  
  }  
}
```

The coforall loop creates a parallel task per iteration

This results in a parallel distributed computation



# ARRAY-BASED PARALLELISM AND LOCALITY



01-basics-distarr.chpl

01-basics-distarr.chpl

```
writeln("Hello from locale ", here.id);


var A: [1..2, 1..2] real;

use BlockDist;

var D = blockDist.createDomain({1..2, 1..2});
var B: [D] real;
B = A;
```

Chapel also supports distributed domains (index sets) and arrays

They also result in parallel distributed computation



# HANDS ON: PARALLELISM ACROSS AND WITHIN LOCALES

## Parallel hello world

- 01-hellopar.chpl

## Things to try

```
chpl 01-hellopar.chpl
./01-hellopar -nl 1 --tasksPerLocale=3
./01-hellopar -nl 2 --tasksPerLocale=3
```

## Key concepts

- 'coforall' over the `Locales` array with an `on` statement
- 'coforall' creating some number of tasks per locale
- configuration constants, 'config const'
- range expression, '0..<tasksPerLocale'
- 'writeln'
- inline comments start with '//'

```
// can be set on the command line with --
tasksPerLocale=2
config const tasksPerLocale = 1;

// parallel loops over nodes and then over threads
coforall loc in Locales do on loc {
    coforall tid in 0..<tasksPerLocale {

        writeln("Hello world! ",
                "(from task ", tid,
                " of ", tasksPerLocale,
                " on locale ", here.id,
                " of ", numLocales, ")");
    }
}
```



# PARALLELISM AND LOCALITY ARE ORTHOGONAL IN CHAPEL



01-parallelism-and-locality.chpl

- This is a parallel, but local program:

```
coforall i in 1..msgs do
  writeln("Hello from task ", i);
```

- This is a distributed, but serial program:

```
writeln("Hello from locale 0!");
on Locales[1] do writeln("Hello from locale 1!");
on Locales[2] {
  writeln("Hello from locale 2!");
  on Locales[0] do writeln("Hello from locale 0!");
}
writeln("Back on locale 0");
```

- This is a distributed parallel program:

```
coforall i in 1..msgs do
  on Locales[i%numLocales] do
    writeln("Hello from task ", i, " running on locale ", here.id);
```

# HANDS ON: PARALLELISM AND LOCALITY IN CHAPEL

## Goals

- Compile and run some of the examples from the last section
- Experiment some with '01-basics-distarr.chpl'

## Compile and run some of the other examples from the last section

```
chpl 01-parallelism-and-locality.chpl  
./01-parallelism-and-locality -nl 1  
./01-parallelism-and-locality -nl 4
```

## Experiment some with '01-basics-distarr.chpl'

1. what happens when you add a 'writeln(D)' to write out the domain 'D'?
2. what happens when you change 'D's initial value to '{0..3,0..3}'?
3. where does the computation on locales other than locale 0 happen?

