



Hewlett Packard
Enterprise

OCT 2023 CHAPEL TUTORIAL UPDATED FOR USE IN SPRING 2025 CS372 CLASS

Chapel Team, edited by Michelle Strout

March 27 and April 1, 2025

PLAN

- **Announcements**

- MT2 and LA2 grades have been posted
- SA6 is due Monday, survey for final project preferences
- Some reading assignment for Chapel have been posted

- **Last time**

- Midterm 2

- **Today**

- Chapel introduction
- Please start the docker pull for Chapel (see next slide)

HANDS ON: HOW TO DO THE HANDS ON



01-hello.chpl

Example codes for Chapel tutorial slides

- <https://github.com/UofA-CSc-372-Spring-2025/CSc372Spring2025-CourseMaterials/tree/main/Sandboxes/ChapelTutorialExamples>

Using a container on your laptop

- First, install docker for your machine and start it up (see the README.md for more info)
- Then, use the chapel-gasnet docker container

```
docker pull docker.io/chapel/chapel-gasnet      # takes about 5 minutes
cd CSc372Spring2025-CourseMaterials/Sandboxes/ChapelTutorial/
docker run --rm -it -v "$PWD":/workspace chapel/chapel-gasnet
root@589405d07f6a:/opt/chapel# cd /workspace
root@xxxxxxxx:/myapp# chpl 01-hello.chpl
root@xxxxxxxx:/myapp# ./01-hello -nl 1
```



OUTLINE: OVERVIEW OF PROGRAMMING IN CHAPEL

- Chapel Goals, Usage, and Comparison with other Tools
- Hello World (Hands On)
- Chapel Execution Model and Parallel Hello World (Hands On)
- kmer counting using file IO, config consts, strings, maps (Hands On)
- Parallelizing a program that processes files (Hands On)
- GPU programming support
- Learning goals for rest of tutorial



CHAPEL GOALS, USAGE, AND COMPARISON WITH OTHER TOOLS

CHAPEL PROGRAMMING LANGUAGE

Chapel is a general-purpose programming language that provides **ease of parallel programming, high performance, and portability.**

And is being used in applications in various ways:


refactoring existing codes,

developing new codes,

serving high performance to Python codes (**Chapel server with Python client**), and **providing distributed and shared memory parallelism** for existing codes.




APPLICATIONS OF CHAPEL: LINKS TO USERS' TALKS (SLIDES + VIDEO)




CHAMPS: 3D Unstructured CFD

[CHIUW 2021](#) [CHIUW 2022](#)




Arkouda: Interactive Data Science at Massive Scale

[CHIUW 2020](#) [CHIUW 2023](#)




ChOp: Chapel-based Optimization

[CHIUW 2021](#) [CHIUW 2023](#)




ChpUltra: Simulating Ultralight Dark Matter

[CHIUW 2020](#) [CHIUW 2022](#)



Lattice-Symmetries: a Quantum Many-Body Toolbox Desk dot chpl: Utilities for Environmental Eng.

[CHIUW 2022](#)




[CHIUW 2022](#)




RapidQ: Mapping Coral Biodiversity

[CHIUW 2023](#)




ChapQG: Layered Quasigeostrophic CFD




Chapel-based Hydrological Model Calibration

[CHIUW 2023](#)




CrayAI HyperParameter Optimization (HPO)

[CHIUW 2021](#)



CHGL: Chapel Hypergraph Library

[CHIUW 2020](#)




Your Application Here?

HIGHLIGHTS OF CHAPEL USAGE

CHAMPS: Computational Fluid Dynamics framework for airplane simulation

- Professor Eric Laurendeau's team at Polytechnique Montreal
- Performance: achieves competitive results w.r.t. established, world-class frameworks from Stanford, MIT, etc.
- Programmability: *"We ask students at the master's degree to do stuff that would take 2 years and they do it in 3 months."*



Arkouda: data analytics framework (<https://github.com/Bears-R-Us/arkouda>)

- Mike Merrill, Bill Reus, et al., US DOD
- Python front end client, Chapel server that processes dozens of terabytes in seconds
- April 2023: 1200 GiB/s for argsort on an HPE EX system



Other recent users

- Marjan Asgari et al, "Development of a knowledge-sharing parallel computing approach for calibrating distributed watershed hydrologic models", Environmental Modeling and Software.
- Scott Bachman has written some coral reef image analysis applications in Chapel.



CHAPEL IS HIGHLY PERFORMANT AND SCALABLE

HPE Apollo (May 2021)



- HDR-100 Infiniband network (100 Gb/s)
- 576 compute nodes
- 72 TiB of 8-byte values
- ~480 GiB/s (~150 seconds)

HPE Cray EX (April 2023)




- Slingshot-11 network (200 Gb/s)
- 896 compute nodes
- 28 TiB of 8-byte values
- ~1200 GiB/s (~24 seconds)

HPE Cray EX (May 2023)



- Slingshot-11 network (200 Gb/s)
- 8192 compute nodes
- 256 TiB of 8-byte values
- ~8500 GiB/s (~31 seconds)

Arkouda Argosort Performance



A notable performance achievement in ~100 lines of Chapel



COMPARE WITH OTHER PARALLEL PROGRAMMING MODELS

- Shared-memory parallelism

- Pthreads: low-level library for creating and managing threads
- OpenMP: pragmas added before loops and other statements
- Rust, Julia: programming languages with some threaded parallelism
- RAJA, Kokkos: C++ libraries that use template metaprogramming

- Distributed-memory parallelism and shared-memory parallelism

- MPI+X:

- MPI stands for message passing interface
- MPI is a library for sending and receiving messages between processes
- All processes allocate their own memory and run the same program
- There are many options for X: OpenMP, Pthreads, Python, Julia, R, etc.

- OpenSHMEM: library for implementing a partitioned global address space

- Spark: Python, Scala, and Java accessible library for especially iterative data processing

- Regent and Legion: programming language and runtime that implements implicit task parallelism

- Kokkos Remote Spaces: extends Kokkos C++ template views to distributed views

Chapel:

- shared memory parallelism,
- distributed-memory parallelism,
- data parallelism,
- task parallelism,
- map-reduce parallelism,
- vector parallelism,
- GPU parallelism, ...

All can be expressed in the same programming language.

COMPARE WITH OTHER PARALLEL PROGRAMMING MODELS (W/OUT CHAPEL BOX)

- Shared-memory parallelism
 - Pthreads: low-level library for creating and managing threads of execution that share memory
 - OpenMP: pragmas added before loops and other statements in C/C++/Fortran programs
 - Rust, Julia: programming languages with some threaded parallelism constructs
 - RAJA, Kokkos: C++ libraries that use template metaprogramming to express parallel policies
- Distributed-memory parallelism and shared-memory parallelism
 - MPI+X:
 - MPI stands for message passing interface
 - MPI is a library for sending and receiving messages between processes
 - All processes allocate their own memory and run the same program, SPMD: Single Program Multiple Data
 - There are many options for X: OpenMP, Pthreads, Python, Julia, RAJA, Kokkos, Chapel, ...
 - OpenSHMEM: library for implementing a partitioned global address space
 - Spark: Python, Scala, and Java accessible library for especially the map-reduce parallelism
 - Regent and Legion: programming language and runtime that implements implicit task parallelism
 - Kokkos Remote Spaces: extends Kokkos C++ template views to distributed views




LET'S COMPARE WITH STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures:




LET'S COMPARE WITH STREAM TRIAD: A PARALLEL COMPUTATION

Given: n -element vectors A, B, C

Compute: $\forall i \in 1..n, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory multicore, global-view):




STREAM TRIAD: IN MPI+OPENMP

```
#include <hpcc.h>
#ifdef _OPENMP
#include <omp.h>
#endif
```

```
static int VectorSize;
static double *a, *b, *c;
```

```
int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
```




```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to
            allocate memory (%d).\n",
                VectorSize );
        fclose( outFile );
    }
}
```

```
#define N 2000000
```

```
int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
```



*HPC suffers from too many distinct notations for expressing parallelism and locality.
This tends to be a result of bottom-up language design.*

```
rv = HPCC_Stream( params, 0, myRank );
MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
    0, comm );

return errCount;
}
```

```
int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;
```

```
VectorSize = HPCC_LocalVectorSize( params, 3,
    sizeof(double), 0 );
```

```
a = HPCC_XMALLOC( double, VectorSize );
b = HPCC_XMALLOC( double, VectorSize );
c = HPCC_XMALLOC( double, VectorSize );
```

```
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 1.0;
}
scalar = 3.0;
```

```
#ifdef _OPENMP
#pragma omp parallel for
#endif
```

```
for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];
```

```
HPCC_free(c);
HPCC_free(b);
HPCC_free(a);

return 0; }
```

```
STREAM_Triad<<<dimGrid,dimBlock>>>>(d_b, d_c, d_a, scalar, N);
cudaThreadSynchronize();
```

```
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
```

```
__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}
```

```
__global__ void STREAM_Triad( float *a, float *b, float *c,
    float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx]; }
```

WHY SO MANY PROGRAMMING MODELS?

HPC tends to approach programming models bottom-up:

Given a system and its core capabilities...

...provide features that permit users to access the available performance.

Type of HW Parallelism	Programming Model	Unit of Parallelism
Inter-node	MPI	executable
Intra-node/multicore	OpenMP / pthreads	iteration/task
Instruction-level vectors/threads	pragmas	iteration
GPU/accelerator	CUDA / Open[MP CL ACC]	SIMD function/task

benefits: lots of control; decent generality; easy to implement


downsides: lots of user-managed detail; brittle to changes



STREAM TRIAD: IN CHAPEL

The special sauce:
How should this index set—
and any arrays and
computations over it—be
mapped to the system?

```
use BlockDist;  
  
config const m = 1000,  
           alpha = 3.0;  
  
const ProblemSpace = blockDist.createDomain({1..m});  
  
var A, B, C: [ProblemSpace] real;  
  
B = 2.0;  
C = 1.0;  
  
A = B + alpha * C;
```



Philosophy: *Top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

HELLO WORLD (HANDS ON)



Example codes for Chapel tutorial slides

- <https://github.com/UofA-CSc-372-Spring-2025/CSc372Spring2025-CourseMaterials/tree/main/Sandboxes/ChapelTutorialExamples>

Using a container on your laptop

- First, install docker for your machine and start it up (see the README.md for more info)
- Then, use the chapel-gasnet docker container

```
docker pull docker.io/chapel/chapel-gasnet      # takes about 5 minutes
cd CSc372Spring2025-CourseMaterials/Sandboxes/ChapelTutorial/
docker run --rm -it -v "$PWD":/workspace chapel/chapel-gasnet
root@589405d07f6a:/opt/chapel# cd /workspace
root@xxxxxxxx:/myapp# chpl 01-hello.chpl
root@xxxxxxxx:/myapp# ./01-hello -nl 1
```



"HELLO WORLD" IN CHAPEL: TWO VERSIONS

- Fast prototyping

```
writeln("Hello, world!");
```



01-hello.chpl

- “Production-grade”

```
module Hello {  
  
    proc main() {  
        writeln("Hello, world!");  
    }  
  
}
```



01-hello-production.chpl



"HELLO WORLD" IN CHAPEL: TWO VERSIONS

- Fast prototyping (configurable)

```
config const audience = "world";  
writeln("Hello, ", audience, "!");
```

 01-hello-configurable.chpl

- “Production-grade” (configurable)

```
module Hello {  
    config const audience = "world";  
  
    proc main() {  
        writeln("Hello, ", audience, "!");  
    }  
}
```

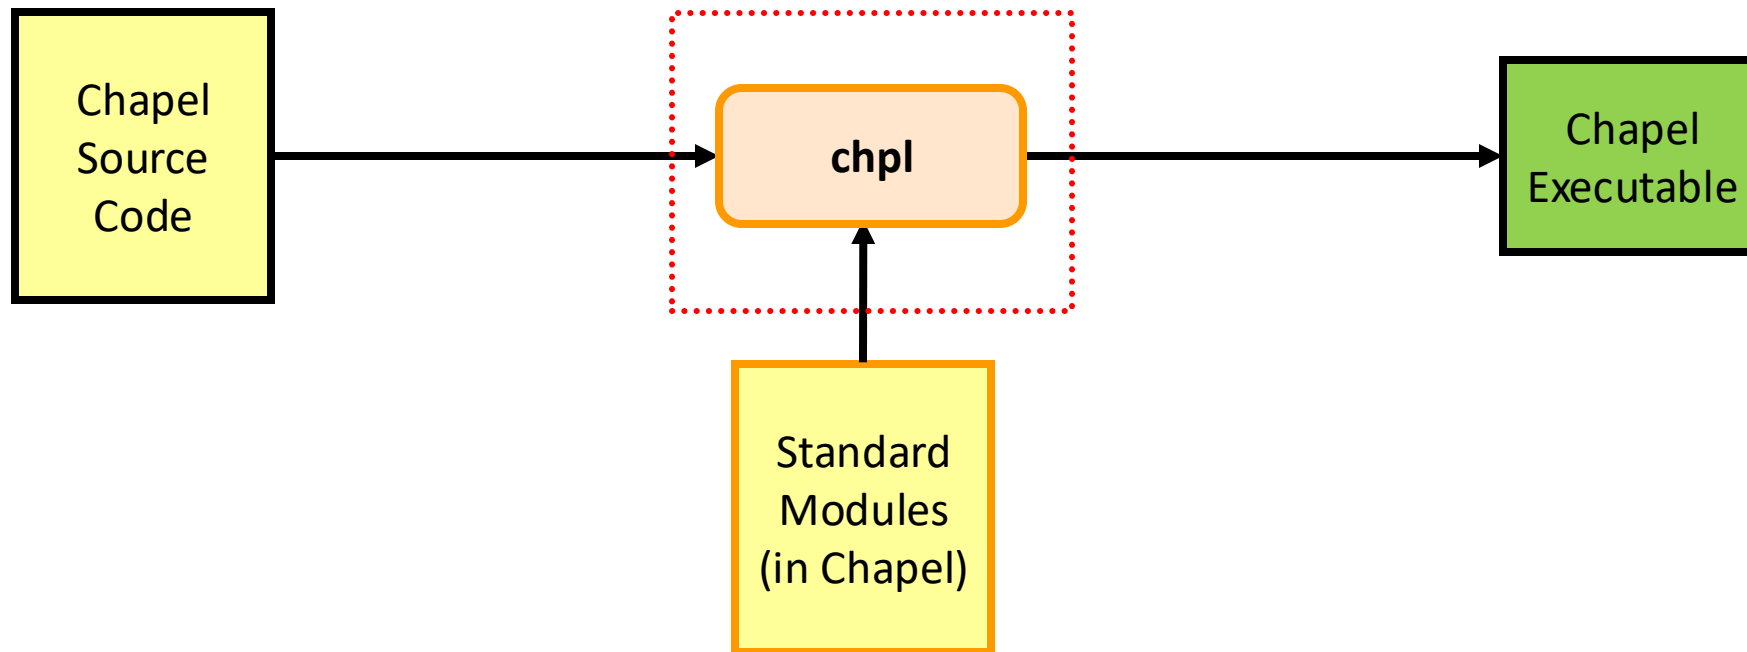
 01-hello-production-configurable.chpl

- To change ‘audience’ for a given run:

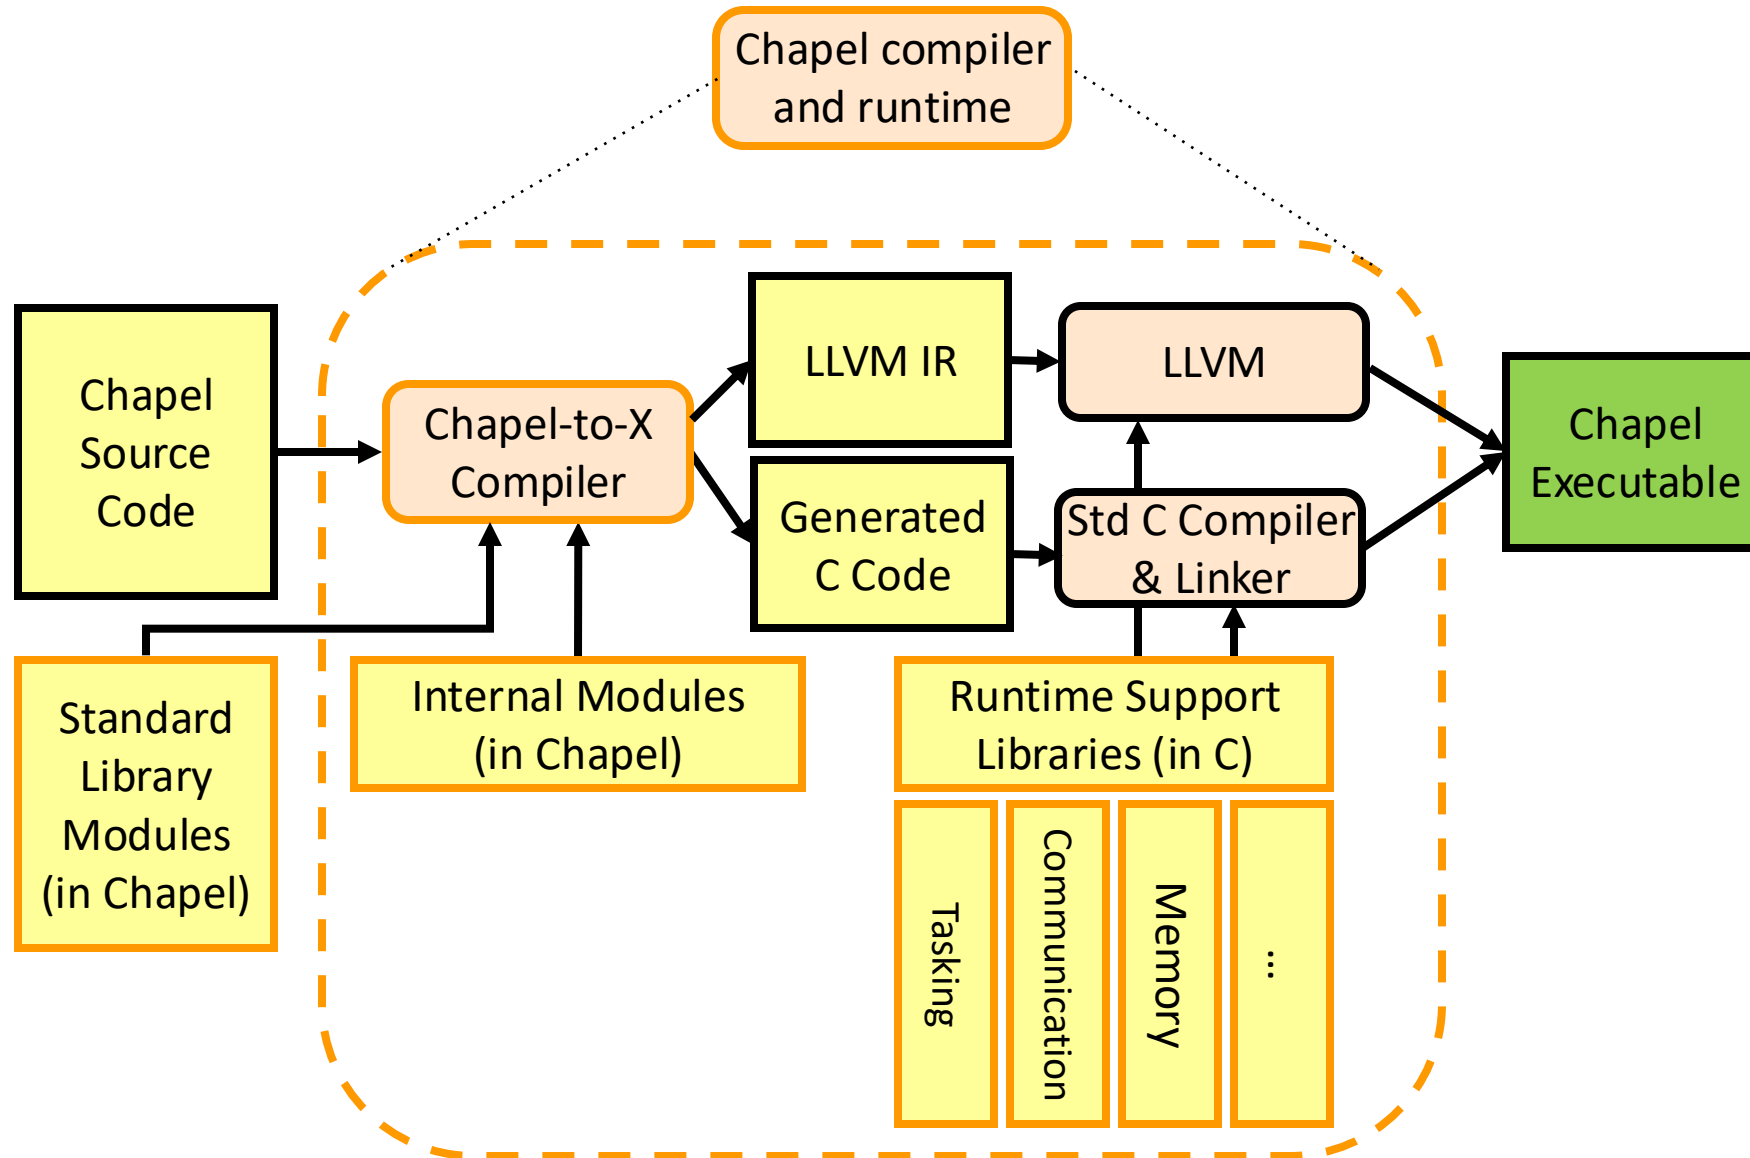
```
./01-hello-configurable -nl 1 --audience="y'all"
```



COMPILING CHAPEL



CHAPEL COMPILER ARCHITECTURE



CHAPEL EXECUTION MODEL AND PARALLEL HELLO WORLD (HANDS ON)

CHAPEL EXECUTION MODEL AND TERMINOLOGY: LOCALES


Locales can run tasks and store variables

- Each locale executes on a “compute node” on a parallel system
- User specifies number of locales on executable’s command-line

```
prompt> ./myChapelProgram --numLocales=4 # or '-nl 4'
```

System has many nodes

Locales array:



User's code starts running as a single task on locale 0

TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
    writef("Hello from task %n of %n on %s\n",  
          tid, numTasks, here.name);
```



TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
  writef("Hello from task %n of %n on %s\n",  
        tid, numTasks, here.name);
```

‘here’ refers to the locale on which we’re currently running

how many concurrent tasks does this node support (typically the number of processor cores)?

what’s my locale’s name?



TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
    writef("Hello from task %n of %n on %s\n",  
          tid, numTasks, here.name);
```

a 'coforall' loop executes each iteration as an independent task

```
> chpl 01-hello-dist-node-names.chpl  
> ./01-hello-dist-node-names -nl 1  
Hello from task 1 of 4 on n1032  
Hello from task 4 of 4 on n1032  
Hello from task 3 of 4 on n1032  
Hello from task 2 of 4 on n1032
```



TASK-PARALLEL “HELLO WORLD”

01-hello-dist-node-names.chpl

```
const numTasks = here.maxTaskPar;  
coforall tid in 1..numTasks do  
    writef("Hello from task %n of %n on %s\n",  
          tid, numTasks, here.name);
```

```
> chpl 01-hello-dist-node-names.chpl  
> ./01-hello-dist-node-names -nl 1  
Hello from task 1 of 4 on n1032  
Hello from task 4 of 4 on n1032  
Hello from task 3 of 4 on n1032  
Hello from task 2 of 4 on n1032
```

So far, this is a shared-memory program

Nothing refers to remote locales,
explicitly or implicitly

TASK-PARALLEL “HELLO WORLD” (DISTRIBUTED VERSION)

01-hello-dist-node-names.chpl

```
coforall loc in Locales {  
  on loc {  
    const numTasks = here.maxTaskPar;  
    coforall tid in 1..numTasks do  
      writef("Hello from task %n of %n on %s\n",  
            tid, numTasks, here.name);  
  }  
}
```

the array of locales we're running
on

Locales array:



TASK-PARALLEL “HELLO WORLD” (DISTRIBUTED VERSION)

01-hello-dist-node-names.chpl

```
coforall loc in Locales {  
  on loc {  
    const numTasks = here.maxTaskPar;  
    coforall tid in 1..numTasks do  
      writef("Hello from task %n of %n on %s\n",  
            tid, numTasks, here.name);  
  }  
}
```

create a task per locale
on which the program is running

have each task run 'on' its
locale

then print a message per
core,
as before


```
> chpl 01-hello-dist-node-names.chpl  
> ./01-hello-dist-node-names -nl=4  
Hello from task 1 of 4 on n1032  
Hello from task 4 of 4 on n1032  
Hello from task 1 of 4 on n1034  
Hello from task 2 of 4 on n1032  
Hello from task 1 of 4 on n1033  
Hello from task 3 of 4 on n1034  
Hello from task 1 of 4 on n1035  
...
```

LOCALES AND EXECUTION MODEL IN CHAPEL

In Chapel, a locale refers to a compute resource with...

- processors, so it can run tasks
- memory, so it can store variables


For now, think of each compute node as having one locale run on it



LOCALES AND EXECUTION MODEL IN CHAPEL

Two key built-in variables for referring to locales in Chapel programs:

- **Locales:** An array of locale values representing the system resources on which the program is running
- **here:** The locale on which the current task is executing



GETTING STARTED WITH LOCALES

- Users specify # of locales when running Chapel programs

```
% a.out --numLocales=8
```

```
% a.out -nl 8
```

- Chapel provides built-in locale variables

```
config const numLocales: int = ...;  
const Locales: [0..#numLocales] locale = ...;
```

- User's `main()` begins executing on locale #0, i.e. 'Locales[0]'



LOCALE OPERATIONS

- Locale methods support queries about the target system:

```
proc locale.physicalMemory(...) { ... }  
proc locale.maxTaskPar { ... }  
proc locale.id { ... }  
proc locale.name { ... }
```

- *On-clauses* support placement of computations:


```
writeln("on locale 0");  
  
on Locales[1] do  
    writeln("now on locale 1");  
  
writeln("on locale 0 again");
```

```
on A[i,j] do  
    bigComputation(A);  
  
on node.left do  
    search(node.left);
```



KEY CONCERNS FOR SCALABLE PARALLEL COMPUTING

1. **parallelism:** Which tasks should run simultaneously?
2. **locality:** Where should tasks run? Where should data be allocated?



BASIC FEATURES FOR LOCALITY



01-basics-on.chpl

01-basics-on.chpl

```
writeln("Hello from locale ", here.id);  
  
var A: [1..2, 1..2] real;  
  
on Locales[1] {  
    var B: [1..2, 1..2] real;  
  
    B = 2 * A;  
}
```


All Chapel programs begin running as a single task on locale 0

Variables are stored using the memory local to the current task

on-clauses move tasks to other locales

remote variables can be accessed directly

This is a serial, but distributed computation



BASIC FEATURES FOR LOCALITY




01-basics-for.chpl

01-basics-for.chpl

```
writeln("Hello from locale ", here.id);  
  
var A: [1..2, 1..2] real;  
  
for loc in Locales {  
  on loc {  
    var B = A;  
  }  
}
```

This loop will serially iterate over the program's locales

This is also a serial, but distributed computation




MIXING LOCALITY WITH TASK PARALLELISM

01-basics-coforall.chpl

```
writeln("Hello from locale ", here.id);  
  
var A: [1..2, 1..2] real;  
  
coforall loc in Llocales {  
  on loc {  
    var B = A;  
  }  
}
```

The coforall loop creates a parallel task per iteration

This results in a parallel distributed computation



ARRAY-BASED PARALLELISM AND LOCALITY

01-basics-distarr.chpl

```
writeln("Hello from locale ", here.id);


var A: [1..2, 1..2] real;

use BlockDist;

var D = blockDist.createDomain({1..2, 1..2});
var B: [D] real;
B = A;
```

Chapel also supports distributed domains (index sets) and arrays

They also result in parallel distributed computation



HANDS ON: PARALLELISM ACROSS AND WITHIN LOCALES

 01-hellopar.chpl

Parallel hello world

- 01-hellopar.chpl

Things to try

```
chpl 01-hellopar.chpl
./01-hellopar -nl 1 --tasksPerLocale=3
./01-hellopar -nl 2 --tasksPerLocale=3
```

Key concepts

- 'coforall' over the `Locales` array with an `on` statement
- 'coforall' creating some number of tasks per locale
- configuration constants, 'config const'
- range expression, '0..<tasksPerLocale'
- 'writeln'
- inline comments start with '//'

```
// can be set on the command line with --
tasksPerLocale=2
config const tasksPerLocale = 1;

// parallel loops over nodes and then over threads
coforall loc in Locales do on loc {
    coforall tid in 0..<tasksPerLocale {

        writeln("Hello world! ",
                "(from task ", tid,
                " of ", tasksPerLocale,
                " on locale ", here.id,
                " of ", numLocales, ")");
    }
}
```


PARALLELISM AND LOCALITY ARE ORTHOGONAL IN CHAPEL



01-parallelism-and-locality.chpl

- This is a parallel, but local program:

```
coforall i in 1..msgs do
  writeln("Hello from task ", i);
```

- This is a distributed, but serial program:

```
writeln("Hello from locale 0!");
on Locales[1] do writeln("Hello from locale 1!");
on Locales[2] {
  writeln("Hello from locale 2!");
  on Locales[0] do writeln("Hello from locale 0!");
}
writeln("Back on locale 0");
```

- This is a distributed parallel program:

```
coforall i in 1..msgs do
  on Locales[i%numLocales] do
    writeln("Hello from task ", i, " running on locale ", here.id);
```

HANDS ON: PARALLELISM AND LOCALITY IN CHAPEL

Goals

- Compile and run some of the examples from the last section
- Experiment some with '01-basics-distarr.chpl'

Compile and run some of the other examples from the last section

```
chpl 01-parallelism-and-locality.chpl  
./01-parallelism-and-locality -nl 1  
./01-parallelism-and-locality -nl 4
```

Experiment some with '01-basics-distarr.chpl'

1. what happens when you add a 'writeln(D)' to write out the domain 'D'?
2. what happens when you change 'D's initial value to '{0..3,0..3}'?
3. where does the computation on locales other than locale 0 happen?

