



OCT 2023 CHAPEL TUTORIAL UPDATED FOR USE IN SPRING 2025 CS372 CLASS



Chapel Team, edited by Michelle Strout

April 1, 2025

PLAN

- **Announcements**

- Final project assignments coming out ASAP
- SA7 will be posted Wednesday April 2nd or Thursday April 3rd

- **Last time**

- Chapel introduction
- Please start the docker pull for Chapel (see next slide)

- **Today**

- TopHat questions about ChapelCon tutorial and last week's class, aka ICA10 prep
- Kmer counting example: file IO, maps, strings
- Parallel processing of files
- Overview of GPU programming support

HANDS ON: HOW TO DO THE HANDS ON



01-hello.chpl

Example codes for Chapel tutorial slides

- <https://github.com/UofA-CSc-372-Spring-2025/CSc372Spring2025-CourseMaterials/tree/main/Sandboxes/ChapelTutorialExamples>

Using a container on your laptop

- First, install docker for your machine and start it up (see the README.md for more info)
- Then, use the chapel-gasnet docker container

```
docker pull docker.io/chapel/chapel-gasnet      # takes about 5 minutes
cd CSc372Spring2025-CourseMaterials/Sandboxes/ChapelTutorialExamples/
docker run --rm -it -v "$PWD":/workspace chapel/chapel-gasnet
root@589405d07f6a:/opt/chapel# cd /workspace
root@xxxxxxxx:/myapp# chpl 01-hello.chpl
root@xxxxxxxx:/myapp# ./01-hello -nl 1
```



HANDS ON: PARALLELISM AND LOCALITY IN CHAPEL

Goals

- Experiment some with '01-basics-distarr.chpl'

```
chpl 01-basics-distarrchpl.chpl
./01-basics-distarr -nl 1
./01-basics-distarr -nl 4
```

Experiment some with '01-basics-distarr.chpl'

1. What happens when you add a 'writeln(D)' to write out the domain 'D'?
2. What happens when you change 'D's initial value to '{0..3,0..3}'?
3. Use a config const for the upper bound of the domain 'D'.
4. Have array A and array B use the same domain.



OUTLINE: OVERVIEW OF PROGRAMMING IN CHAPEL

- Chapel Goals, Usage, and Comparison with other Tools
- Hello World (Hands On)
- Chapel Execution Model and Parallel Hello World (Hands On)
- kmer counting using file IO, config consts, strings, maps (Hands On)
- Parallelizing a program that processes files (Hands On)
- GPU programming support
- Learning goals for rest of Chapel unit



KMER COUNTING USING FILE IO, CONFIG CONSTS, AND STRINGS (HANDS ON)

SERIAL CODE USING MAP/Dictionary: K-MER COUNTING



kmer.chpl

kmer.chpl

```
use Map, IO;

config const infilename = "kmer_large_input.txt";
config const k = 4;

var sequence, line : string;
var f = open(infilename, ioMode.r);
var infile = f.reader();
while infile.readLine(line) {
    sequence += line.strip();
}

var nkmerCounts : map(string, int);

for ind in 0..<(sequence.size-k) {
    nkmerCounts[sequence[ind..#k]] += 1;
}
```

'Map' and 'IO' are two of the standard libraries provided in Chapel. A 'map' is like a dictionary in python.

'config const' indicates a configuration constant, which result in built-in command-line parsing

Reading all of the lines from the input file into the string 'sequence'.

The variable 'nkmerCounts' is being declared as a dictionary mapping strings to ints

Counting up each kmer in the sequence



Some things to try out with 'kmer.chpl'

```
chpl kmer.chpl
./kmer -nl 1
./kmer -nl 1 --k=10 # can change k
./kmer -nl 1 --infilename="kmer.chpl" # changing infilename
./kmer -nl 1 --k=10 --infilename="kmer.chpl" # can change both
```

Experiment some with kmer.chpl

1. When k=5, what is the most common kmer in the file from doing “wget https://www.bioinformatics.nl/tools/crab_fasta.html”
2. When k=8?

Key concepts

- 'use' command for including modules
- configuration constants, 'config const'
- reading from a file
- 'map' data structure



PARALLELIZING A PROGRAM THAT PROCESSES FILES (HANDS ON)

ANALYZING MULTIPLE FILES USING PARALLELISM

 parfilekmer.chpl

parfilekmer.chpl

```
use FileSystem, BlockDist;
config const dir = "DataDir";
var fList = findFiles(dir);
var filenames =
    blockDist.createArray(0..<fList.size, string);
filenames = fList;

// per file word count
forall f in filenames {
    ...
    // code from kmer.chpl
    ...
}
```

```
prompt> chpl --fast parfilekmer.chpl
prompt> ./parfilekmer -nl 1
prompt> ./parfilekmer -nl 4
```

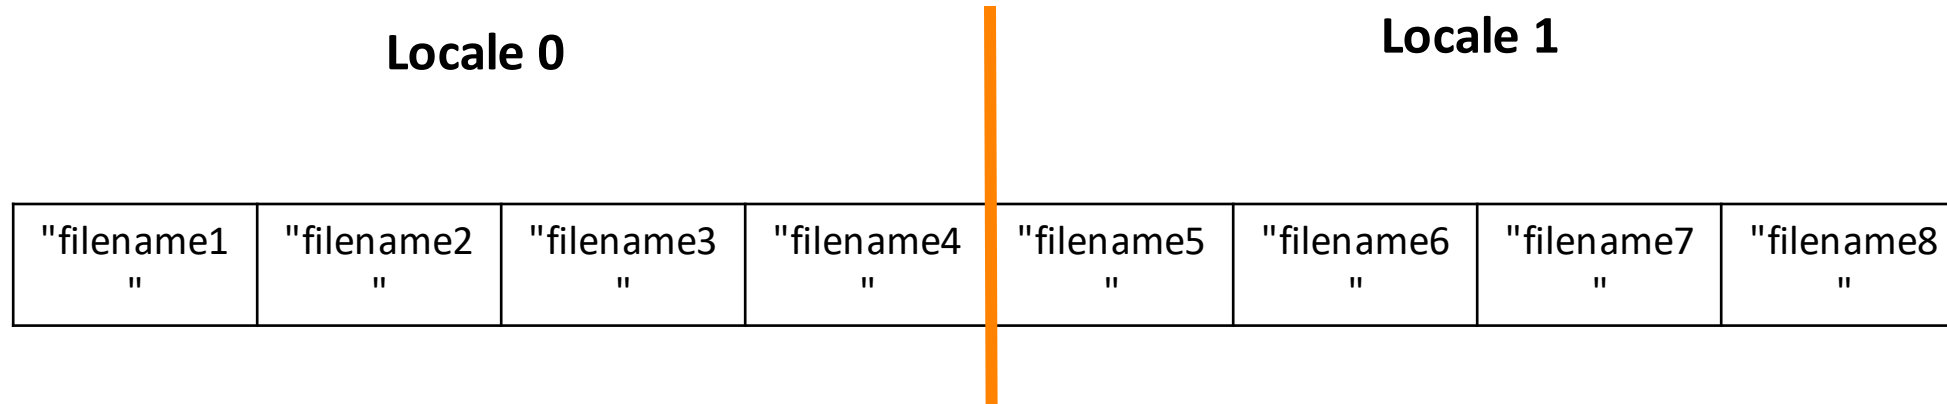
- shared and distributed-memory parallelism using 'forall'
 - in other words, parallelism within the locale/node and across locales/nodes
- a distributed array
- command line options to indicate number of locales

Experiment some with parfilekmer.chpl

1. Using 'writeln', edit parfilekmer.chpl so that 'fList' is printed to the screen.
2. Now print out the 'filenames' array.



BLOCK DISTRIBUTION OF ARRAY OF STRINGS



- Array of strings for filenames is distributed across locales
- 'forall' will do parallelism across locales and then within each locale to take advantage of multicore

HANDS ON: PROCESSING FILES IN PARALLEL

 parfilekmer.chpl

Some things to try out with 'parfilekmer.chpl'

```
chpl parfilekmer.chpl --fast
./parfilekmer -nl 2 --dir="SomethingElse/"      # change dir with inputs files

./parfilekmer -nl 2 --k=10                       # can also change k
```

Concepts illustrated

- 'forall' provides distributed and shared memory parallelism when do a 'forall' over the Block distributed array
- No remote puts and gets



GPU PROGRAMMING SUPPORT

GPU SUPPORT IN CHAPEL

Generate code for GPUs

- Support for NVIDIA and AMD GPUs
- Exploring Intel support

Key concepts

- Using the 'locale' concept to indicate execution and data allocation on GPUs
- 'forall' and 'foreach' loops are converted to kernels
- Arrays declared within GPU sublocale code blocks are allocated on the GPU

Chapel code calling CUDA examples

- <https://github.com/chapel-lang/chapel/blob/main/test/gpu/interop/stream/streamChpl.chpl>
- <https://github.com/chapel-lang/chapel/blob/main/test/gpu/interop/cuBLAS/cuBLAS.chpl>

For more info...

– <https://chapel-lang.org/docs/technotes/gpu.html>

gpuExample.chp

```
use GpuDiagnostics;
startGpuDiagnostics();

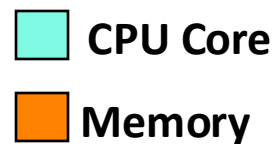

var operateOn =
if here.gpus.size>0 then here.gpus
    else [here,];

// Same code can run on GPU or CPU
coforall loc in operateOn do on loc {
    var A : [1..10] int;
    foreach a in A do a+=1;
    writeln(A);
}

stopGpuDiagnostics();
writeln(getGpuDiagnostics());
```

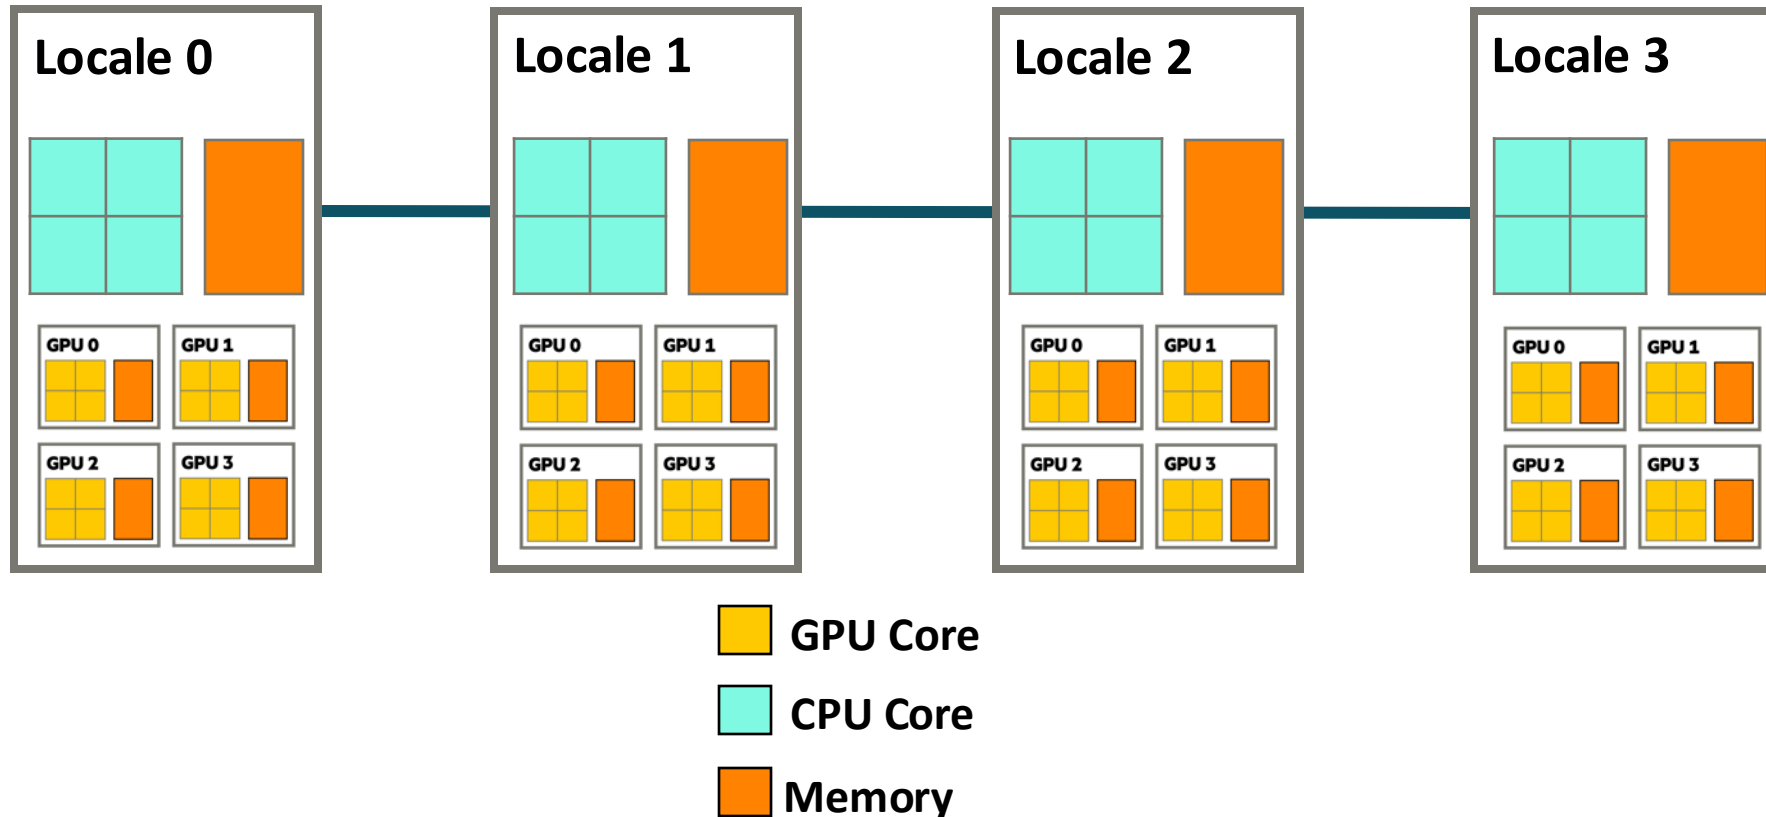
KEY CONCERNS FOR SCALABLE PARALLEL COMPUTING

1. **parallelism:** Which tasks should run simultaneously?
2. **locality:** Where should tasks run? Where should data be allocated?
 - complicating matters, compute nodes now often have GPUs with their own processors and memory



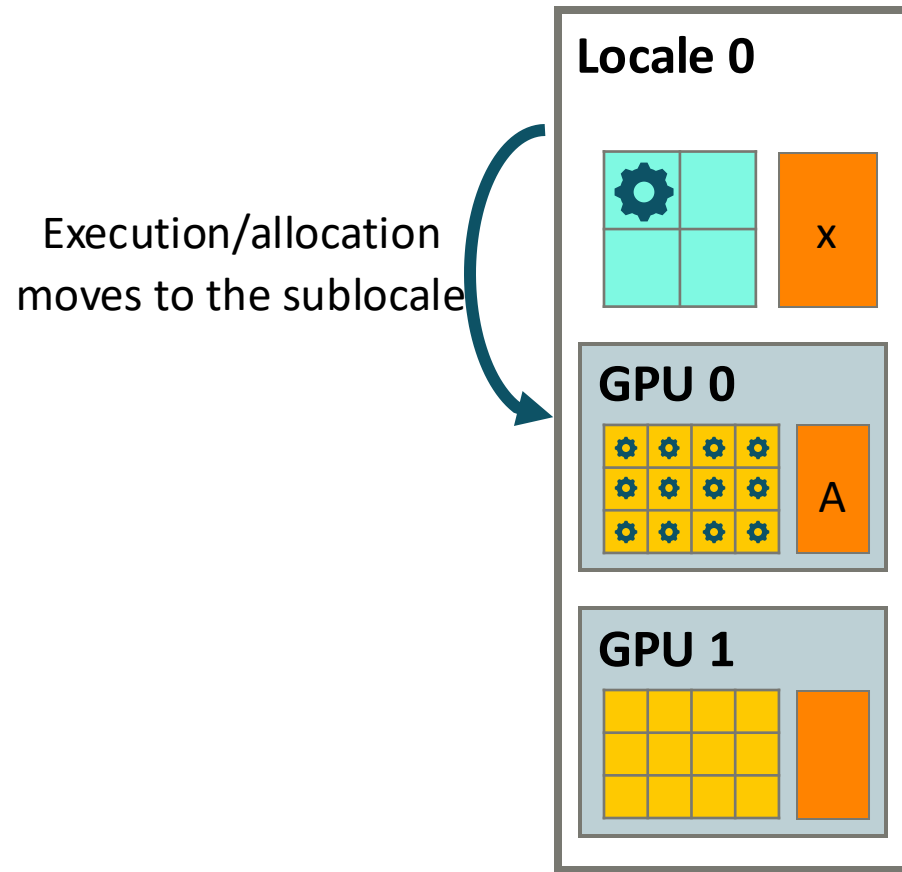
KEY CONCERNS FOR SCALABLE PARALLEL COMPUTING


1. **parallelism:** Which tasks should run simultaneously?
2. **locality:** Where should tasks run? Where should data be allocated?
 - complicating matters, compute nodes now often have GPUs with their own processors and memory
 - we represent these as *sub-locales* in Chapel






PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

 CPU Core  GPU Core  Memory



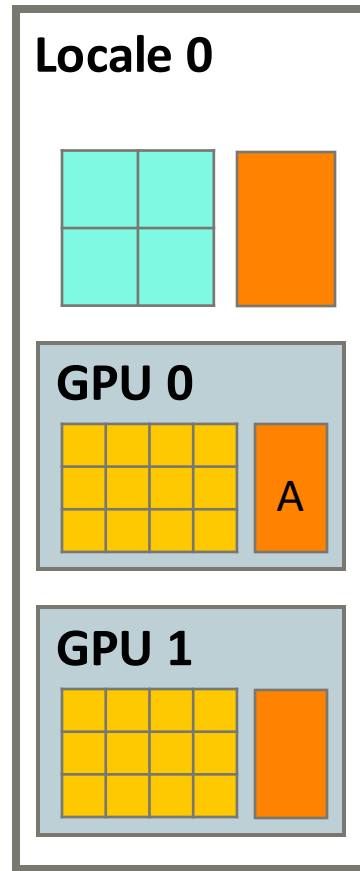
 `var x = 10;`



 `on here.gpus[0] {`
 `var A = [1, 2, 3, 4, 5, ...];`
 `foreach a in A do a += 1;`
}

 `writeln(x);`

PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

 CPU Core  GPU Core  Memory



```
var x = 10;
```

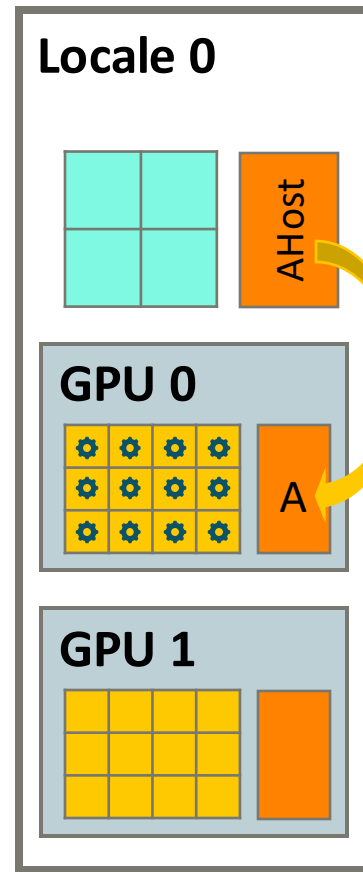


```
on here.gpus[0] {  
  var A = [1, 2, 3, 4, 5, ...];  
  foreach a in A do a += 1;  
}
```

```
writeln(x);
```

PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

■ CPU Core ■ GPU Core ■ Memory



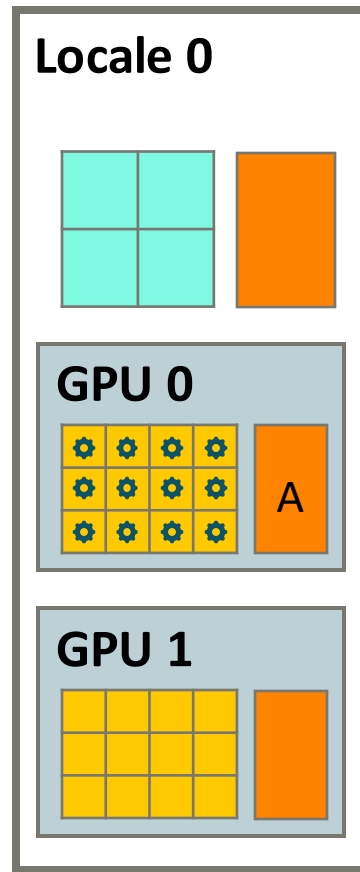
```
var x = 10;  
var AHost = [1, 2, 3, 4, 5, ...];
```

```
on here.gpus[0] {  
  var A = AHost;  
  foreach a in A do a += 1;  
}
```

```
writeln(x);
```

PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

 CPU Core  GPU Core  Memory



```
var x = 10;
```

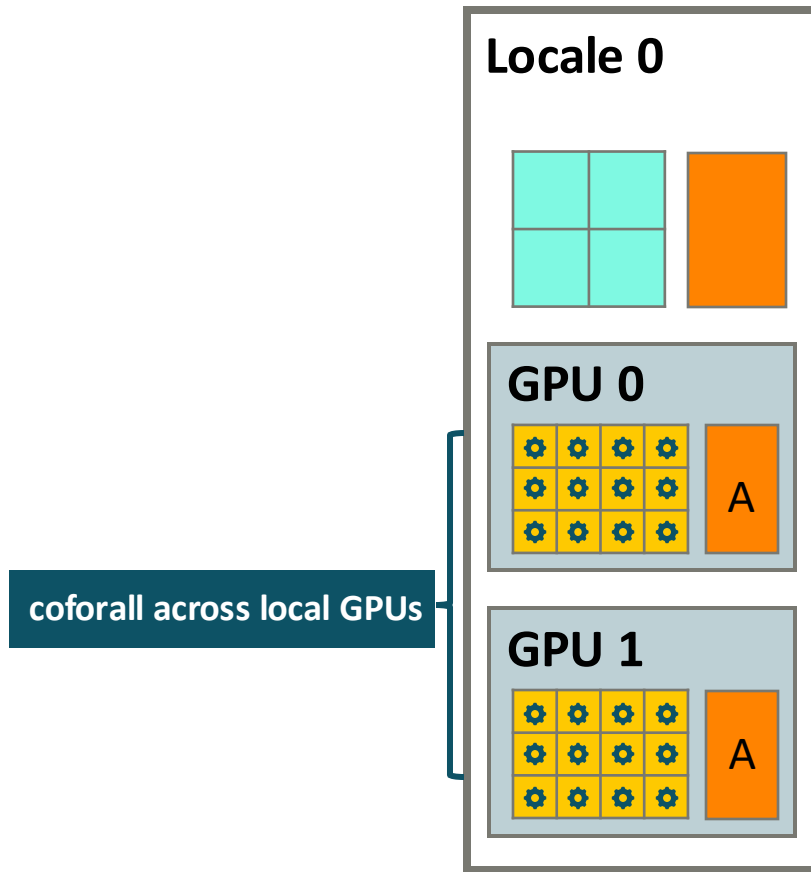


```
on here.gpus[0] {  
  var A = [1, 2, 3, 4, 5, ...];  
  foreach a in A do a += 1;  
}
```

```
writeln(x);
```

PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

 CPU Core  GPU Core  Memory



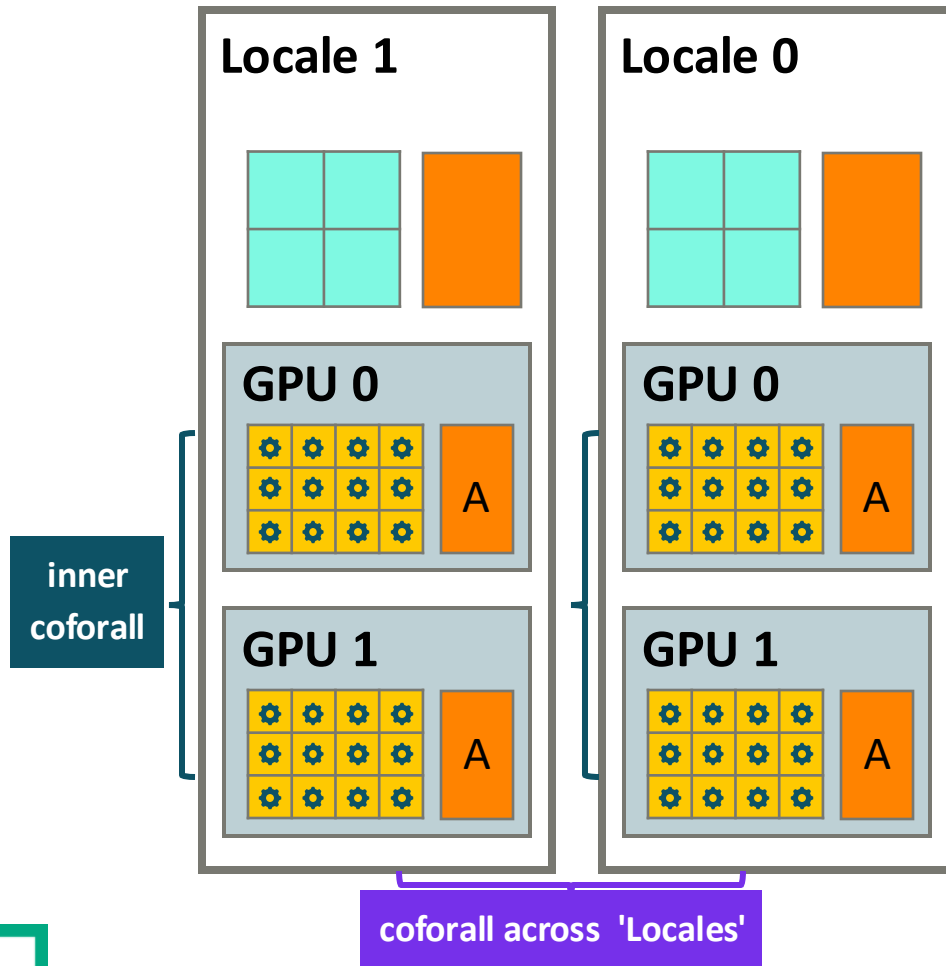
```
var x = 10;
```

```
coforall g in here.gpus do on g {  
  var A = [1, 2, 3, 4, 5, ...];  
  foreach a in A do a += 1;  
}
```

```
writeln(x);
```

PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

 CPU Core  GPU Core  Memory



```
var x = 10;
```

```
coforall l in Locales do on l {
```

```
coforall g in here.gpus do on g {
```

```
  var A = [1, 2, 3, 4, 5, ...];
```

```
  foreach a in A do a += 1;
```

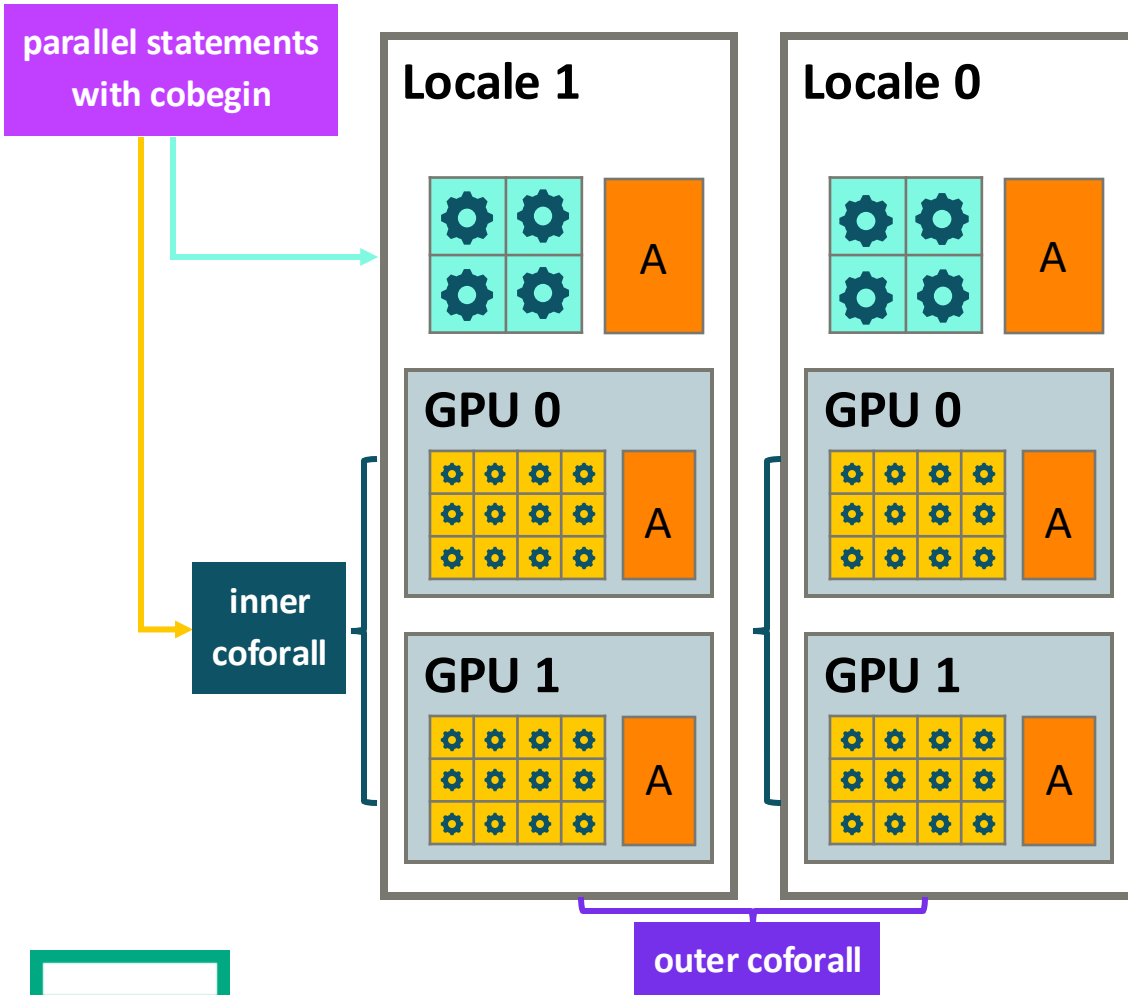
```
}
```

```
}
```

```
writeln(x);
```

PARALLELISM AND LOCALITY IN THE CONTEXT OF GPUS

■ CPU Core ■ GPU Core ■ Memory



```
var x = 10;  
coforall l in Locales do on l {  
  cobegin {  
    coforall g in here.gpus do on g {  
      var A = [1, 2, 3, 4, 5, ...];  
      foreach a in A do a += 1;  
    }  
    {  
      var A = [1, 2, 3, 4, 5, ...];  
      foreach a in A do a += 1;  
    }  
  }  
}  
writeln(x);
```

LEARNING OBJECTIVES FOR THE REST OF THE CHAPEL UNIT

LEARNING OBJECTIVES FOR CHAPEL UNIT

- Familiarity with the Chapel execution model including how to run codes in parallel on a single node, across nodes, and both
- Learn Chapel concepts by compiling and running provided code examples
 - ✓ Serial code using map/dictionary, (k-mer counting from bioinformatics)
 - ✓ Parallelism and locality in Chapel
 - ✓ Distributed parallelism and 1D arrays, (processing files in parallel)
- Chapel basics in the context of an n-body code
- Distributed parallelism and 2D arrays, (heat diffusion problem)
- How to parallelize histogram
- Using CommDiagnostics for counting remote reads and writes
- Chapel and Arkouda best practices including avoiding races and performance gotchas
- Where to get help and how you can participate in the Chapel community
- Memory safety in Chapel and other languages like Rust

OTHER CHAPEL EXAMPLES & PRESENTATIONS

Primers

- <https://chapel-lang.org/docs/primers/index.html>

Blog posts for Advent of Code

- <https://chapel-lang.org/blog/index.html>

Examples people have written

- <https://chapel-lang.org/examples/>

Test directory in main repository

- <https://github.com/chapel-lang/chapel/tree/main/test>

Presentations

- <https://chapel-lang.org/presentations.html>



CHAPEL RESOURCES

Chapel homepage: <https://chapel-lang.org>

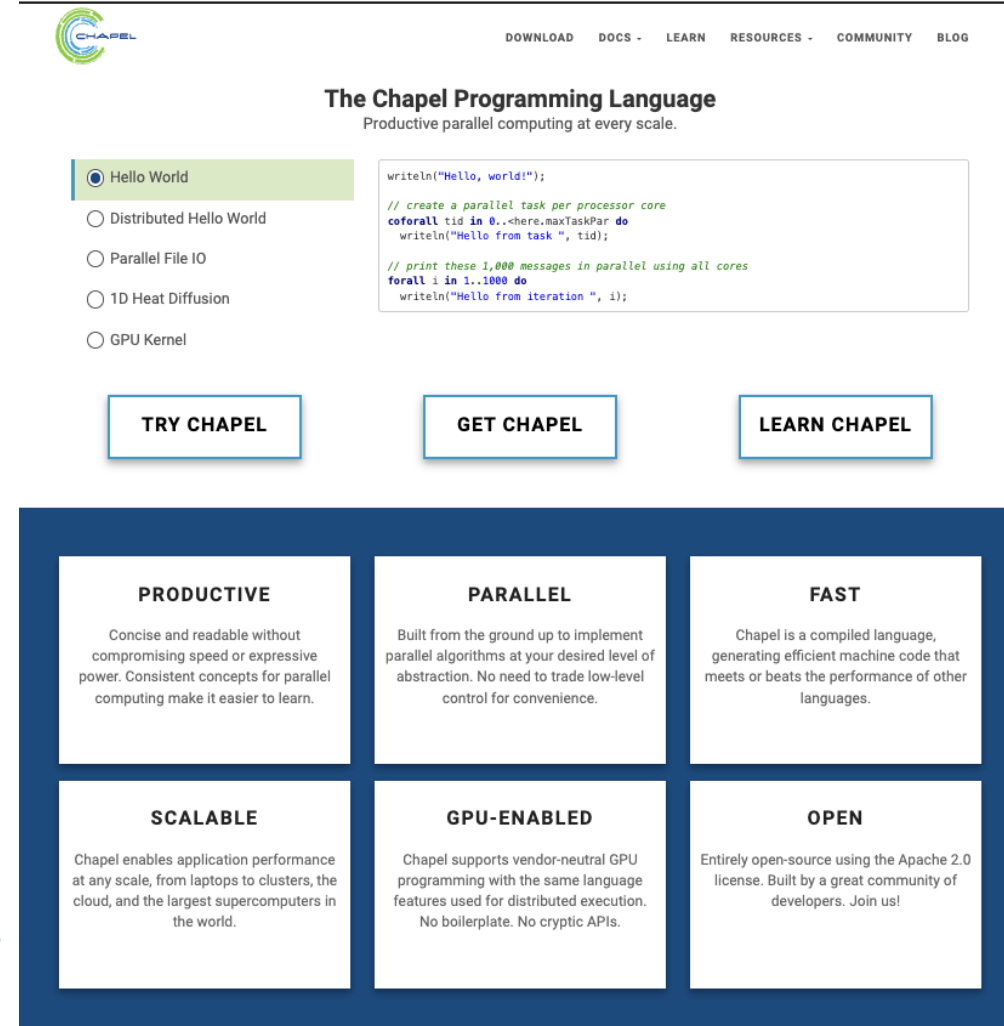
- (points to all other resources)

Social Media:

- Twitter: [@ChapelLanguage](https://twitter.com/ChapelLanguage)
- Facebook: [@ChapelLanguage](https://facebook.com/ChapelLanguage)
- YouTube: <http://www.youtube.com/c/ChapelParallelProgrammingLanguage>

Community Discussion / Support:

- Discord: <https://discord.com/invite/xu2xg45yqH>
- Stack Overflow: <https://stackoverflow.com/questions/tagged/chapel>
- GitHub Issues: <https://github.com/chapel-lang/chapel/issues>



The screenshot shows the Chapel Programming Language homepage. At the top, there's a navigation bar with links: DOWNLOAD, DOCS, LEARN, RESOURCES, COMMUNITY, and BLOG. The main heading is "The Chapel Programming Language" with the tagline "Productive parallel computing at every scale." Below this, there's a section for "Hello World" with a radio button selected, and four other options: Distributed Hello World, Parallel File IO, 1D Heat Diffusion, and GPU Kernel. To the right of these options is a code editor showing a "Hello World" program in Chapel. Below the code editor are three buttons: TRY CHAPEL, GET CHAPEL, and LEARN CHAPEL. At the bottom, there's a dark blue section with six white boxes, each containing a feature: PRODUCTIVE, PARALLEL, FAST, SCALABLE, GPU-ENABLED, and OPEN. Each box has a brief description of the feature.

PRODUCTIVE
Concise and readable without compromising speed or expressive power. Consistent concepts for parallel computing make it easier to learn.

PARALLEL
Built from the ground up to implement parallel algorithms at your desired level of abstraction. No need to trade low-level control for convenience.

FAST
Chapel is a compiled language, generating efficient machine code that meets or beats the performance of other languages.

SCALABLE
Chapel enables application performance at any scale, from laptops to clusters, the cloud, and the largest supercomputers in the world.

GPU-ENABLED
Chapel supports vendor-neutral GPU programming with the same language features used for distributed execution. No boilerplate. No cryptic APIs.

OPEN
Entirely open-source using the Apache 2.0 license. Built by a great community of developers. Join us!