



PARALLELISM IN CHAPEL, PART I

Chapel Team, edited by Michelle Strout

April 10, 2025

PLAN

- **Announcements**

- SA7 is due Friday April 11th
- Final projects are due Friday May 2nd (3 weeks left)

- **Last time**

- TopHat question about graduation
- Chapel programming basics in the context of an Nbody simulation, part II

- **Today**

- TopHat questions about Chapel basics
- Data parallelism in Chapel
- Domain decomposition in Chapel

OUTLINE: OVERVIEW OF PROGRAMMING IN CHAPEL

- Recall processing files in parallel
- Data parallelism concepts and examples including multi-locale parallelism with distributions
- Domains
- Forall Loops
- Domain Distributions
- Using a Different Domain Distribution



RECALL PROCESSING FILES IN PARALLEL

RECALL: ANALYZING MULTIPLE FILES USING PARALLELISM

 parfilekmer.chpl

parfilekmer.chpl

```
use FileSystem;
config const dir = "DataDir";
var fList = findFiles(dir);
var filenames =
    blockDist.createArray(0..<fList.size, string);
filenames = fList;

// per file word count
forall f in filenames {
    ...
    // code from kmer.chpl
    ...
}
```

```
prompt> chpl --fast parfilekmer.chpl
prompt> ./parfilekmer -nl 1
prompt> ./parfilekmer -nl 4
```

- shared and distributed-memory parallelism using 'forall'
 - in other words, parallelism within the locale/node and across locales/nodes
- a distributed array
- command line options to indicate number of locales



RECALL: BLOCK DISTRIBUTION OF ARRAY OF STRINGS

Locale 0

Locale 1

"filename1"	"filename2"	"filename3"	"filename4"	"filename5"	"filename6"	"filename7"	"filename8"
-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------

```
prompt> chpl --fast parfilekmer.chpl
prompt> ./parfilekmer -nl 2
```

- Array of strings for filenames is distributed across locales
- 'forall' will do parallelism across locales and then within each locale to take advantage of multicore




DATA PARALLELISM CONCEPTS AND EXAMPLES INCLUDING MULTI- LOCALE PARALLELISM WITH DISTRIBUTIONS

STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures:




STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (shared memory / multicore):




STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory):




STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory multicore):




STREAM TRIAD: CHAPEL (SEE TOPHAT QUESTION)

```
use BlockDist;  
  
config const m = 1000,  
           alpha = 3.0;  
  
const ProblemSpace = blockDist.createDomain({1..m});  
  
var A, B, C: [ProblemSpace] real;  
  
B = 2.0;  
C = 1.0;  
  
A = B + alpha * C;
```

The special sauce:

How should this index set—
and any arrays and
computations over it—be
mapped to the system?



Philosophy: Good, *top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Question: What happens when you remove the "with (ref A)"?

```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```


```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

DOMAINS

DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Domains (Index Sets)



```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel -nl 1 --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```


DOMAINS

Domain:

- A first-class index set
- The fundamental Chapel concept for data parallelism

```
config const m = 4, n = 8;  
  
const D = {1..m, 1..n};  
const Inner = {2..m-1, 2..n-1};
```

Question: What do 'D' and 'Inner' look like when you print them out?




DOMAINS

Domain:

- A first-class index set
- The fundamental Chapel concept for data parallelism
- Useful for declaring arrays and computing with them

```
config const m = 4, n = 8;  
  
const D = {1..m, 1..n};  
const Inner = {2..m-1, 2..n-1};  
  
var A, B, C: [D] real;
```



DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Arrays



```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

FORALL LOOPS

DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Data-Parallel Forall Loops

```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

FORALL LOOPS

Forall loops: Central concept for data parallel computation

- Like for-loops, but parallel
- Implementation details determined by iterand (e.g., D below)
 - specifies number of tasks, which tasks run which iterations, ...
 - in practice, typically uses a number of tasks appropriate for target HW

```
forall (i,j) in D with (ref A) do  
    A[i,j] = i + j/10.0;
```

Forall loops assert...

- ...**parallel safety**: OK to execute iterations simultaneously
- ...**order independence**: iterations could occur in any order
- ...**serializability**: all iterations could be executed by one task
 - e.g., can't have synchronization dependences between iterations

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

COMPARISON OF LOOPS: FOR, FORALL, AND COFORALL (TOPHAT QUESTION)

For loops: executed using one task

- use when a loop must be executed serially
- or when one task is sufficient for performance

Forall loops: typically executed using $1 < \#tasks \ll \#iters$

- use when a loop *should* be executed in parallel...
- ...but *can* legally be executed serially
- use when desired $\# tasks \ll \# of iterations$

Coforall loops: executed using a task per iteration

- use when the loop iterations *must* be executed in parallel
- use when you want $\# tasks == \# of iterations$
- use when each iteration has substantial work



DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

This is a shared memory program

Nothing has referred to remote
locales, explicitly or implicitly

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel -nl 1 --n=5
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

DOMAIN DISTRIBUTIONS

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Domain Distribution
(Map Data Parallelism to the System)

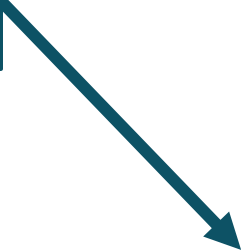
```
use CyclicDist;  
config const n = 1000;  
var D = cyclicDist.createDomain({1..n, 1..n});  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel --n=5 -nl 4  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

High-level distributed and shared
memory parallelism



```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});

var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

Provides programmability and control


- Lowering of code is well-defined
- User can control details
- Part of Chapel's *multiresolution philosophy*...

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel --n=5 --nl 4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

CHAPEL'S MULTIREOLUTION PHILOSOPHY

Multiresolution Design: Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for greater degrees of control
- build the higher-level concepts in terms of the lower
- permit users to intermix layers arbitrarily



DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

⇒ invoke and inline D's
default parallel iterator

- defined by D's type /
domain distribution

default domain distribution

- create a task per local core
- block indices across tasks

```
config const n = 1000;
var D = {1..n, 1..n};

var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel --n=5 -nl 1
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

⇒ invoke and inline D's
default parallel iterator

- defined by D's type /
domain distribution

cyclic domain distribution

on each target locale...

- create a task per core
- Round robin local indices
across tasks

```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});

var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel --n=5 -nl=4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

What if I don't like D's
iteration strategy?

```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});
var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

Write and call your own parallel iterator:

```
forall (i,j) in myParIter(D) do...
```



DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

What if I don't like D's
iteration strategy?

```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});
var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

Write and call your own parallel iterator:

```
forall (i,j) in myParIter(D) do...
```


Or use a different domain distribution:

```
var D = blockDist.createDomain({1..n, 1..n});
```

USING A DIFFERENT DOMAIN DISTRIBUTION


DOMAIN DISTRIBUTIONS: A MULTIREOLUTION FEATURE

Domain distributions are “recipes” that instruct the compiler how to map the global view of a computation...




$$A = B + \alpha * C;$$

...to the target locales' memory and processors:

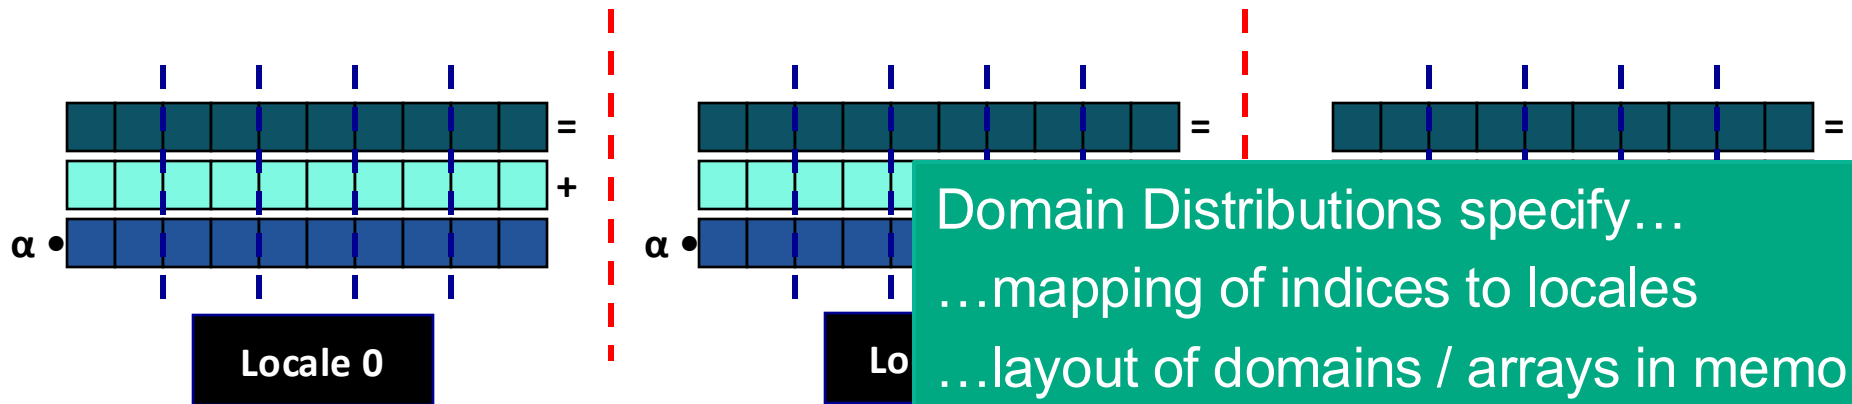


DOMAIN DISTRIBUTIONS: A MULTIREOLUTION FEATURE

Domain distributions are “recipes” that instruct the compiler how to map the global view of a computation...



...to the target locales' memory and processors:



Domain Distributions specify...

...mapping of indices to locales


...layout of domains / arrays in memory

...parallel iteration strategies

...core operations on arrays / domains

SAMPLE DOMAIN DISTRIBUTIONS: BLOCK AND CYCLIC (TOPHAT QUESTION)


```
var Dom = blockDist.createDomain({1..4, 1..8});
```



distributed to



```
var Dom = cyclicDist.createDomain({1..4, 1..8});
```



distributed to

