



PARALLELISM IN CHAPEL, PART I

Chapel Team, edited by Michelle Strout

April 15, 2025

PLAN

- **Announcements**

- SA7 grades are posted
- LA3 is due on Friday April 25th (1.5 weeks left)
- Final projects are due Friday May 2nd (2.5 weeks left)

- **Last time**

- TopHat questions about Chapel basics
- Data parallelism in Chapel
- Domain decomposition in Chapel

- **Today**

- TopHat questions about Chapel data parallelism
- LA3 parallelism suggestions
- Domain decomposition in Chapel

OUTLINE: OVERVIEW OF PARALLELISM IN CHAPEL, PART II

- LA3: Parallelizing a Chapel program
- Extra, Extra Credit: Leaderboard for performance improvement and/or UA HPC system
- Implicit Communication: Remote writes/Puts and Reads/Gets
- Parallelizing a 1D heat diffusion solver (Hands On)
- Heat 2D example with CommDiagnostics (Hands On)



LA3: PARALLELIZING A COMPUTATION IN CHAPEL

LA3: PARALLELIZING CHAPEL

Basics for LA3

- la3_parallel.chpl is a copy of la3_serial.chpl
- Improve the performance of la3_parallel.chpl so it is at least 10% faster than la3_serial.chpl
- Describe what you did in the comment header

Ideas

- Parallelize one or more loops
- Fuse the loops and remove temp file write and read
- Distributed data parallelism, probably need to run on the UofA HPC system
- Deal with load imbalance issues
- Limit number of tasks to when scaling improves

la3_serial.chpl

```
// convert all files to gray scale
for fname in files {
    ...
    var imageArray = readImage(fName, imageType.png);
    var grayImage = rgbToGrayscale(imageArray);
    writeImage(grayImage, ..., grayImage);
}

// do edge detection on all of the grayscale files
for fname in files {
    ...
    var grayArray = readImage(fName, imageType.png);
    var sobelImage = sobelEdgeDetection(grayArray);
    writeImage(edgefName, ..., sobelImage);
}
```

Using the UofA HPC system for Extra, Extra Credit

- Log into HPC system and run a Chapel program

```
Laptop_prompt> ssh netid@hpc.arizona.edu
[netid@gatekeeper ~]$ shell
(puma) [netid@wentletrap ~]$ ocelote
(ocelote) [netid@wentletrap ~]$ interactive -a cs372spring2025 -n 8 -N 4 -t 1:0:0
Run "interactive -h for help customizing interactive use"
Submitting with /usr/local/bin/salloc --job-name=interactive --mem-per-cpu=4GB --
nodes=4      --ntasks=8 --time=1:0:0 --account=cs372spring2025 --partition=standard
// set up an ssh key for use with GitHub
// clone your LA3 github repository
// cd into a directory you have with Chapel code
(ocelote) [netid@i7n# Chapel]$ module load chapel-ibv
(ocelote) [netid@i7n# Chapel]$ chpl --version
chpl version 2.4.0
...
(ocelote) [netid@i7n# Chapel]$ chpl hello6-taskpar.chpl
(ocelote) [netid@i7n# Chapel]$ export GASNET_PHYSMEM_MAX=6GB
(ocelote) [netid@i7n# Chapel]$ ./hello6-taskpar -nl 2
```

- References

- https://hpcdocs.hpc.arizona.edu/registration_and_access/system_access/#command-line-access

IMPLICIT COMMUNICATION:
REMOTE WRITES/PUTS AND READS/GETS

CHAPEL SUPPORTS A GLOBAL NAMESPACE WITH PUTS AND GETS

Note 1: Variables are allocated on the locale where the task is running

 03-onClause.chpl

03-

onClause.chpl

```
config const verbose = false;
```

```
var total = 0,  
    done = false;
```

```
...
```

```
on Locales[1] {  
    var x, y, z: int;
```

```
...
```

```
}
```

verbose false

total 0

done false

locale 0

x 0

y 0

z 0

locale 1

CHAPEL SUPPORTS A GLOBAL NAMESPACE WITH PUTS AND GETS

Note 2: Tasks can refer to lexically visible variables, whether local or remote

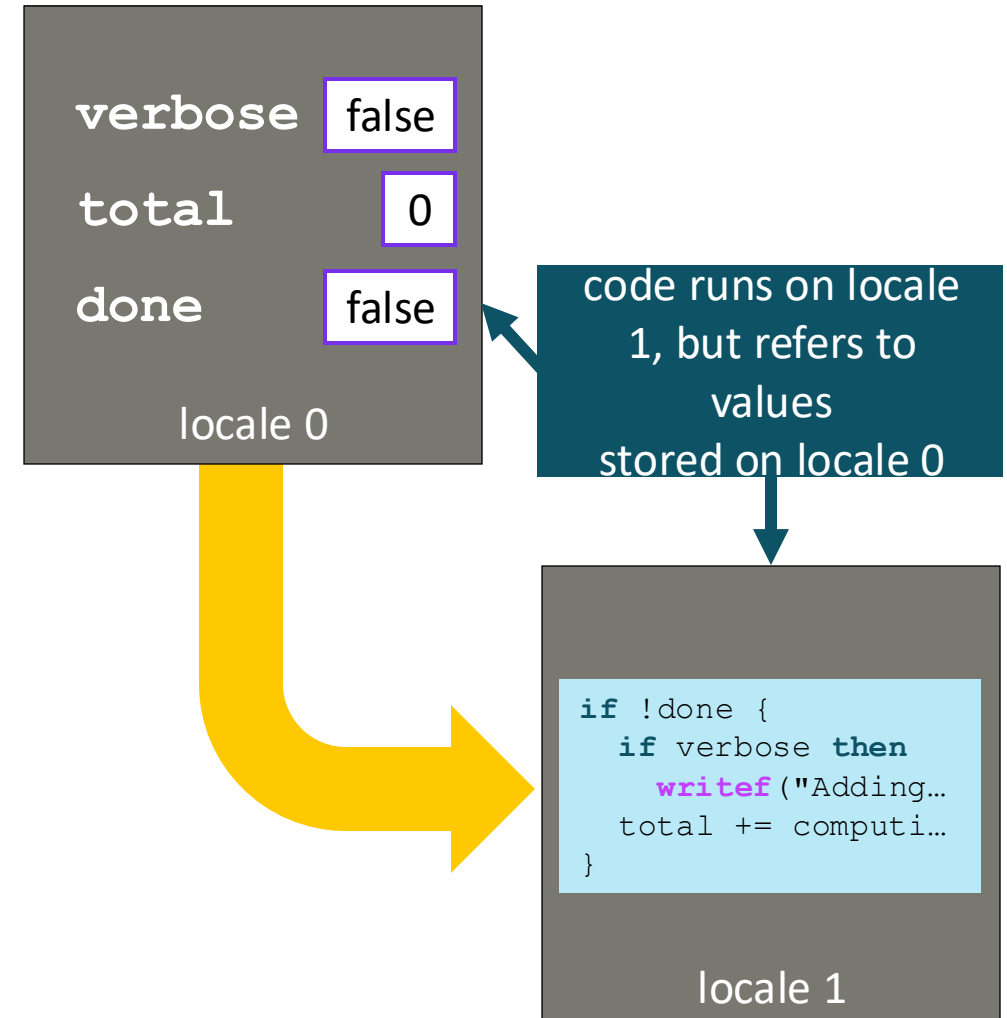
03-onClause.chpl

03-onClause.chpl

```
config const verbose = false;
var total = 0,
    done = false;

...

on Locales[1] {
  if !done {
    if verbose then
      writef("Adding locale 1's contribution");
    total += computeMyContribution();
  }
}
```



ARRAY-BASED PARALLELISM AND LOCALITY

03-basics-distarr.chpl

```
writeln("Hello from locale ", here.id);

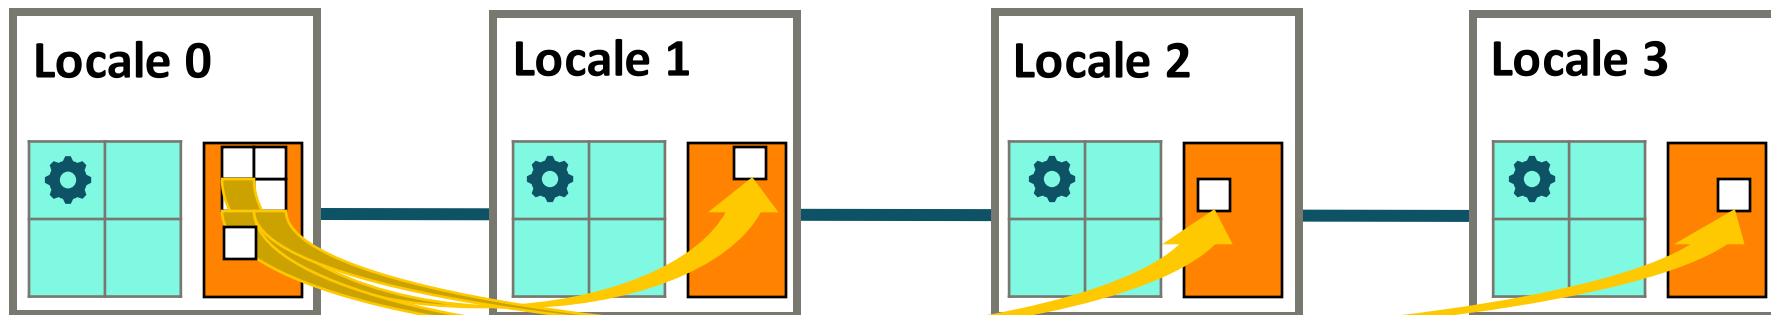
var A: [1..2, 1..2] real;

use BlockDist;

var D = blockDist.createDomain({1..2, 1..2});
var B: [D] real;
B = A;
```

Chapel also supports distributed domains (index sets) and arrays

They also result in parallel distributed computation

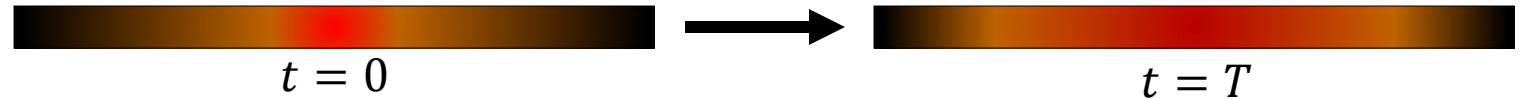


PARALLELIZING A 1D HEAT DIFFUSION SOLVER (HANDS ON)

Also read <https://github.com/jeremiah-corrado/Chapel-Heat1D-PPA>

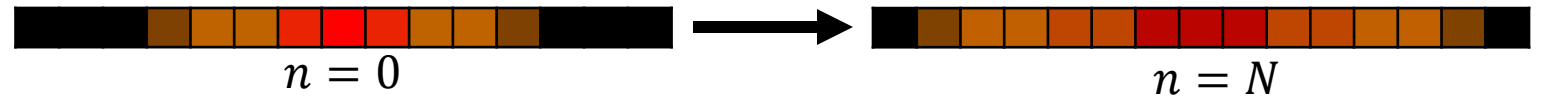
1D HEAT EQUATION EXAMPLE

Differential equation: $\frac{\partial u}{\partial t} = \alpha \frac{\partial^2 u}{\partial x^2}$



Discretized (finite difference) equation: $u_i^{n+1} = u_i^n + \alpha (u_{i-1}^n - 2u_i^n + u_{i+1}^n)$

- where $i \in \Omega \subset \mathbb{R}^1$ are discrete points in space, and $(n, n+1, \dots)$ are discrete instances in time



Finite difference algorithm:

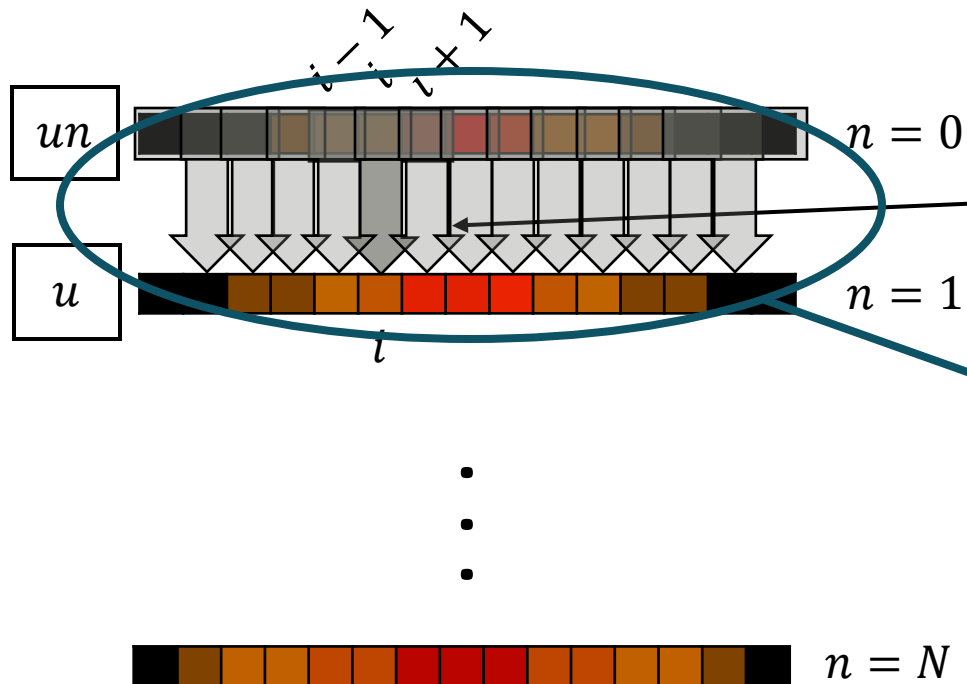
- define Ω to be a set of discrete points along the x-axis
- define $\hat{\Omega}$ over the same points, excluding the boundaries
- define an array u over Ω
- set some initial conditions
- create a temporary copy of u , named un
- for N timesteps:
 - (1) swap u and un
 - (2) compute u in terms of un over $\hat{\Omega}$

```
1  const omega = {0..<nx},
2      omegaHat = omega.expand(-1);
3  var u: [omega] real = 1.0;
4  u[nx/4..3*nx/4] = 2.0;
5  var un = u;
6  for 1..N {
7      un <=> u;
8      forall i in omegaHat do
9          u[i] = un[i] + alpha *
10              (un[i-1] - 2*un[i] + un[i+1]);
11  }
```

1D HEAT EQUATION EXAMPLE

This pattern is often referred to as a *Stencil Computation*

- The values in the array can be computed by applying a "stencil" to its previous state
- Note that in this case, the stencil can be applied to the entire array in parallel
 - each value in u^n depends strictly on values in u



"stencil"

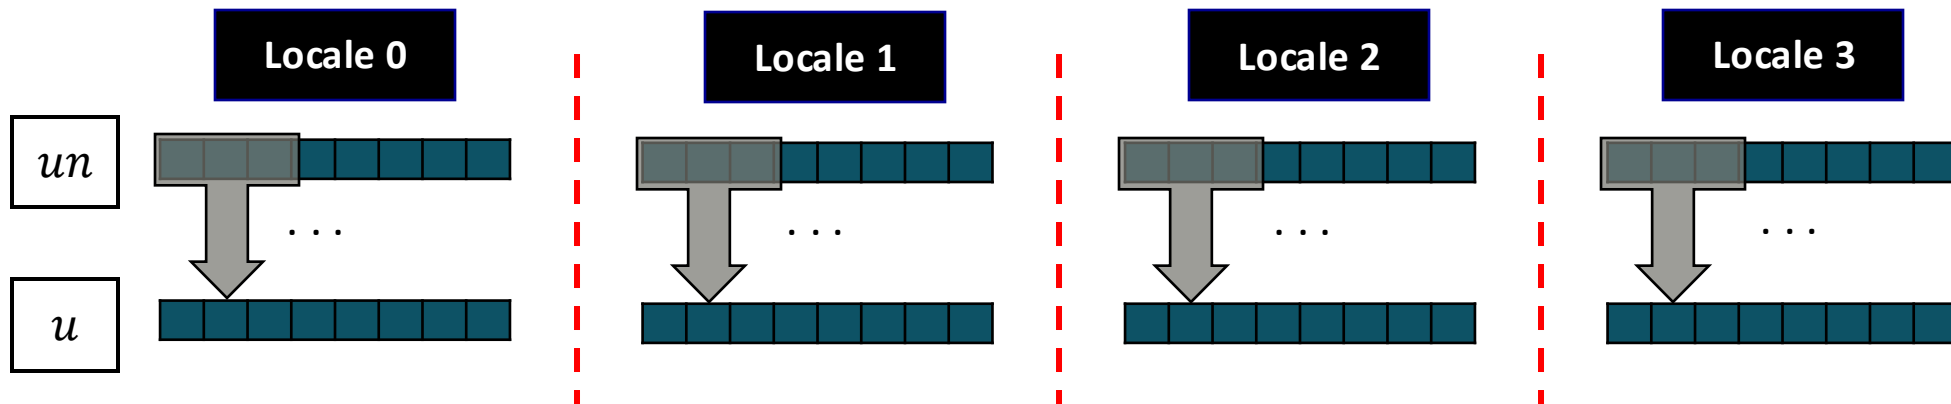
$$u_i^{n+1} = u_i^n + \alpha (u_{i-1}^n - 2u_i^n + u_{i+1}^n)$$

```
7  ...  
8  forall i in omegaHat do  
9      u[i] = un[i] + alpha *  
10         (un[i-1] - 2*un[i] + un[i+1]);  
11  ...
```

HANDS ON: DISTRIBUTING THE 1D HEAT EQUATION

Imagine we want to simulate a very large domain

- We could use the Block distribution to distribute u and un across multiple locales
 - taking advantage of their memory and compute resources



Look at **heat-1D-block.chpl** and fill in the blanks to make the arrays block-distributed

Hint | Define a block-distributed domain:

```
use BlockDist;  
...  
const myBlockDom = blockDist.createDomain({1..10});
```

HANDS ON: DISTRIBUTING THE 1D HEAT EQUATION

heat-1D-block-solution.chpl

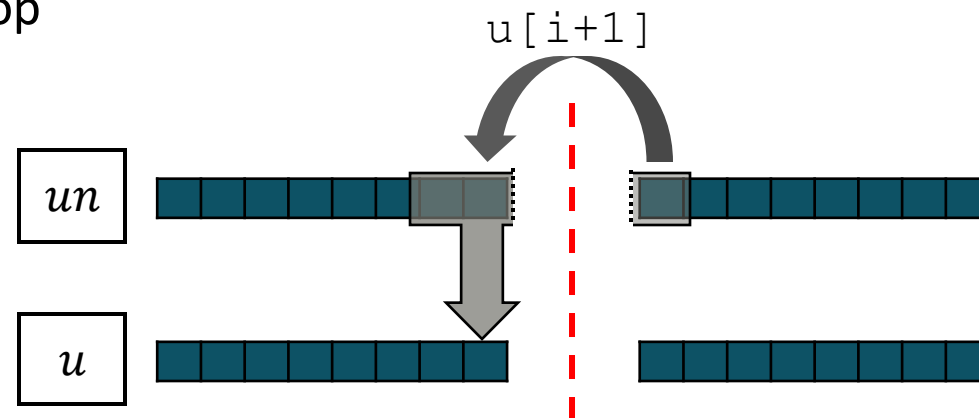
Solution: make 'omega' block-distributed:

```
omega = blockDist.createDomain({0.. $\text{nx}$ });
```

Why does this work?

- 'omegaHat' inherits 'omega's distribution
- 'u' is block-distributed
- 'un' inherits 'u's domain (and distribution)
- 'omegaHat' invokes 'blockDist's parallel/distr. iterator
 - the body of the loop is automatically split across multiple tasks on each locale
- Communication occurs automatically when a loop references a value stored on a remote locale

```
1  const omega =  
2      blockDist.createDomain({0.. $\text{nx}$ }),  
3      omegaHat = omega.expand(-1);  
4  var u: [omega] real = 1.0;  
5  u[nx/4.. $3*\text{nx}/4$ ] = 2.0;  
6  var un = u;  
7  for 1.. $N$  {  
8      un <=> u;  
9      forall i in omegaHat do  
10         u[i] = un[i] + alpha *  
11             (un[i-1] - 2*un[i] + un[i+1]);  
12 }
```



HEAT 2D EXAMPLE WITH COMMDIAGNOSTICS (HANDS ON)

2D HEAT EQUATION EXAMPLE

2D and 3D stencil codes are more common and practical

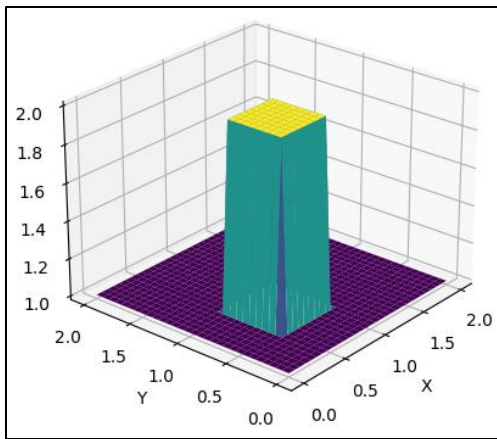
- They also present more interesting considerations for parallelization and distribution

2D heat / diffusion PDE:

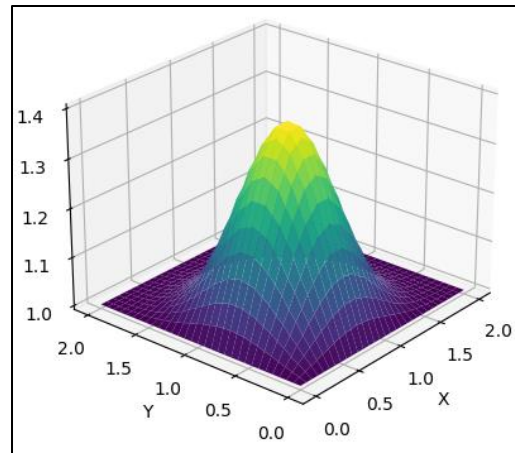
$$\frac{\partial u}{\partial t} = \alpha \Delta u = \alpha \left(\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} \right)$$

Discretized (finite-difference) form:

$$u_{i,j}^{n+1} = u_{i,j}^n + \alpha (u_{i+1,j}^n + u_{i-1,j}^n - 4u_{i,j}^n + u_{i,j+1}^n + u_{i,j-1}^n)$$



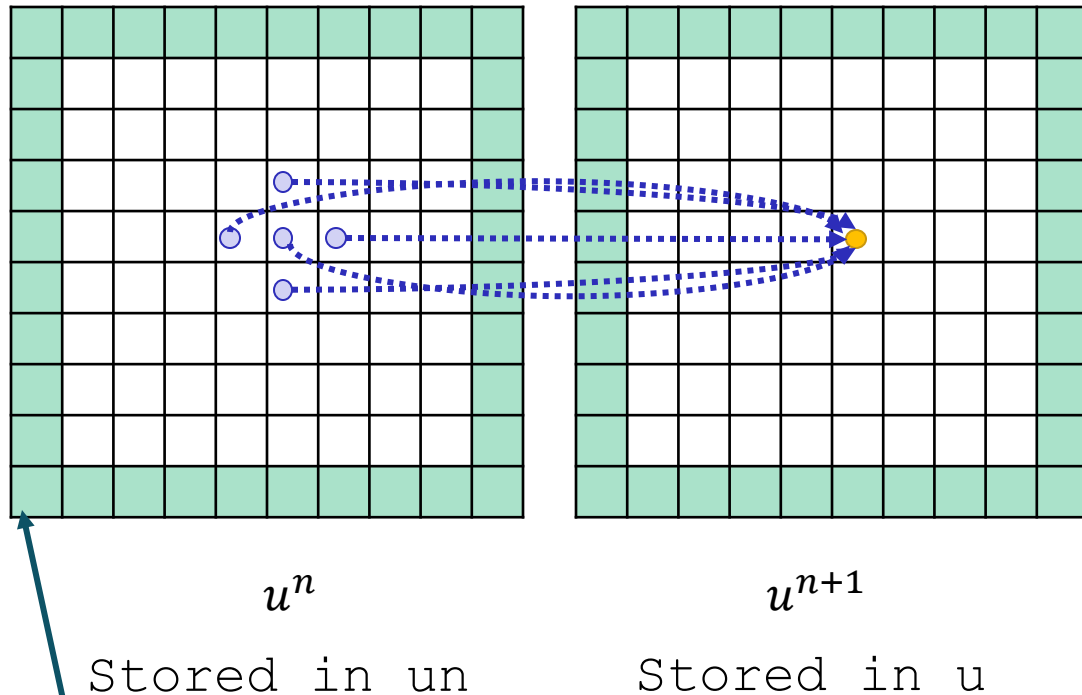
$n = 0$



$n = N$

```
1  const omega = {0..
```

PARALLEL 2D HEAT EQUATION



Fixed
boundary
values

- This computation uses a "5 point stencil"
- Each point in 'u' can be computed in parallel
 - this is accomplished using a 'forall' loop

```
7  ...
8  forall (i, j) in omegaHat do
9      u[i, j] = un[i, j] + alpha * (
10         un[i-1, j] + un[i, j-1] +
11         un[i+1, j] + un[i, j+1] -
12         4 * un[i, j]);
13  ...
```

$$u_{i,j}^{n+1} = u_{i,j}^n + \alpha(u_{i-1,j}^n + u_{i,j-1}^n + u_{i+1,j}^n + u_{i,j+1}^n - 4u_{i,j}^n)$$

BLOCK DISTRIBUTED & PARALLEL 2D HEAT EQUATION

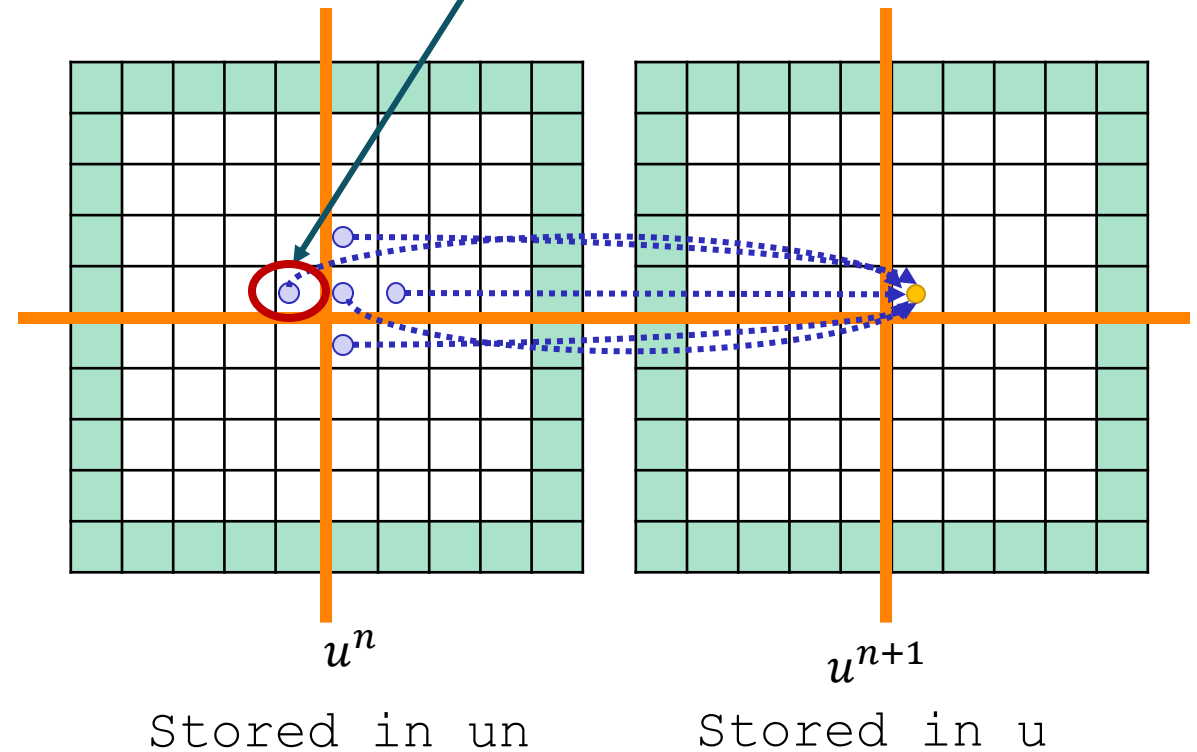
- Declaring distributed domains with the block distribution

```
const Omega = blockDist.createDomain(0.. $nx$ , 0.. $ny$ ),  
      OmegaHat = Omega.expand(-1);
```

- Distributed & Parallel loop over 'OmegaHat'

```
for 1.. $nt$  {  
  u <=> un;  
  
  forall (i, j) in OmegaHat do  
    u[i, j] = un[i, j] + alpha * (  
      un[i-1, j] + un[i, j-1] +  
      un[i+1, j] + un[i, j+1] -  
      4 * un[i, j]);  
}
```

Array access across locale boundaries automatically invokes communication



STENCIL DISTRIBUTED & PARALLEL 2D HEAT EQUATION

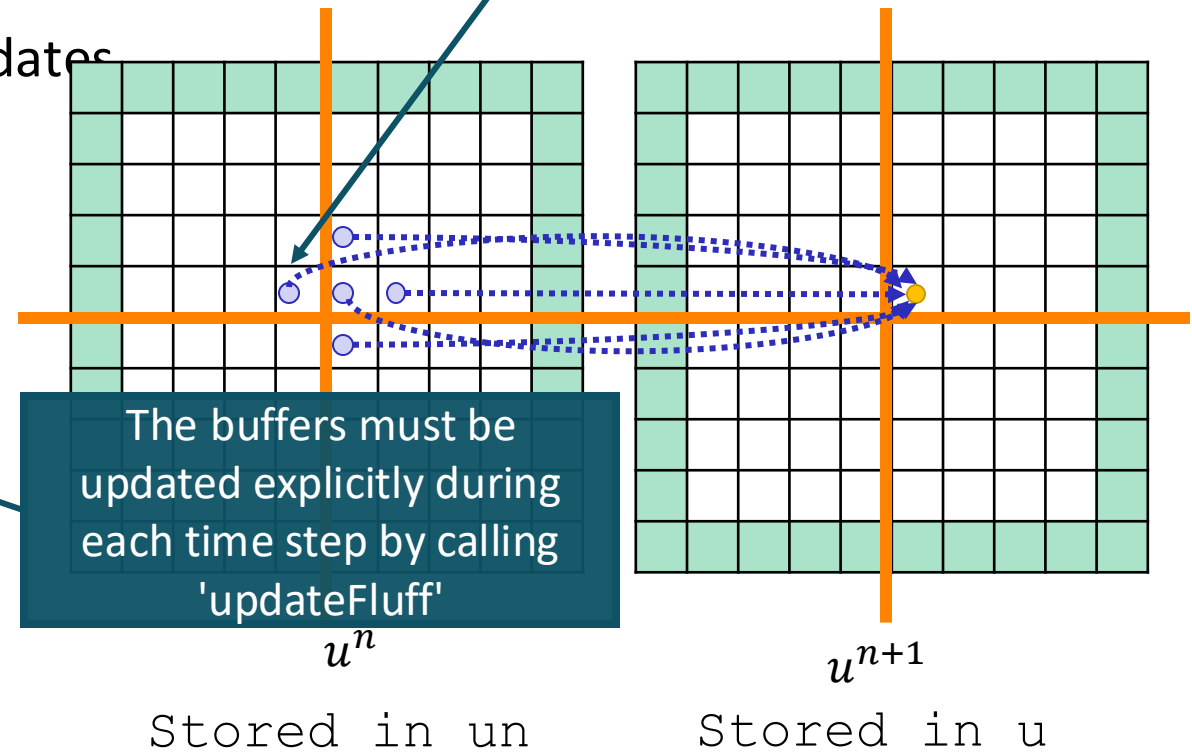
- Declaring distributed domains with the stencil distribution

```
const Omega = stencilDist.createDomain(  
    {0..  
nx, 0..  
ny}, fluff=(1,1)),  
OmegaHat = Omega.expand(-1);
```

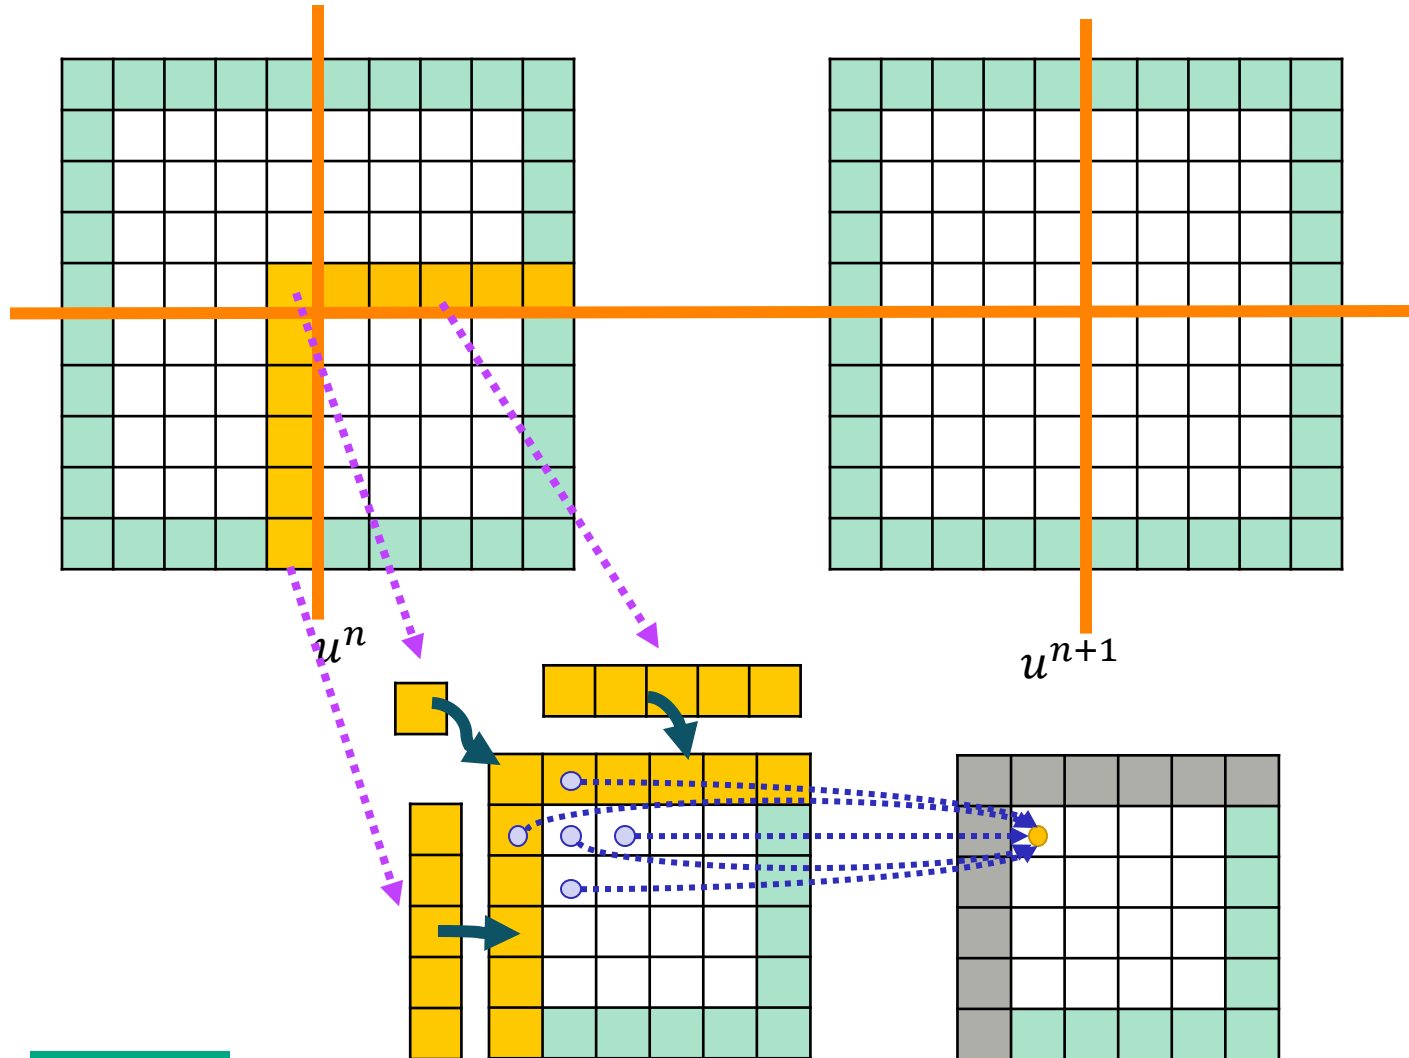
- Distributed & Parallel loop including buffer updates

```
for 1..  
nt {  
    u <=> un;  
    un.updateFluff();  
    forall (i, j) in OmegaHat do  
        u[i, j] = un[i, j] + alpha * (  
            un[i-1, j] + un[i, j-1] +  
            un[i+1, j] + un[i, j+1] -  
            4 * un[i, j]);  
}
```

Array access across locale boundaries (within the fluff region) results in a local buffer access — no communication is required



STENCIL DISTRIBUTED & PARALLEL 2D HEAT EQUATION



- Each locale owns a region of the array surrounded by a "fluff" (buffer) region
- Calling 'updateFluff' copies values from neighboring regions of the array into the local buffered region
- Subsequent accesses of those values result in a local memory access, rather than a remote communication

COMM DIAGNOSTICS

The 'CommDiagnostics' module provides functions for tracking comm between locales

- the following is a common pattern:

```
use CommDiagnostics;
...
startCommDiagnostics();
potentiallyCommHeavyOperation();
stopCommDiagnostics();
...
printCommDiagnosticsTable();
```


- which results in a table summarizing comm counts between the **start** and **stop** calls, e.g.,

locale	get	put	execute_on	execute_on_nb
-----:	--:	--:	-----:	-----:
0	10	0	6	12
1	105	5	0	0
2	105	4	0	0
3	105	7	0	0

- Compiling with '--no-cache-remote' before collecting comm diagnostics is recommended



HANDS ON: HEAT 2D COMM DIAGNOSTICS RESULTS

 heat-2D-block.chpl, heat-2D-stencil.chpl

- Comparing comm diagnostics for:

- heat-2D-block.chpl
- heat-2D-stencil.chpl

- *Compilation:*

```
chpl heat-2D-block.chpl --fast
    --no-cache-remote -sRunCommDiag=true

chpl heat-2D-stencil.chpl -fast
    --no-cache-remote -sRunCommDiag=true
```

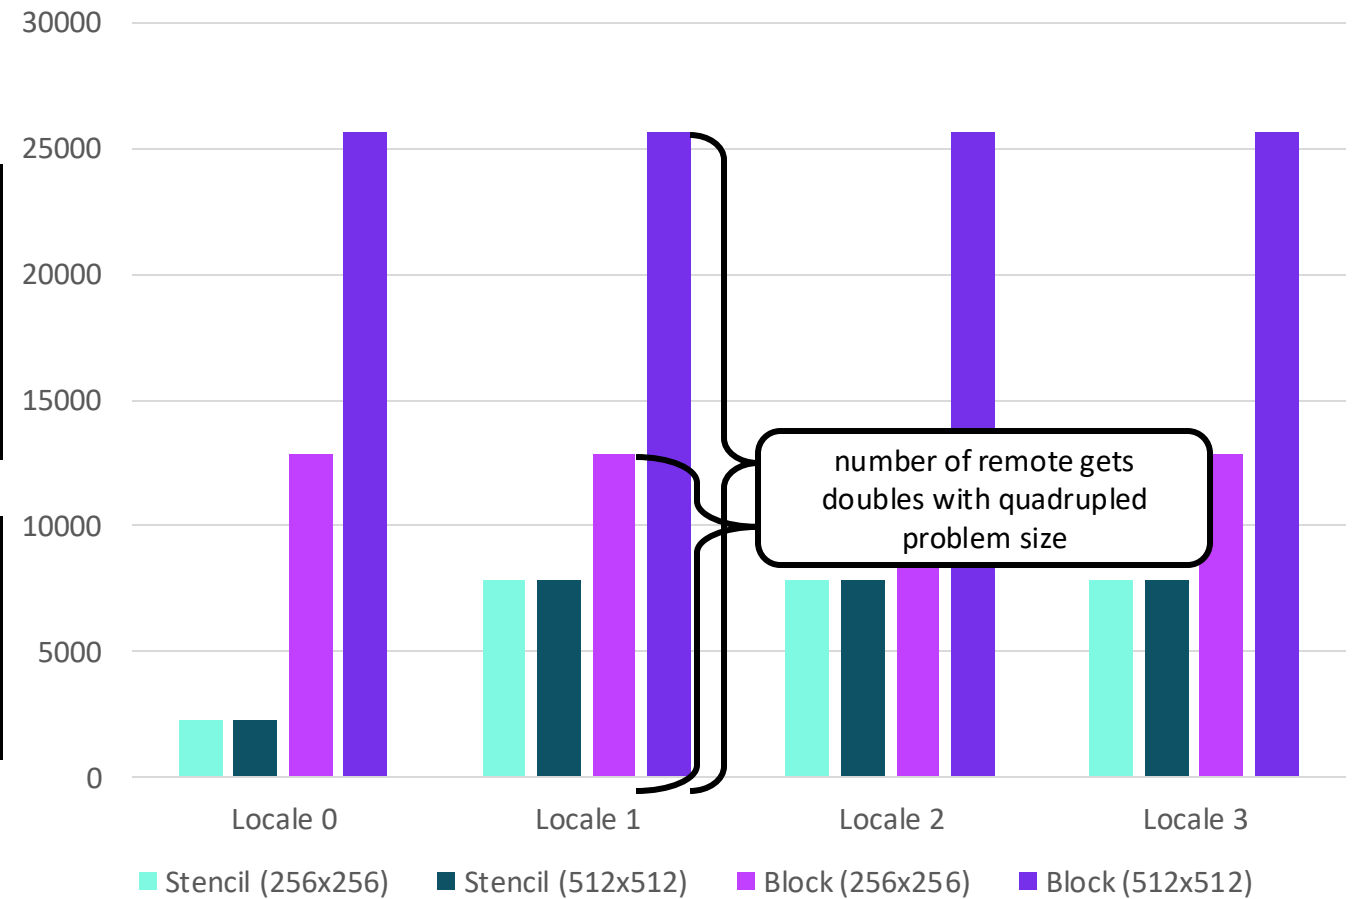
- *Execution:*

```
./heat-2D-block -nl4 --nx=256 --ny=256

./heat-2D-stencil -nl4 --nx=512 --ny=512
```

- **Block:** number of gets scales with size
- **Stencil:** static number of gets per iteration

Number of Gets on 4 Locales – Block vs. Stencil



SUMMARIZING WHAT WE LEARNED ABOUT PARALLELISM IN CHAPEL

- Data parallelism session
 - Provides shared memory and distributed memory parallelism
 - Distributions like block and cyclic can be applied to arrays of any dimension
 - Main control abstraction is the 'forall' loop
 - 'forall' loop uses default iterator over provided array or domain, but can use own iterator
 - This is an example of multi-resolution design in Chapel, i.e., the 'forall' loop is mapped down to lower-level abstractions like 'coforall'
 - CommDiagnostics module can be used to observe the number of remote puts/writes and gets/reads at runtime



CHAPEL RESOURCES

Chapel homepage: <https://chapel-lang.org>

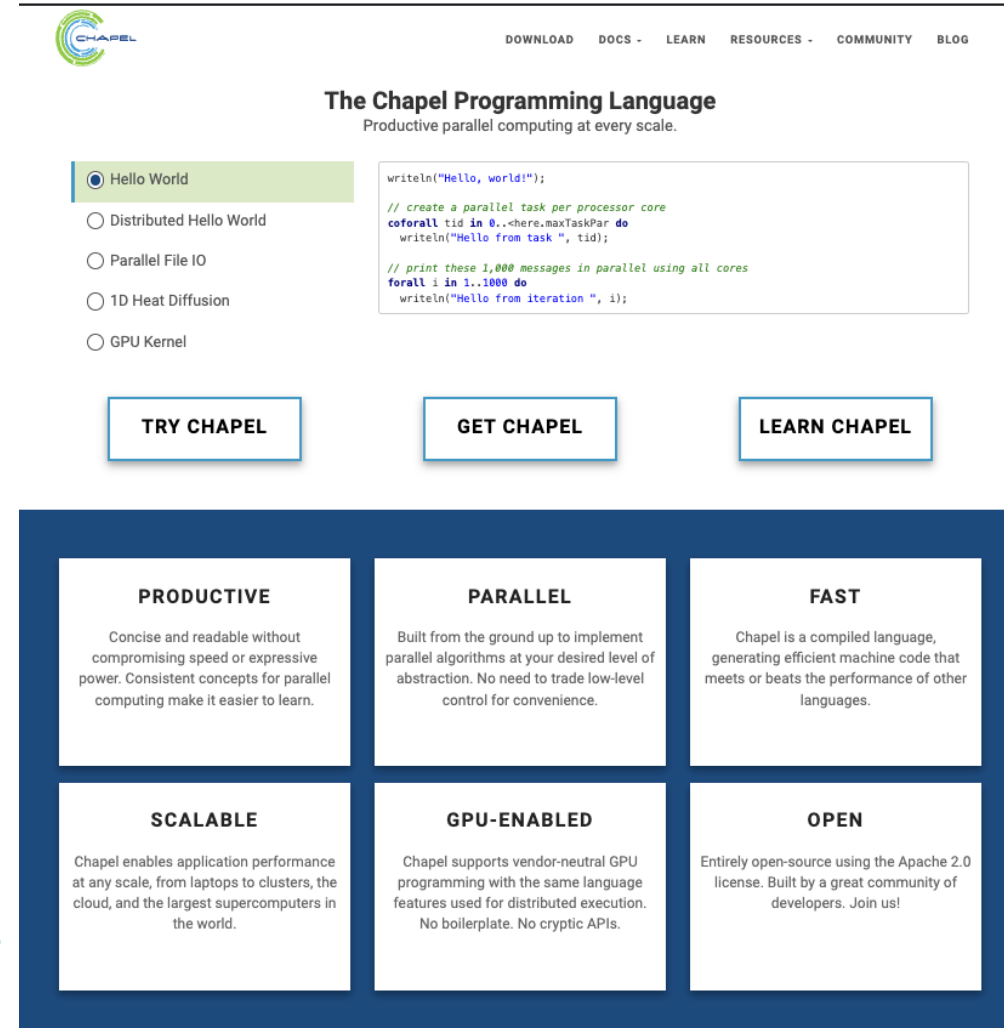
- (points to all other resources)

Social Media:

- Twitter: [@ChapelLanguage](https://twitter.com/ChapelLanguage)
- Facebook: [@ChapelLanguage](https://facebook.com/ChapelLanguage)
- YouTube: <http://www.youtube.com/c/ChapelParallelProgrammingLanguage>

Community Discussion / Support:

- Discord: <https://discord.com/invite/xu2xg45yqH>
- Stack Overflow: <https://stackoverflow.com/questions/tagged/chapel>
- GitHub Issues: <https://github.com/chapel-lang/chapel/issues>



The screenshot shows the Chapel Programming Language homepage. At the top, there's a navigation bar with links: DOWNLOAD, DOCS, LEARN, RESOURCES, COMMUNITY, and BLOG. The main heading is "The Chapel Programming Language" with the tagline "Productive parallel computing at every scale." Below this, there's a section for "Hello World" with a radio button selected, and four other options: Distributed Hello World, Parallel File IO, 1D Heat Diffusion, and GPU Kernel. To the right of these options is a code editor showing a "Hello, world!" program with comments. Below the code editor are three buttons: TRY CHAPEL, GET CHAPEL, and LEARN CHAPEL. At the bottom, there's a dark blue section with six white boxes, each containing a feature: PRODUCTIVE, PARALLEL, FAST, SCALABLE, GPU-ENABLED, and OPEN, each with a brief description.

PRODUCTIVE
Concise and readable without compromising speed or expressive power. Consistent concepts for parallel computing make it easier to learn.

PARALLEL
Built from the ground up to implement parallel algorithms at your desired level of abstraction. No need to trade low-level control for convenience.

FAST
Chapel is a compiled language, generating efficient machine code that meets or beats the performance of other languages.

SCALABLE
Chapel enables application performance at any scale, from laptops to clusters, the cloud, and the largest supercomputers in the world.

GPU-ENABLED
Chapel supports vendor-neutral GPU programming with the same language features used for distributed execution. No boilerplate. No cryptic APIs.

OPEN
Entirely open-source using the Apache 2.0 license. Built by a great community of developers. Join us!