



MEMORY SAFETY: CHAPEL, RUST, C/C++, PYTHON, MPI, OPENSHPMEM

Slides based off blog post by Michael Ferguson

<https://chapel-lang.org/blog/posts/memory-safety/>

Michelle Strout put together slides for CSc 372 at UArizona

April 17, 2025

PLAN

- **Announcements**

- Brad Chamberlain, Chapel co-creator, visiting Thursday April 24th
- LA3 is due on Friday April 25th (1 weeks left)
- Final projects are due Friday May 2nd (2 weeks left)

- **Last time**

- TopHat questions about Chapel data parallelism
- LA3 parallelism suggestions
- Heat diffusion in Chapel and implicit communication

- **Today**

- Recall extra credit opportunities
- TopHat questions about Chapel implicit communication
- Memory safety comparison between languages



OUTLINE: MEMORY SAFETY

- Visualizing
 - Use-after-free and memory leak bugs
 - How ownership types prevent them
 - How garbage collection prevents them
- Punchline from blog: comparison of languages
- Kinds of errors memory safety help prevent
 - Variable not initialized
 - Mishandling strings
 - Use-after-free
 - Out-of-bounds array access

Memory Safety in Chapel

Posted on April 10, 2025.

Tags:

Safety

Language Comparison

By: [Michael Ferguson](#)

VISUALIZING SOME MEMORY BUGS

IN CLASS ACTIVITY: VISUALIZING SOME MEMORY BUGS

- Learning Objectives

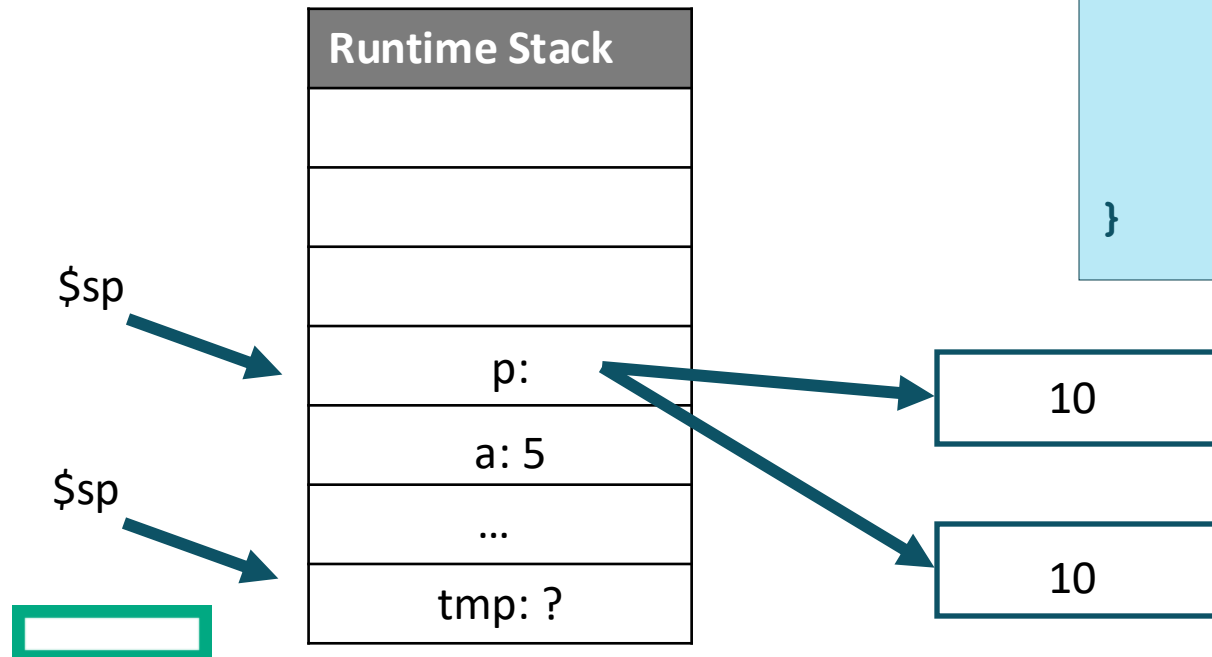
- Understand how memory is managed using stack frames and dynamic allocation
- Identify common memory safety issues: use-after-free and memory leaks
- See how different languages (Rust, Chapel, Java) handle memory safety
- Role-play borrow checking to understand ownership models



RUNTIME STACK AND HEAP DIAGRAMS

Memory allocation

- Local variables are allocated on the run-time stack
- Malloc allocates memory in the heap
- Malloc returns a pointer to that memory in the heap that is typically then stored in the runtime stack

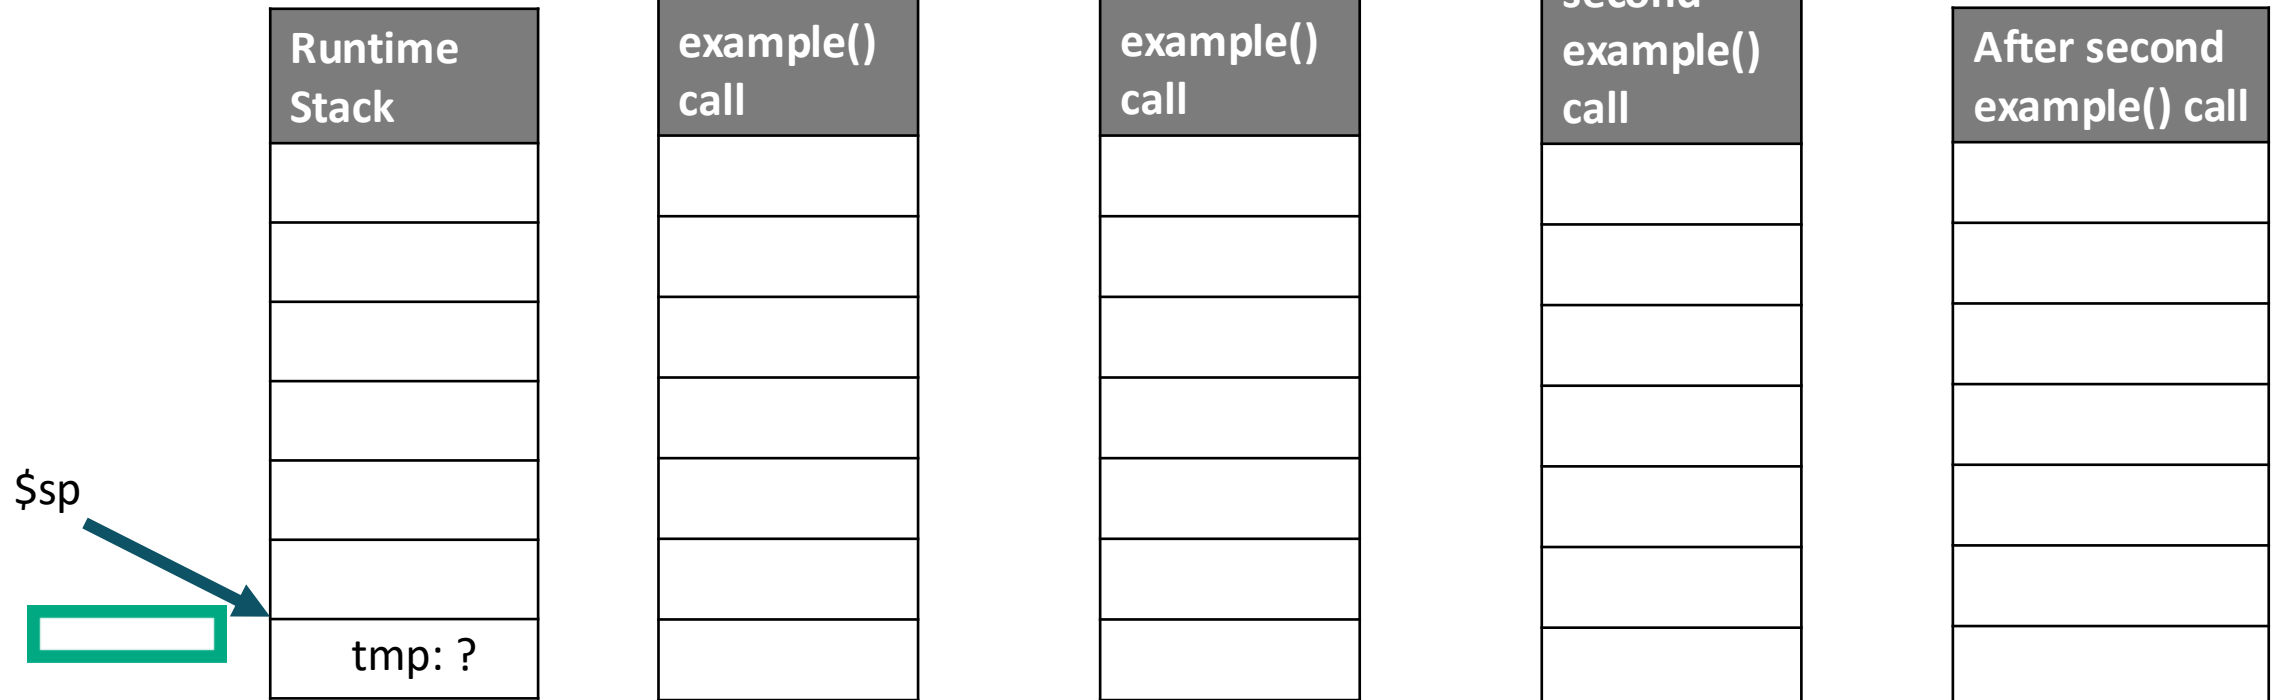


```
// example C code
void example() {
    int a = 5;           // local variable
    int* p = malloc(4); // heap allocation
    *p = 10;
} // p goes out of scope, but memory is still allocated unless freed

void main() {
    int tmp;           // local variable
    example();
    example();
}
```

ACTIVITY: SHOW THE STEPS DURING AND AFTER EACH EXAMPLE() CALL

```
void example() {  
    int a = 5;  
    int* p = malloc(4);  
    *p = 10;  
}  
void main() {  
    int tmp;  
    example();  
    example();  
}
```



WHAT HAPPENS WITH THE FOLLOWING TWO EXAMPLES?

```
int* x = malloc(sizeof(int));  
*x = 42;  
free(x);  
printf("%d\n", *x); // use-after-free
```

```
void leak() {  
    int* y = malloc(sizeof(int));  
    *y = 99;  
    // no free(y);  
} // memory leak
```



OWNERSHIP TYPES IN RUST AND CHAPEL

Ownership types

- Only one reference var owns the reference
- Can move ownership
- The compiler generates an error if trying to access through a reference that no longer owns something
- When that reference var goes out of scope, the heap allocation is freed

```
// example Rust code
struct MyBox {
    value: i32,
}
fn main() {
    let a = MyBox { value: 10 };
    let b = a; //ownership moves, a no longer valid
    println!("{}", a.value); //compile error
}

// example Chapel code
class MyClass {
    var x: int;
}
proc main() {
    var c = new owned MyClass();
    var d = c; // transfers ownership
    writeln(c); // error: c no longer owns
} // memory automatically freed
```

GARBAGE COLLECTION IN LANGUAGES LIKE JAVA AND PYTHON

Garbage collection

- Mark and sweep variants will traverse the active runtime stack and mark everything in the heap that is reachable
- During the sweep, anything in the heap that is not marked will be swept up and freed

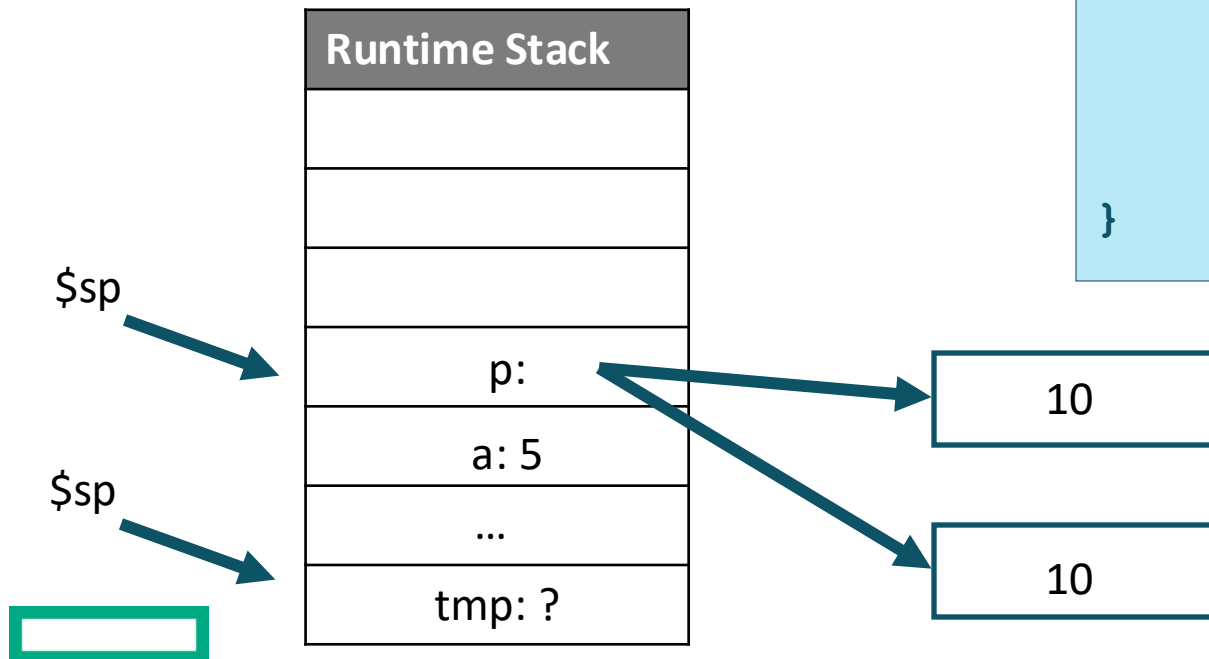
```
// example Java code  
void foo() {  
    MyClass obj = new MyClass();  
    obj.doSomething();  
} // obj goes out of scope; object becomes  
unreachable and is eligible for GC
```



HOW WOULD GC HELP WITH THIS EXAMPLE?

```
// example C code
void example() {
    int a = 5;           // local variable
    int* p = malloc(4); // heap allocation
    *p = 10;
} // p goes out of scope, but memory is still allocated unless freed

void main() {
    int tmp;           // local variable
    example();
    example();
}
```





OUTLINE: MEMORY SAFETY

- Visualizing
 - Use-after-free and memory leak bugs
 - How ownership types prevent them
 - How garbage collection prevents them
- Punchline from blog: comparison of languages
- Kinds of errors memory safety help prevent
 - Variable not initialized
 - Mishandling strings
 - Use-after-free
 - Out-of-bounds array access

Memory Safety in Chapel

Posted on April 10, 2025.

Tags:

Safety

Language Comparison

By: [Michael Ferguson](#)

MEMORY SAFETY IN CHAPEL

Memory Safety

- Properties of programming languages that help reduce bugs
- Red means not much help, yellow means some help, green means protection from error

Error	C	C++	Rust	Python	Chapel
Variable Not Initialized	✗	✗	✓	✓	✓
Mishandling Strings	✗	⚠	✓	✓	✓
Use-After-Free	✗	⚠	⚠	✓	⚠
Out-of-Bounds Array Access	✗	✗	✓	✓	⚠

Tradeoffs

- Garbage collection makes it really easy for users to have memory safe code
- Unfortunately, it has performance impacts
- This second grid shows how various languages balance these tradeoffs

	C/C++	Rust	Python	MPI	OpenSHMEM	Chapel
Productivity	−	✓	+	−	−	+
Performance	+	+	−	+	+	+
Scalability				+	+	+
Safety	−	+	+	−	−	✓

Key: +: great; ✓: good; −: drawback

VARIABLE NOT INITIALIZED

Languages that don't help

- C/C++ you can access junk left in stack from previous stack frames
- What if instead of junk it is confidential data?
- Leads to hard to find bugs and potential security issues

Other approaches

- Rust will generate a compile-time error
- Chapel variables are initialized to a default value for their type
- How does Python work?

```
1  #include <stdio.h>
2  int main() {
3      int x;
4
5      printf("x is %i\n", x);
6      return 0;
7  }
```

```
$ gcc unset-variable.c
$ ./a.out
x is 32764
```

unset-variable.rs

```
1  fn main() {
2      let mut x: i64; // OOPS! forgot to initialize x
3      let y = x;
4      println!("y is {}", y);
5  }
```

unset-int-variable.chpl

```
1  proc main() {
2      var x: int; // integer variables are set to 0 if not initialized
3      writeln(x);
4  }
```

MISHANDLING STRINGS IN C

string-greeting.c

```
1  #include <stdio.h>
2  #include <string.h>
3
4  #define MAX_GREETING 16
5
6  int main(int argc, char** argv) {
7      char greeting[MAX_GREETING]; // C doesn't really have string support;
8          // here we allocate an array to store the
9          // greeting
10
11          // OOPS! allocated array might not be big enough
12
13      strcpy(greeting, "Hello "); // copy "Hello " into 'greeting'
14      strcat(greeting, argv[1]); // append the passed name to the greeting
15      printf("%s\n", greeting);
16      return 0;
17 }
```

```
$ gcc string-greeting.c
$ ./a.out abcdefghijklmnopqrstuv
Hello abcdefghijklmnopqrstuv
*** stack smashing detected ***: terminated
Aborted (core dumped)
```



BETTER STRING HANDLING

C and C++

- In C, best practice is to use `strncpy` instead of `strcpy` and `strlcat` instead of `strcat`
- In C++, use the standard string type

Rust, Python, and Chapel

- All of these languages have built in string types

▼ [string-greeting.cpp](#)

```
1  #include <string>
2  #include <iostream>
3
4  int main(int argc, char** argv) {
5      std::string greeting = "Hello ";
6      greeting += argv[1];           // append the passed name to the greeting
7      std::cout << greeting << std::endl; // print the greeting
8
9      return 0;
10 }
```

▼ [string-greeting.chpl](#)

```
1  config const who = ""; // enable command-line options like --who=world
2
3  var greeting = "Hello ";
4  greeting += who;
5  writeln(greeting);
```


USE AFTER FREE

C and C++

- In C, use-after-free is very easy to accidentally write
- In C++, `unique_ptr` and `shared_ptr` help to some extent, but use-after-free is still possible

Rust, Python, and Chapel

- Rust, compile-time checking prevents a use-after-free in safe code, still possible in unsafe code
- Python, garbage collector avoids this issue
- Chapel, automatic memory management for most types means free is only needed when using unmanaged classes

[use-after-free-scoped.chpl](#)

```
1  class C { var x: int; }
2
3  proc main() {
4      // create a reference to a 'C' instance on the heap
5      var b: borrowed C? = nil;
6
7      {
8          var instance = new owned C(0);
9          b = instance.borrow();
10     }
11
12     b!.x = 42;
13
14     // do other heap operations to make heap corruption more visible
15     {
16         var x = new owned C(2);
17         var y = new owned C(3);
18     }
19
20     writeln(b);
21 }
```

ARRAY BOUNDS CHECKING

C and C++

- Out-of-bounds array accesses aren't checked unless running with a memory-checking tool like valgrind
- If you haven't learned valgrind, check it out, it is a life changing tool for debugging C/C++ code

Rust, Python, and Chapel

- Rust, out-of-bounds array accesses cause the program to halt with an out-of-bounds error
- Python, out-of-bounds array accesses raise an error
- Chapel, by default out-of-bounds will cause program to halt, compiling with --fast turns checking off

```
$ g++ out-of-bounds.cpp
$ ./a.out 123456789
zsh: segmentation fault ./a.out 123456789
```

```
$ python3 out-of-bounds.py 123456789
Traceback (most recent call last):
  File "/Users/mferguson/chapel-blog/content/
    main(sys.argv)
    ~~~~^
  File "/Users/mferguson/chapel-blog/content/
    x = array[idx]
    ~~~~^
IndexError: list index out of range
```

```
$ chpl out-of-bounds.chpl
$ ./out-of-bounds --idx=123456789
out-of-bounds.chpl:5: error: halt reached - array index out of bounds
note: index was 123456789 but array bounds are 0..9
```