



# MORE PARALLELISM IN CHAPEL



Chapel Team, edited by Michelle Strout for CSc 372 at UArizona

April 22, 2025

# PLAN

- **Announcements**

- Brad Chamberlain, Chapel co-creator, visiting Thursday April 24th
- LA3 is due on Friday April 25<sup>th</sup> (3 days)
- Final projects are due Friday May 2<sup>nd</sup> (10 days)

- **Last time**

- Recall extra credit opportunities
- Memory safety comparison between languages

- **Today**

- TopHat questions about Questions for Brad, Memory safety, and Chapel implicit communication
- Parallelizing histogram in Chapel
- Other parallel constructs

## OUTLINE: MORE PARALLELISM

---

- Spectrum of Chapel loops
- Task intents including reduce intents
- Parallelizing histogram
- Story of index gather parallelization
- Other parallel constructs: 'cobegin', 'begin', 'sync`
- Avoiding races with task intents and task-private variables
- Performance gotchas



# SPECTRUM OF CHAPEL LOOPS

# SPECTRUM OF CHAPEL FOR-LOOP STYLES

See <https://chapel-lang.org/docs/primers/loops.html> for more details on loops.

**for loop:** each iteration is executed serially by the current task

- predictable execution order, similar to conventional languages

**foreach loop:** all iterations executed by the current task, but in no specific order

- a candidate for vectorization, SIMD execution on GPUs

**forall loop:** all iterations are executed by one or more tasks in no specific order

- implemented using one or more tasks, locally or distributed, as determined by the iterand expression

```
forall i in 1..n do ...           // forall loops over ranges use local tasks only
forall (i,j) in {1..n, 1..n} do ... // ditto for local domains...
forall elem in myLocArr do ...    // ...and local arrays
forall elem in myDistArr do ...   // distributed arrays use tasks on each locale owning part of the array
forall i in myParIter(...) do ... // you can also write your own iterators that use the policy you want
```

**coforall loop:** each iteration is executed concurrently by a distinct task

- explicit parallelism; supports synchronization between iterations (tasks)

# IMPLICIT LOOPS: PROMOTION OF SCALAR SUBROUTINES & ARRAY OPS

- Any function or operator that takes scalar arguments can be called with array expressions instead

```
proc foo(x: real, y: real, z: real) {  
    return x**y + 10*z;  
}
```

- Interpretation is similar to that of a zippered forall loop, thus:

```
C = foo(A, 2, B);
```

is equivalent to:

```
forall (c, a, b) in zip(C, A, B) do  
    c = foo(a, 2, b);
```

as is:

```
C = A**2 + 10*B;
```



## IN CLASS ACTIVITY: UNDERSTANDING PROMOTION AND ZIP

- Questions to answer with Chapel code examples from previous slide
  - What sizes can arrays A, B, and C be declared with relative to each other for the following to work? What happens if C is bigger? A and B are bigger?

```
C = foo(A, 2, B);
```

- Using zip and unbounded ranges, how could we get an array of tuples (#, elem), where # is the index of each element in an array?

```
<TODO in class>
```

- Using zip and unbounded ranges, how could we get an array of tuples (#,elem), where # is the count of each element in an array?

```
<TODO in class>
```



TASK INTENTS INCLUDING REDUCE INTENTS



# USING TASK INTENTS IN LOOPS

**Recall Procedure argument intents** (<https://chapel-lang.org/docs/primers/procedures.html?highlight=intents#argument-intents>)

- Tell how to pass a symbol actual argument into a formal parameter
- Default intent is 'const', which means formal can't be modified in procedure body
- 'ref' means formal can be changed AND that change will be visible elsewhere, e.g., at the callsite
- Others: 'in', 'out', and 'inout' refer to copying the actual argument in, the formal out, or both

## Task intents in loops

- Similar to argument intents in syntax and philosophy
- Also have a 'reduce' intent similar to OpenMP
- 'reduce' intent means each task has its own copy and specified operation like '+' will combine at end of loop

## Design principles

- Avoid common race conditions
- Avoid copies of (potentially) large data structures



# TASK INTENTS IN FORALL LOOPS: SCALARS

```
var sum: real;  
forall i in 1..n do  
    sum += computeMyResult(i);
```

Default intent of scalars is 'const in' so this is illegal (and avoids a race)

```
var sum: real;  
forall i in 1..n with (ref sum) do  
    sum += computeMyResult(i);
```

With 'ref' intent, we are requesting a race

```
var sum: real;  
forall i in 1..n with (+ reduce sum) do  
    sum += computeMyResult(i);
```

Override default intent so that each task accumulates its own copy. On loop exit, all tasks combine their results into original 'sum'

# FORALL INTENT EXAMPLES: ARRAYS

```
var bucketCount: [0.. $m$ ] real;  
forall i in 1.. $n$  with (ref bucketCount) do  
    bucketCount[i %  $m$ ] += 1;
```

'ref' intent avoids array copies,  
but can result in data races

```
var bucketCount: [0.. $m$ ] real;  
forall i in 1.. $n$  with (in bucketCount) do  
    bucketCount[i %  $m$ ] += 1;
```

*'in' intent will result in  
each task having its  
own copy*

```
var bucketCount : [0.. $m$ ] real;  
forall i in 1.. $n$  with (+ reduce bucketCount) do  
    bucketCount[i %  $m$ ] += 1;
```

*'reduce' intent will result in  
each task having own copy,  
but then on loop exit tasks  
combine their results into the  
original 'bucketCount'*

# ATOMIC VARIABLES

## Meaning

- Atomic means 'indivisible'
- An atomic operation is indivisible.
- A thread of computation cannot interfere with another thread that is doing an atomic operation.

## Atomic Type Semantics in Chapel

- Supports operations on variable atomically w.r.t. other tasks
- Based on C/C++ atomic operations

## Example: Counting barrier

```
var count: atomic int, done: atomic bool;  
  
proc barrier(numTasks) {  
    const myCount = count.fetchAdd(1);  
    if (myCount < numTasks - 1) then  
        done.waitFor(true);  
    else  
        done.testAndSet();  
}
```

# ARRAY OF ATOMIC

```
var bucketCount: [0..<m] atomic real;  
forall i in 1..n with (ref bucketCount) do  
  bucketCount[i % m].add(1);
```

Make the 'bucketCount' array  
contain 'atomic real's

Use the atomic 'add' operation

```
var bucketCount: [0..<m] atomic real;  
forall i in 1..n do  
  bucketCount[i % m].add(1);
```

Can leave off 'ref' intent, since  
that is the default for 'atomic'  
types

## PARALLELIZING HISTOGRAM (HANDS ON)

# HANDS ON: PARALLELIZING HISTOGRAM

## Goals

- Parallelize a program that computes a histogram using reductions
- Parallelize it using an array of atomic integers
- Compare the performance of both versions versus each other and the serial version

## Parallelize 'histogram-serial.chpl' using a 'forall' loop and a 'reduction' intent

1. Copy 'histogram-serial.chpl' into 'histogram-reduce.chpl'
2. Parallelize the serial 'for' loop using concepts from '04-task-intents.chpl'

## Parallelize 'histogram-serial.chpl' using an array of atomic integers

1. Copy 'histogram-serial.chpl' into 'histogram-atomic.chpl'
2. Parallelize the serial 'for' loop using concepts from '04-atomic-type.chpl'

## Compare the performance of all three

```
./histogram-serial --numNumbers=100000000 --printRandomNumbers=false --useRandomSeed=false  
./histogram-reduce --numNumbers=100000000 --printRandomNumbers=false --useRandomSeed=false  
./histogram-atomic --numNumbers=100000000 --printRandomNumbers=false --useRandomSeed=false
```





# HANDS ON FURTHER INVESTIGATION: PARALLELIZE N-BODY



nbody.chpl

## Goals and Questions to Answer

- Parallelize as many loops in n-body as possible
- Determine when a 'reduce' intent or 'atomic' variable type is needed
- How can you check if you got the same answer?
- Is it possible for floating-point roundoff differences to change what the answers are slightly? For which loops?
- Did you get a performance improvement by doing the parallelization?



# ATOMIC METHODS

- `read() : t` return current value
- `write(v : t)` store *v* as current value
- `exchange(v : t) : t` store *v*, returning previous value
- `compareExchange(old : t, new : t) : bool`  
store *new* iff previous value was *old*; returns true on success
- `waitFor(v : t)` wait until the stored value is *v*
- `add(v : t)` add *v* to the value atomically
- `fetchAdd(v : t) : t` same, returning pre-sum value  
(*sub, or, and, xor* also supported similarly)
- `testAndSet()` like *exchange(true)* for atomic bool
- `clear()` like *write(false)* for atomic bool



# REDUCTIONS IN CHAPEL

- Recall the following snippet of code from the histogram exercise

```
// verify number of items in histogram is equal to number of random
// numbers and output timing results
if + reduce histogram != numNumbers then
    halt("Number of items in histogram does not match number of random numbers");
writeln("Histogram computed in ", timer.elapsed( ), " seconds\n");
```

- Standard reductions supported by default:

```
+, *, min, max, &, |, &&, ||, minloc, maxloc, ...
```

- Reductions can reduce arbitrary iterable expressions:

```
const total = + reduce Arr,
    factN = * reduce 1..n,
    biggest = max reduce (forall i in myIter() do foo(i));
```



# STORY OF INDEX-GATHER PARALLELIZATION

# STORY ABOUT PARALLELIZING INDEX GATHER

- Computation in Bale that gathers spread out data into a packed array

```
for i in D do
  Dest[i] = Src[Inds[i]];
```

- Parallelize over threads using a 'forall'

```
forall i in D with (ref Dest) do
  Dest[i] = Src[Inds[i]];
```

- Parallelize by distributing the D2 domain and using a 'forall'

```
const D = blockDist.createDomain({0..numUpdates-1});
var Inds: [D] int;

forall i in D with (ref Dest) do
  Dest[i] = Src[Inds[i]];
```



# CHAPEL TENDS TO BE COMPACT, CLEAN, AND FAST (BALE INDEX-GATHER)

## Exstack version

```
i=0;
while( exstack_proceed(ex, (i==l_num_req)) ) {
  i0 = i;
  while(i < l_num_req) {
    l_idx = pckindx[i] >> 16;
    pe = pckindx[i] & 0xffff;
    if(!exstack_push(ex, &l_idx, pe))
      break;
    i++;
  }

  exstack_exchange(ex);

  while(exstack_pop(ex, &idx, &fromth)) {
    idx = ltable[idx];
    exstack_push(ex, &idx, fromth);
  }
  lgp_barrier();
  exstack_exchange(ex);

  for(j=i0; j<i; j++) {
    fromth = pckindx[j] & 0xffff;
    exstack_pop_thread(ex, &idx, (uint64_t)fromth);
    tgt[j] = idx;
  }
  lgp_barrier();
}
```

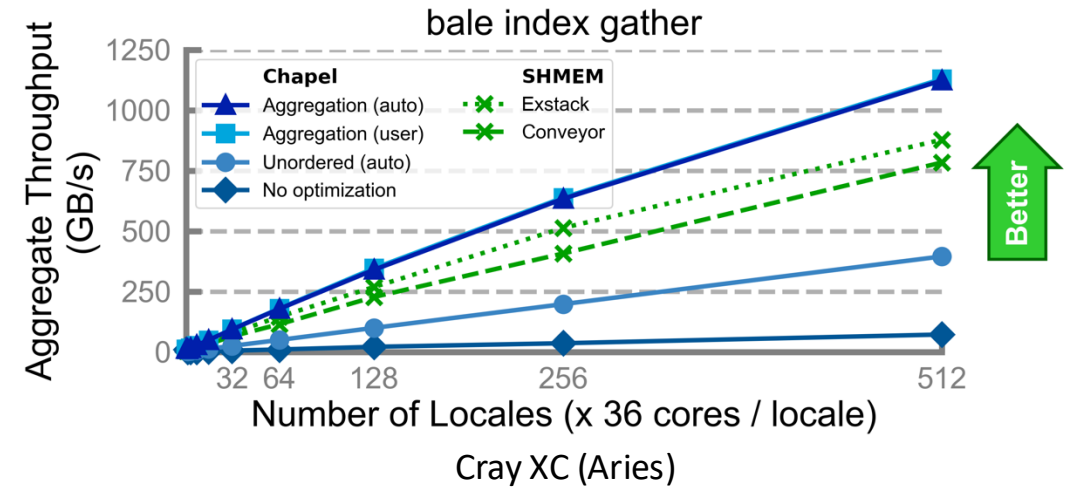
## Conveyors version

```
i = 0;
while (more = convey_advance(requests, (i == l_num_req)),
       more | convey_advance(replies, !more)) {

  for (; i < l_num_req; i++) {
    pkg.idx = i;
    pkg.val = pckindx[i] >> 16;
    pe = pckindx[i] & 0xffff;
    if (!convey_push(requests, &pkg, pe))
      break;
  }

  while (convey_pull(requests, ptr, &from) == convey_OK) {
    pkg.idx = ptr->idx;
    pkg.val = ltable[ptr->val];
    if (!convey_push(replies, &pkg, from)) {
      convey_unpull(requests);
      break;
    }
  }

  while (convey_pull(replies, ptr, NULL) == convey_OK)
    tgt[ptr->idx] = ptr->val;
}
```



## Manually Tuned Chapel version (using explicit aggregator type)

```
forall (d, i) in zip(Dest, Inds) with (var agg = new SrcAggregator(int)) do
  agg.copy(d, Src[i]);
```

## Elegant Chapel version (compiler-optimized w/ '--auto-aggregation')

```
forall (d, i) in zip(Dst, Inds) do
  d = Src[i];
```

OTHER PARALLEL CONSTRUCTS

## DEFINING OUR TERMS

---

**Task:** a unit of computation that can/should execute in parallel with other tasks

**Thread:** a system resource that executes tasks

- not exposed in the language
- occasionally exposed in the implementation

**Task Parallelism:** a style of parallel programming in which parallelism is driven by programmer-specified tasks

(in contrast with):

**Data Parallelism:** a style of parallel programming in which parallelism is driven by computations over collections of data elements or their indices





# PARALLELISM SUPPORTED BY CHAPEL

## Synchronous task parallelism

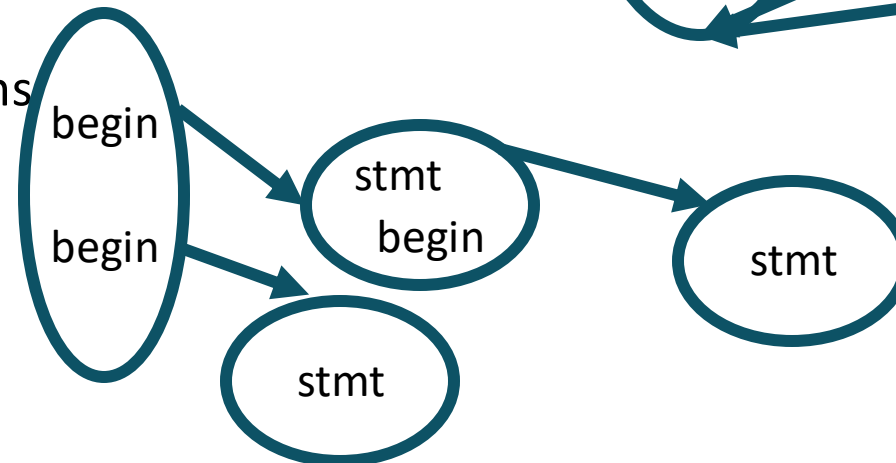
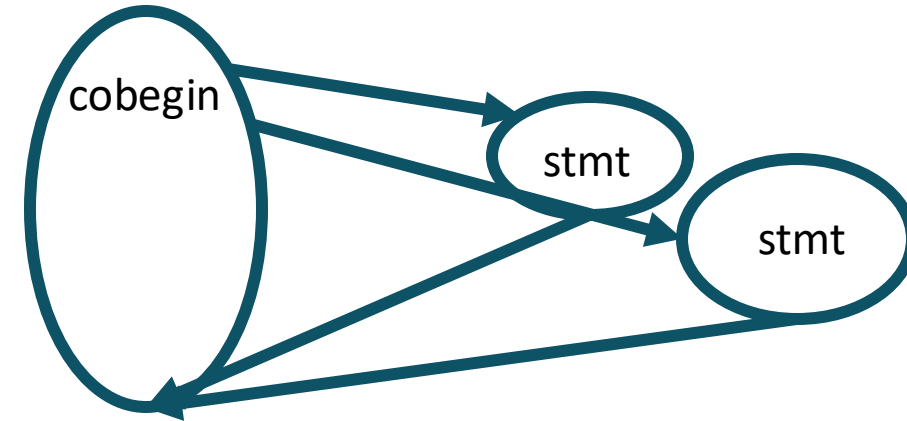
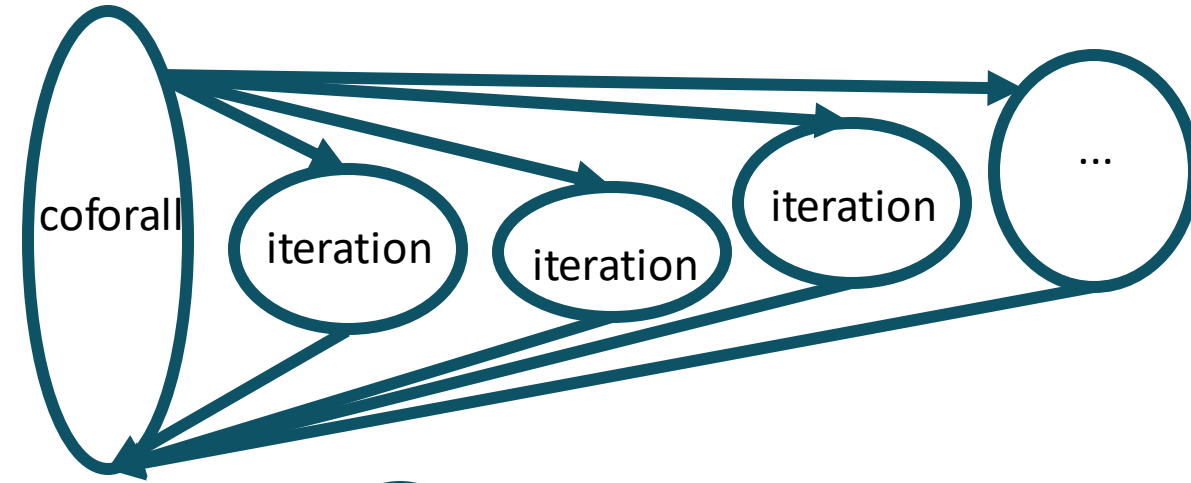
- 'coforall', parallel task per iteration
- 'cobegin', executes all statements in block in parallel

## Asynchronous task parallelism

- 'begin', creates an asynchronous task
- 'sync' and 'atomic' vars for task coordination

## Higher-level parallelism abstractions

- 'forall', data parallelism and iterator abstraction
- 'foreach', SIMD parallelism
- 'scan', operations such as cumulative sums
- 'reduce', operations such as summation



# PARALLELISM SUPPORTED BY CHAPEL

## Synchronous task parallelism

- 'coforall', parallel task per iteration
- 'cobegin', executes all statements in block in parallel

## Asynchronous task parallelism

- 'begin', creates an asynchronous task
- 'sync' and 'atomic' vars for task coordination

## Higher-level parallelism abstractions

- 'forall', data parallelism and iterator abstraction
- 'foreach', SIMD parallelism
- 'scan', operations such as cumulative sums
- 'reduce', operations such as summation

```
coforall loc in Locales do on loc { /* ... */ }
coforall tid in 0..<numTasks { /* ... */ }

cobegin { doTask0(); doTask1(); ... doTaskN(); }

var x : atomic int = 0, y : sync int;
sync {
    begin x.add(1);
    begin y.writeEF(1);
    begin x.sub(1);
    begin { y.readFE(); y.writeEF(0); }
}
assert(x.read() == 0);
assert(y.readFE() == 0);

var n = [i in 1..10] i*i;
forall x in n do x += 1;

var nPartialSums = + scan n;
var nSum = + reduce n;
```

## OTHER TASK PARALLEL FEATURES

- **begin / cobegin statements:** the two other ways of creating tasks

```
begin stmt;    // fire off an asynchronous task to run 'stmt'
```

```
cobegin {      // fire off a task for each of 'stmt1', 'stmt2', ...  
  stmt1;  
  stmt2;  
  stmt3;  
  ...  
}              // wait here for these tasks to complete before proceeding
```

- **atomic / synchronized variables:** types for safe data sharing & coordination between

```
var sum: atomic int;    // supports various atomic methods like .add(), .compareExchange(), ...  
var cursor: sync int;   // stores a full/empty bit governing reads/writes, supporting .readFE(), .writeEF()
```

- **coforall** *i* in 1..*n*iters **with** (**ref** *x*, + **reduce** *y*, **var** *z*: **int**) { ... }



# USE OF PARALLELISM IN SOME APPLICATIONS AND BENCHMARKS

Application	Distributed 'coforall'	Threaded 'coforall'	Asynchronous 'begin'	'cobegin'	sync or atomic	forall	scan
Arkouda	✓	✓			✓	✓	✓
CHAMPS	✓	✓			✓		
ChOp	✓		✓		✓	✓	
ParFlow						✓	
Coral Reef	✓	✓		✓		✓	



# TASK PARALLELISM: BEGIN STATEMENTS

```
// create a fire-and-forget task for a statement  
begin writeln("hello world");  
writeln("goodbye");
```

Possible outputs:

```
hello world  
goodbye
```

```
goodbye  
hello world
```



# JOINING SUB-TASKS: SYNC-STATEMENTS

## Syntax

```
sync-statement:  
  sync stmt
```

## Definition

- Executes *stmt*
- Waits for all *dynamically-scoped* begins to complete

## Examples

```
sync {  
  for i in 1..numConsumers {  
    begin consumer(i);  
  }  
  producer();  
}
```

```
proc search(node: TreeNode) {  
  if (node != nil) {  
    begin search(node.left);  
    begin search(node.right);  
  }  
}  
sync { search(root); }
```

# TASK PARALLELISM: COBEGIN STATEMENTS

```
// create a task per child statement  
cobegin {  
    producer (1) ;  
    producer (2) ;  
    consumer (1) ;  
} // implicit join of the three tasks here
```



# COBEGINS/SERIAL BY EXAMPLE: QUICKSORT



04-quicksort.chpl

'cobegin' will start both  
'quickSort' calls in parallel  
unless the number of  
running tasks would exceed  
the available HW parallelism

```
proc quickSort(arr: [?D],
               low: int = D.low,
               high: int = D.high) {
  if high - low < 8 {
    bubbleSort(arr, low, high);
  } else {
    const pivotLoc = partition(arr, low, high);
    serial (here.runningTasks() > here.maxTaskPar) do
      cobegin {
        quickSort(arr, low, pivotLoc-1);
        quickSort(arr, pivotLoc+1, high);
      }
    }
  }
}
```



# TASK PARALLELISM: COFORALL LOOPS

```
// create a task per iteration  
coforall t in 0..#numTasks {  
    writeln("Hello from task ", t, " of ", numTasks);  
} // implicit join of the numTasks tasks here  
  
writeln("All tasks done");
```

Sample output:

```
Hello from task 2 of 4  
Hello from task 0 of 4  
Hello from task 3 of 4  
Hello from task 1 of 4  
All tasks done
```



# COMPARISON OF BEGIN, COBEGIN, AND COFORALL

## **begin:**

- Use to create a dynamic task with an unstructured lifetime
- “fire and forget” (or at least “leave running for awhile”)

## **cobegin:**

- Use to create a related set of heterogeneous tasks  
...or a small, fixed set of homogenous tasks
- The parent task depends on the completion of the tasks

## **coforall:**

- Use to create a fixed or dynamic # of homogenous tasks
- The parent task depends on the completion of the tasks

**Note:** All these concepts can be composed arbitrarily



# SYNCHRONIZATION VARIABLES

## TASK PARALLELISM: DATA-DRIVEN SYNCHRONIZATION

---

- **sync variables:** store full-empty state along with value
- **atomic variables:** support atomic operations
  - e.g., compare-and-swap; atomic sum, multiply, etc.
  - similar to C/C++



# BOUNDED BUFFER PRODUCER/CONSUMER EXAMPLE

```
// 'sync' types store full/empty state along with value
var buff: [0..#buffersize] sync real;

begin producer();
consumer();

proc producer() {
  var i = 0;
  for ... {
    i = (i+1) % buffersize;
    buff[i].writeEF( ... ); // wait for empty, write, leave full
  }
}

proc consumer() {
  var i = 0;
  while ... {
    i = (i+1) % buffersize;
    ...buff[i].readFE()...; // wait for full, read, leave empty
  }
}
```

## Syntax

```
sync-type:  
  sync type
```

## Semantics

- Stores *full/empty* state along with normal value
- Initially *full* if initialized, *empty* otherwise

## Examples: Critical sections and futures

```
var lock: sync bool;  
  
lock.writeEF(true);  
critical();  
lock.readFE();
```

```
var future: sync real;  
  
begin future.writeEF(compute());  
res = computeSomethingElse();  
useComputedResults(future.readFE(), res);
```

## SYNCHRONIZATION VARIABLE METHODS

---

- **readFE** () : t      block until *full*, leave *empty*, return value
- **readFF** () : t      block until *full*, leave *full*, return value
- **writeEF** (v : t)    block until *empty*, set value to  $v$ , leave *full*



# COMPARISON OF SYNCHRONIZATION TYPES

---

## **sync:**

- Best for producer/consumer style synchronization
  - “this task should block until something happens”
  - use single for write-once values

## **atomic:**

- Best for uncoordinated accesses to shared state
  - “these tasks are unlikely to interfere with each other, at least for very long...”





# AVOIDING RACES WITH TASK INTENTS AND TASK PRIVATE VARIABLES

# TASK INTENTS

---

- Tells how to “pass” variables from outer scopes to tasks
  - Similar to argument intents in syntax and philosophy
    - also adds a “reduce intent”, similar to OpenMP
  - Design principles:
    - “principle of least surprise”
    - avoid simple race conditions
    - avoid copies of (potentially) expensive data structures
    - support coordination via sync/atomic variables



# TASK INTENT EXAMPLES



04-task-intents-coforall.chpl

```
var sum: real;  
coforall i in 1..n do  
    sum += computeMyResult(i);
```

Default task intent of scalars is 'const in' so this is illegal (and avoids a race)

```
var sum: real;  
coforall i in 1..n with (ref sum) do  
    sum += computeMyResult(i);
```

Use a 'ref' task intent for 'sum' variable. We've now requested a race.

```
var sum: real;  
coforall i in 1..n with (+ reduce sum) do  
    sum += computeMyResult(i);
```

Use a 'reduce' task intent. Per-task sums will be reduced on task exit.

```
var sum: atomic real;  
coforall i in 1..n do  
    sum.add(computeMyResult(i));
```

Default task intent of atomics is 'ref' so this is legal, meaningful, and safe

# TASK-PRIVATE VARIABLES

- Task-parallel features support task-private variables easily

```
coforall i in 1..numTasks {  
    var mySum: real; // each task gets its own copy of mySum  
    for j in 1..n do  
        mySum += A[i][j];  
}
```

- Forall loops need special support for task-private variables

```
var oneSingleVariable: real;  
forall i in 1..n {  
    var onePerIteration: real;  
}
```



# TASK-PRIVATE VARIABLES

- Task-parallel features support task-private variables easily

```
coforall i in 1..numTasks {  
    var mySum: real; // each task gets its own copy of mySum  
    for j in 1..n do  
        mySum += A[i][j];  
}
```

- Forall loops need special support for task-private variables

```
var oneSingleVariable: real;  
forall i in 1..n with (var onePerTask: real) {  
    var onePerIteration: real;  
}
```

# TASK-PRIVATE VARIABLES

- Task-parallel features support task-private variables easily

```
coforall i in 1..numTasks {  
    var mySum: real; // each task gets its own copy of mySum  
    for j in 1..n do  
        mySum += A[i][j];  
}
```

- Forall loops need special support for task-private variables

```
var oneSingleVariable: real;  
forall i in 1..n with (var onePerTask = 3.14) {  
    var onePerIteration: real;  
}
```



# CHAPEL RESOURCES

Chapel homepage: <https://chapel-lang.org>

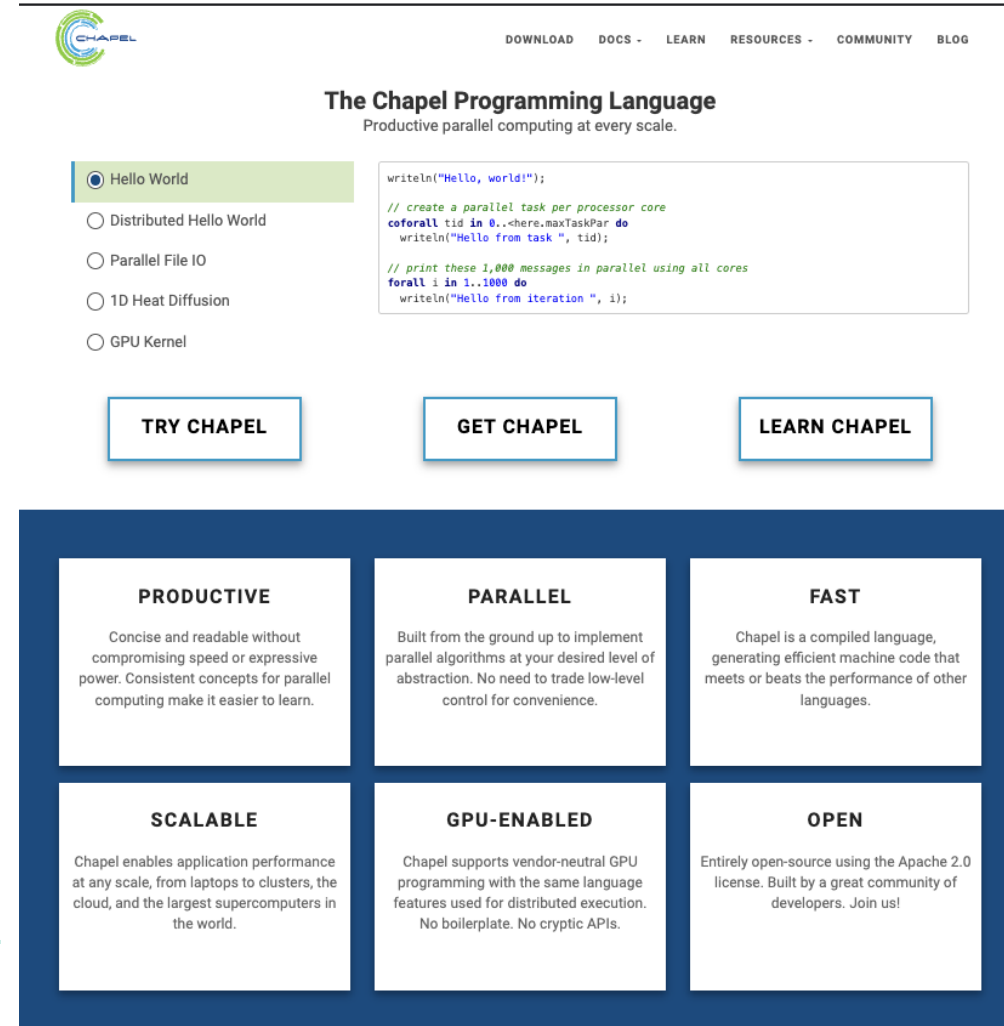
- (points to all other resources)

## Social Media:

- Twitter: [@ChapelLanguage](https://twitter.com/ChapelLanguage)
- Facebook: [@ChapelLanguage](https://facebook.com/ChapelLanguage)
- YouTube: <http://www.youtube.com/c/ChapelParallelProgrammingLanguage>

## Community Discussion / Support:

- Discord: <https://discord.com/invite/xu2xg45yqH>
- Stack Overflow: <https://stackoverflow.com/questions/tagged/chapel>
- GitHub Issues: <https://github.com/chapel-lang/chapel/issues>



The screenshot shows the Chapel Programming Language homepage. At the top, there's a navigation bar with links: DOWNLOAD, DOCS, LEARN, RESOURCES, COMMUNITY, and BLOG. The main heading is "The Chapel Programming Language" with the tagline "Productive parallel computing at every scale." Below this, there's a section for "Hello World" with a radio button selected, and four other options: Distributed Hello World, Parallel File IO, 1D Heat Diffusion, and GPU Kernel. To the right of these options is a code editor showing a "Hello World" program in Chapel. Below the code editor are three buttons: TRY CHAPEL, GET CHAPEL, and LEARN CHAPEL. At the bottom, there's a grid of six boxes, each describing a feature of Chapel: PRODUCTIVE, PARALLEL, FAST, SCALABLE, GPU-ENABLED, and OPEN.

**PRODUCTIVE**  
Concise and readable without compromising speed or expressive power. Consistent concepts for parallel computing make it easier to learn.

**PARALLEL**  
Built from the ground up to implement parallel algorithms at your desired level of abstraction. No need to trade low-level control for convenience.

**FAST**  
Chapel is a compiled language, generating efficient machine code that meets or beats the performance of other languages.

**SCALABLE**  
Chapel enables application performance at any scale, from laptops to clusters, the cloud, and the largest supercomputers in the world.

**GPU-ENABLED**  
Chapel supports vendor-neutral GPU programming with the same language features used for distributed execution. No boilerplate. No cryptic APIs.

**OPEN**  
Entirely open-source using the Apache 2.0 license. Built by a great community of developers. Join us!