



PARALLELISM II AND 372 OUTRO

Chapel Team, edited by Michelle Strout for CSc 372 at UArizona

April 22, 2025

PLAN

- **Announcements**

- Tuesday May 6th is the last day of class
 - ICA12 quiz will be Tuesday May 6th
 - SA8 is due May 6th
- Final exam is in this room, Wednesday May 14th from 3:30 to 5:30pm

- **Last time**

- Please go complete the survey at <https://scsonline.ucatt.arizona.edu>
- Group work time

- **Today**

- Outro
- Parallelizing histogram in Chapel
- Other parallel constructs

- Introduction to several major high-level programming languages and their characteristics. Programming projects are required in at least three languages.
- Main Goal: Be able to compare programming languages as tools to solve problems

- Compare different languages and paradigms in the context of real-world problems
 - You will need a vocabulary and an understanding of a range of programming language features to do this
 - Use these skills and concepts to build confidence in reading and understanding code in any programming language
- Some example language features we will discuss
 - First-class functions, i.e., lambda functions
 - Pattern matching
 - Type inference
 - Generics/templates (didn't get to)
 - Concurrency and parallelism

Outro: Recommendations for Studying for the Final Exam

- **Practice previous example problems**
 - Gather questions from the midterms, pre-assessment, quizzes, TopHat, gradescope online reviews (SA3, SA5, and SA8), and slides
 - Create questions and answers about the code you worked with in all the small programming assignments and large assignments
 - Ask questions in Piazza if there are concepts that are not clear to you
- **Concepts ~ 30 Questions Will Cover**
 - SML, Prolog, and Chapel code understanding
 - Comparisons between programming languages
 - Chapel parallelism including data races
 - Which language features are useful in particular application domains
 - Lexing and parsing as was done in LA1
 - Type inference as was done in LA2

OUTLINE: PARALLELISM II

- Parallelizing histogram
- Other parallel constructs: 'cobegin', 'begin', 'sync`
- Avoiding races with task intents and task-private variables



TASK INTENTS INCLUDING REDUCE INTENTS

USING TASK INTENTS IN LOOPS

Recall Procedure argument intents (<https://chapel-lang.org/docs/primers/procedures.html?highlight=intents#argument-intents>)

- Tell how to pass a symbol actual argument into a formal parameter
- Default intent is 'const', which means formal can't be modified in procedure body
- 'ref' means formal can be changed AND that change will be visible elsewhere, e.g., at the callsite
- Others: 'in', 'out', and 'inout' refer to copying the actual argument in, the formal out, or both

Task intents in loops

- Similar to argument intents in syntax and philosophy
- Also have a 'reduce' intent similar to OpenMP
- 'reduce' intent means each task has its own copy and specified operation like '+' will combine at end of loop

Design principles

- Avoid common race conditions
- Avoid copies of (potentially) large data structures



TASK INTENTS IN FORALL LOOPS: SCALARS

```
var sum: real;  
forall i in 1..n do  
    sum += computeMyResult(i);
```

Default intent of scalars is 'const in' so this is illegal (and avoids a race)

```
var sum: real;  
forall i in 1..n with (ref sum) do  
    sum += computeMyResult(i);
```

With 'ref' intent, we are requesting a race

```
var sum: real;  
forall i in 1..n with (+ reduce sum) do  
    sum += computeMyResult(i);
```

Override default intent so that each task accumulates its own copy. On loop exit, all tasks combine their results into original 'sum'

FORALL INTENT EXAMPLES: ARRAYS

```
var bucketCount: [0.. $m$ ] real;  
forall i in 1.. $n$  with (ref bucketCount) do  
    bucketCount[i %  $m$ ] += 1;
```

'ref' intent avoids array copies,
but can result in data races

```
var bucketCount: [0.. $m$ ] real;  
forall i in 1.. $n$  with (in bucketCount) do  
    bucketCount[i %  $m$ ] += 1;
```

*'in' intent will result in
each task having its
own copy*

```
var bucketCount : [0.. $m$ ] real;  
forall i in 1.. $n$  with (+ reduce bucketCount) do  
    bucketCount[i %  $m$ ] += 1;
```

*'reduce' intent will result in
each task having own copy,
but then on loop exit tasks
combine their results into the
original 'bucketCount'*

ATOMIC VARIABLES

Meaning

- Atomic means 'indivisible'
- An atomic operation is indivisible.
- A thread of computation cannot interfere with another thread that is doing an atomic operation.

Atomic Type Semantics in Chapel

- Supports operations on variable atomically w.r.t. other tasks
- Based on C/C++ atomic operations

Example: Counting barrier

```
var count: atomic int, done: atomic bool;  
  
proc barrier(numTasks) {  
    const myCount = count.fetchAdd(1);  
    if (myCount < numTasks - 1) then  
        done.waitFor(true);  
    else  
        done.testAndSet();  
}
```

ARRAY OF ATOMIC

```
var bucketCount: [0..<m] atomic real;  
forall i in 1..n with (ref bucketCount) do  
  bucketCount[i % m].add(1);
```

Make the 'bucketCount' array
contain 'atomic real's

Use the atomic 'add' operation

```
var bucketCount: [0..<m] atomic real;  
forall i in 1..n do  
  bucketCount[i % m].add(1);
```

Can leave off 'ref' intent, since
that is the default for 'atomic'
types

PARALLELIZING HISTOGRAM (HANDS ON)

HANDS ON: PARALLELIZING HISTOGRAM

Goals

- Parallelize a program that computes a histogram using reductions
- Parallelize it using an array of atomic integers
- Compare the performance of both versions versus each other and the serial version

Parallelize 'histogram-serial.chpl' using a 'forall' loop and a 'reduction' intent

1. Copy 'histogram-serial.chpl' into 'histogram-reduce.chpl'
2. Parallelize the serial 'for' loop using concepts from '04-task-intents.chpl'

Parallelize 'histogram-serial.chpl' using an array of atomic integers

1. Copy 'histogram-serial.chpl' into 'histogram-atomic.chpl'
2. Parallelize the serial 'for' loop using concepts from '04-atomic-type.chpl'

Compare the performance of all three

```
./histogram-serial --numNumbers=100000000 --printRandomNumbers=false --useRandomSeed=false  
./histogram-reduce --numNumbers=100000000 --printRandomNumbers=false --useRandomSeed=false  
./histogram-atomic --numNumbers=100000000 --printRandomNumbers=false --useRandomSeed=false
```



HANDS ON FURTHER INVESTIGATION: PARALLELIZE N-BODY



nbody.chpl

Goals and Questions to Answer

- Parallelize as many loops in n-body as possible
- Determine when a 'reduce' intent or 'atomic' variable type is needed
- How can you check if you got the same answer?
- Is it possible for floating-point roundoff differences to change what the answers are slightly? For which loops?
- Did you get a performance improvement by doing the parallelization?



ATOMIC METHODS

- `read() : t` return current value
- `write(v : t)` store *v* as current value
- `exchange(v : t) : t` store *v*, returning previous value
- `compareExchange(old : t, new : t) : bool`
store *new* iff previous value was *old*; returns true on success
- `waitFor(v : t)` wait until the stored value is *v*
- `add(v : t)` add *v* to the value atomically
- `fetchAdd(v : t) : t` same, returning pre-sum value
(*sub, or, and, xor* also supported similarly)
- `testAndSet()` like *exchange(true)* for atomic bool
- `clear()` like *write(false)* for atomic bool



REDUCTIONS IN CHAPEL

- Recall the following snippet of code from the histogram exercise

```
// verify number of items in histogram is equal to number of random
// numbers and output timing results
if + reduce histogram != numNumbers then
    halt("Number of items in histogram does not match number of random numbers");
writeln("Histogram computed in ", timer.elapsed( ), " seconds\n");
```

- Standard reductions supported by default:

```
+, *, min, max, &, |, &&, ||, minloc, maxloc, ...
```

- Reductions can reduce arbitrary iterable expressions:

```
const total = + reduce Arr,
    factN = * reduce 1..n,
    biggest = max reduce (forall i in myIter() do foo(i));
```

OTHER PARALLEL CONSTRUCTS

DEFINING OUR TERMS

Task: a unit of computation that can/should execute in parallel with other tasks

Thread: a system resource that executes tasks

- not exposed in the language
- occasionally exposed in the implementation

Task Parallelism: a style of parallel programming in which parallelism is driven by programmer-specified tasks

(in contrast with):

Data Parallelism: a style of parallel programming in which parallelism is driven by computations over collections of data elements or their indices



PARALLELISM SUPPORTED BY CHAPEL

Synchronous task parallelism

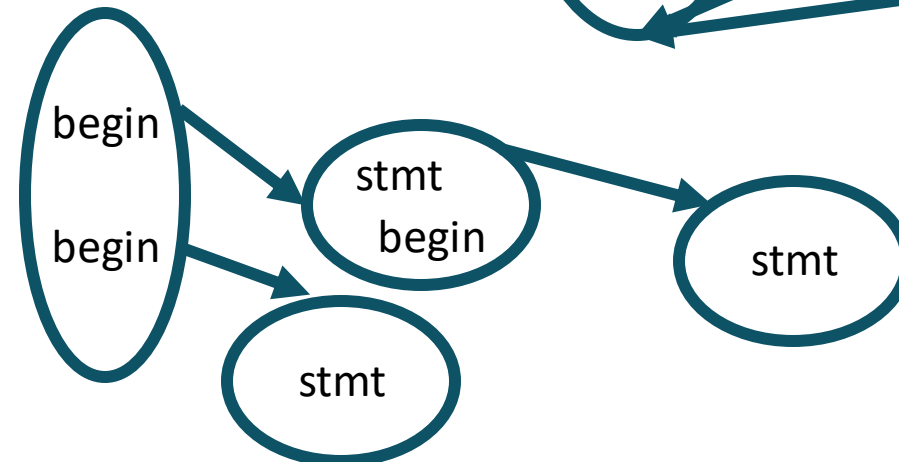
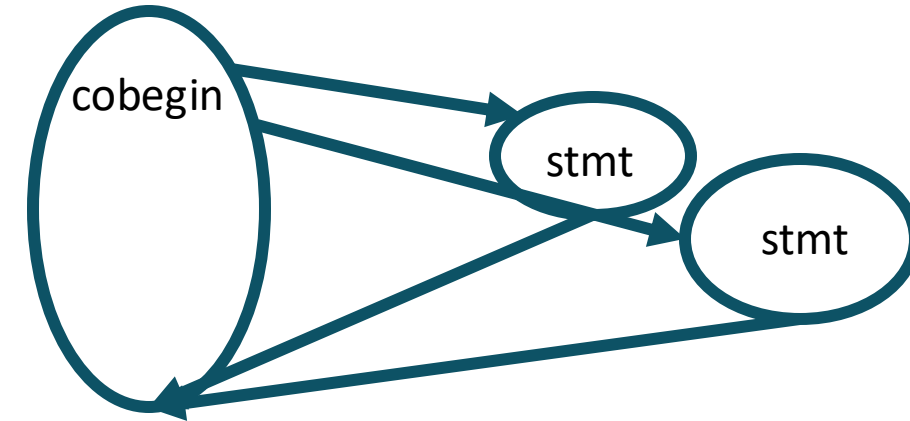
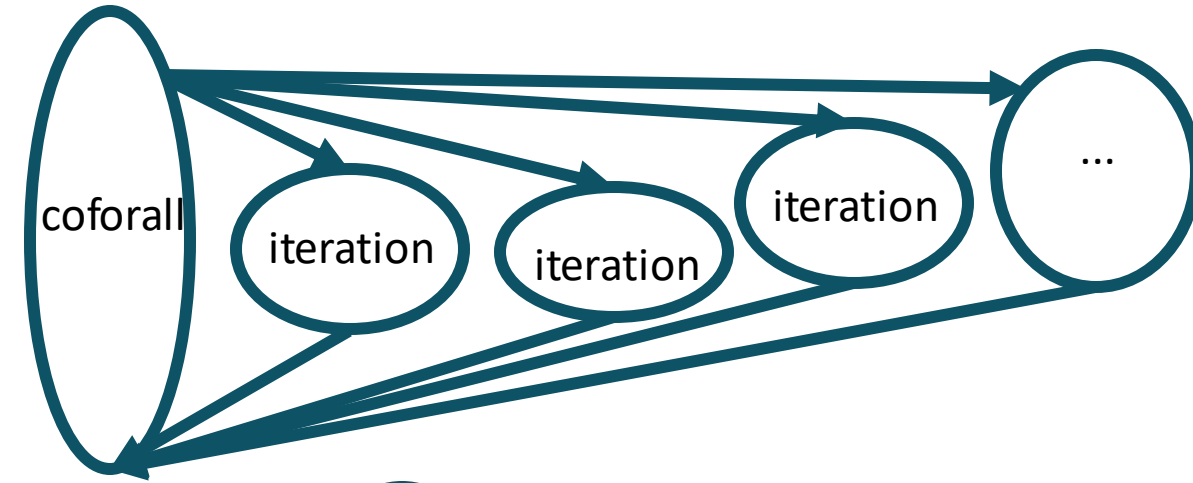
- 'coforall', parallel task per iteration
- 'cobegin', executes all statements in block in parallel

Asynchronous task parallelism

- 'begin', creates an asynchronous task
- 'sync' and 'atomic' vars for task coordination

Higher-level parallelism abstractions

- 'forall', data parallelism and iterator abstraction
- 'foreach', SIMD parallelism
- 'scan', operations such as cumulative sums
- 'reduce', operations such as summation



PARALLELISM SUPPORTED BY CHAPEL

Synchronous task parallelism

- 'coforall', parallel task per iteration
- 'cobegin', executes all statements in block in parallel

Asynchronous task parallelism

- 'begin', creates an asynchronous task
- 'sync' and 'atomic' vars for task coordination

Higher-level parallelism abstractions

- 'forall', data parallelism and iterator abstraction
- 'foreach', SIMD parallelism
- 'scan', operations such as cumulative sums
- 'reduce', operations such as summation

```
coforall loc in Locales do on loc { /* ... */ }
coforall tid in 0..<numTasks { /* ... */ }

cobegin { doTask0(); doTask1(); ... doTaskN(); }

var x : atomic int = 0, y : sync int;
sync {
    begin x.add(1);
    begin y.writeEF(1);
    begin x.sub(1);
    begin { y.readFE(); y.writeEF(0); }
}
assert(x.read() == 0);
assert(y.readFE() == 0);

var n = [i in 1..10] i*i;
forall x in n do x += 1;

var nPartialSums = + scan n;
var nSum = + reduce n;
```

OTHER TASK PARALLEL FEATURES

- **begin / cobegin statements:** the two other ways of creating tasks

```
begin stmt;    // fire off an asynchronous task to run 'stmt'
```

```
cobegin {      // fire off a task for each of 'stmt1', 'stmt2', ...  
  stmt1;  
  stmt2;  
  stmt3;  
  ...  
}              // wait here for these tasks to complete before proceeding
```

- **atomic / synchronized variables:** types for safe data sharing & coordination between

```
var sum: atomic int;    // supports various atomic methods like .add(), .compareExchange(), ...  
var cursor: sync int;   // stores a full/empty bit governing reads/writes, supporting .readFE(), .writeEF()
```

- **coforall** *i* in 1..*n*iters **with** (**ref** *x*, + **reduce** *y*, **var** *z*: **int**) { ... }



USE OF PARALLELISM IN SOME APPLICATIONS AND BENCHMARKS

Application	Distributed 'coforall'	Threaded 'coforall'	Asynchronous 'begin'	'cobegin'	sync or atomic	forall	scan
Arkouda	✓	✓			✓	✓	✓
CHAMPS	✓	✓			✓		
ChOp	✓		✓		✓	✓	
ParFlow						✓	
Coral Reef	✓	✓		✓		✓	



TASK PARALLELISM: BEGIN STATEMENTS

```
// create a fire-and-forget task for a statement  
begin writeln("hello world");  
writeln("goodbye");
```

Possible outputs:

```
hello world  
goodbye
```

```
goodbye  
hello world
```



JOINING SUB-TASKS: SYNC-STATEMENTS

Syntax

```
sync-statement:  
  sync stmt
```

Definition

- Executes *stmt*
- Waits for all *dynamically-scoped* begins to complete

Examples

```
sync {  
  for i in 1..numConsumers {  
    begin consumer(i);  
  }  
  producer();  
}
```

```
proc search(node: TreeNode) {  
  if (node != nil) {  
    begin search(node.left);  
    begin search(node.right);  
  }  
}  
sync { search(root); }
```

TASK PARALLELISM: COBEGIN STATEMENTS

```
// create a task per child statement  
cobegin {  
    producer (1) ;  
    producer (2) ;  
    consumer (1) ;  
} // implicit join of the three tasks here
```



COBEGINS/SERIAL BY EXAMPLE: QUICKSORT



04-quicksort.chpl

'cobegin' will start both
'quickSort' calls in parallel
unless the number of
running tasks would exceed
the available HW parallelism

```
proc quickSort(arr: [?D],
               low: int = D.low,
               high: int = D.high) {
  if high - low < 8 {
    bubbleSort(arr, low, high);
  } else {
    const pivotLoc = partition(arr, low, high);
    serial (here.runningTasks() > here.maxTaskPar) do
      cobegin {
        quickSort(arr, low, pivotLoc-1);
        quickSort(arr, pivotLoc+1, high);
      }
    }
  }
}
```

TASK PARALLELISM: COFORALL LOOPS

```
// create a task per iteration  
coforall t in 0..#numTasks {  
    writeln("Hello from task ", t, " of ", numTasks);  
} // implicit join of the numTasks tasks here  
  
writeln("All tasks done");
```

Sample output:

```
Hello from task 2 of 4  
Hello from task 0 of 4  
Hello from task 3 of 4  
Hello from task 1 of 4  
All tasks done
```



COMPARISON OF BEGIN, COBEGIN, AND COFORALL

begin:

- Use to create a dynamic task with an unstructured lifetime
- “fire and forget” (or at least “leave running for awhile”)

cobegin:

- Use to create a related set of heterogeneous tasks
...or a small, fixed set of homogenous tasks
- The parent task depends on the completion of the tasks

coforall:

- Use to create a fixed or dynamic # of homogenous tasks
- The parent task depends on the completion of the tasks

Note: All these concepts can be composed arbitrarily



SYNCHRONIZATION VARIABLES

TASK PARALLELISM: DATA-DRIVEN SYNCHRONIZATION

- **sync variables:** store full-empty state along with value
- **atomic variables:** support atomic operations
 - e.g., compare-and-swap; atomic sum, multiply, etc.
 - similar to C/C++



BOUNDED BUFFER PRODUCER/CONSUMER EXAMPLE

```
// 'sync' types store full/empty state along with value
var buff: [0..#buffersize] sync real;

begin producer();
consumer();

proc producer() {
  var i = 0;
  for ... {
    i = (i+1) % buffersize;
    buff[i].writeEF( ... ); // wait for empty, write, leave full
  }
}

proc consumer() {
  var i = 0;
  while ... {
    i = (i+1) % buffersize;
    ...buff[i].readFE()...; // wait for full, read, leave empty
  }
}
```

Syntax

```
sync-type:  
  sync type
```

Semantics

- Stores *full/empty* state along with normal value
- Initially *full* if initialized, *empty* otherwise

Examples: Critical sections and futures

```
var lock: sync bool;  
  
lock.writeEF(true);  
critical();  
lock.readFE();
```

```
var future: sync real;  
  
begin future.writeEF(compute());  
res = computeSomethingElse();  
useComputedResults(future.readFE(), res);
```

SYNCHRONIZATION VARIABLE METHODS

- **readFE** () : t block until *full*, leave *empty*, return value
- **readFF** () : t block until *full*, leave *full*, return value
- **writeEF** (v : t) block until *empty*, set value to v , leave *full*



COMPARISON OF SYNCHRONIZATION TYPES

sync:

- Best for producer/consumer style synchronization
 - “this task should block until something happens”
 - use single for write-once values

atomic:

- Best for uncoordinated accesses to shared state
 - “these tasks are unlikely to interfere with each other, at least for very long...”



AVOIDING RACES WITH TASK INTENTS AND TASK PRIVATE VARIABLES

TASK INTENTS

- Tells how to “pass” variables from outer scopes to tasks
 - Similar to argument intents in syntax and philosophy
 - also adds a “reduce intent”, similar to OpenMP
 - Design principles:
 - “principle of least surprise”
 - avoid simple race conditions
 - avoid copies of (potentially) expensive data structures
 - support coordination via sync/atomic variables



TASK INTENT EXAMPLES



04-task-intents-coforall.chpl

```
var sum: real;  
coforall i in 1..n do  
    sum += computeMyResult(i);
```

Default task intent of scalars is 'const in' so this is illegal (and avoids a race)

```
var sum: real;  
coforall i in 1..n with (ref sum) do  
    sum += computeMyResult(i);
```

Use a 'ref' task intent for 'sum' variable. We've now requested a race.

```
var sum: real;  
coforall i in 1..n with (+ reduce sum) do  
    sum += computeMyResult(i);
```

Use a 'reduce' task intent. Per-task sums will be reduced on task exit.

```
var sum: atomic real;  
coforall i in 1..n do  
    sum.add(computeMyResult(i));
```

Default task intent of atomics is 'ref' so this is legal, meaningful, and safe

TASK-PRIVATE VARIABLES

- Task-parallel features support task-private variables easily

```
coforall i in 1..numTasks {  
    var mySum: real; // each task gets its own copy of mySum  
    for j in 1..n do  
        mySum += A[i][j];  
}
```

- Forall loops need special support for task-private variables

```
var oneSingleVariable: real;  
forall i in 1..n {  
    var onePerIteration: real;  
}
```


TASK-PRIVATE VARIABLES

- Task-parallel features support task-private variables easily

```
coforall i in 1..numTasks {  
    var mySum: real; // each task gets its own copy of mySum  
    for j in 1..n do  
        mySum += A[i][j];  
}
```

- Forall loops need special support for task-private variables

```
var oneSingleVariable: real;  
forall i in 1..n with (var onePerTask: real) {  
    var onePerIteration: real;  
}
```



TASK-PRIVATE VARIABLES

- Task-parallel features support task-private variables easily

```
coforall i in 1..numTasks {  
    var mySum: real; // each task gets its own copy of mySum  
    for j in 1..n do  
        mySum += A[i][j];  
}
```

- Forall loops need special support for task-private variables

```
var oneSingleVariable: real;  
forall i in 1..n with (var onePerTask = 3.14) {  
    var onePerIteration: real;  
}
```



CHAPEL RESOURCES

Chapel homepage: <https://chapel-lang.org>

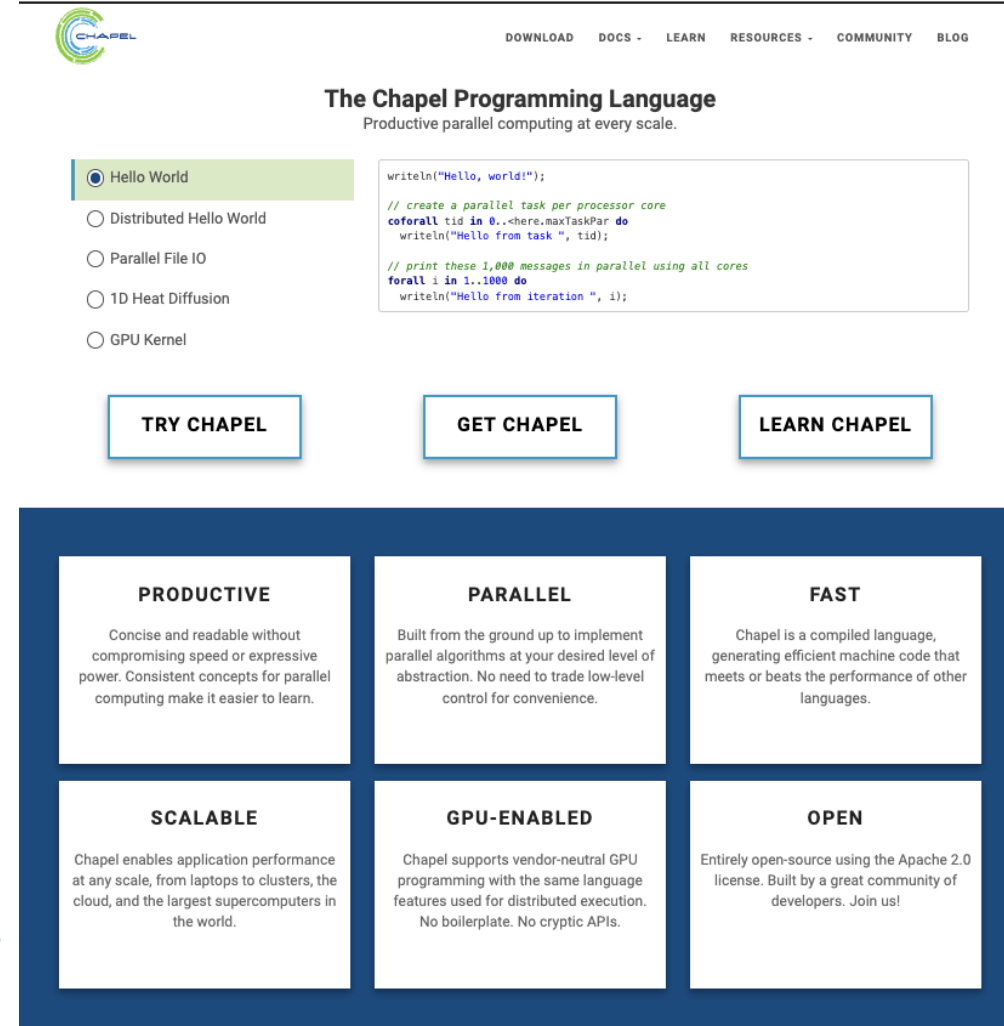
- (points to all other resources)

Social Media:

- Twitter: [@ChapelLanguage](https://twitter.com/ChapelLanguage)
- Facebook: [@ChapelLanguage](https://facebook.com/ChapelLanguage)
- YouTube: <http://www.youtube.com/c/ChapelParallelProgrammingLanguage>

Community Discussion / Support:

- Discord: <https://discord.com/invite/xu2xg45yqH>
- Stack Overflow: <https://stackoverflow.com/questions/tagged/chapel>
- GitHub Issues: <https://github.com/chapel-lang/chapel/issues>



The screenshot shows the Chapel Programming Language homepage. At the top, there's a navigation bar with links: DOWNLOAD, DOCS, LEARN, RESOURCES, COMMUNITY, and BLOG. The main heading is "The Chapel Programming Language" with the tagline "Productive parallel computing at every scale." Below this, there's a section for "Hello World" with a radio button selected, and four other options: Distributed Hello World, Parallel File IO, 1D Heat Diffusion, and GPU Kernel. To the right of these options is a code editor showing a "Hello, world!" program. Below the code editor are three buttons: TRY CHAPEL, GET CHAPEL, and LEARN CHAPEL. At the bottom, there's a grid of six boxes, each describing a feature of Chapel: PRODUCTIVE, PARALLEL, FAST, SCALABLE, GPU-ENABLED, and OPEN.

PRODUCTIVE
Concise and readable without compromising speed or expressive power. Consistent concepts for parallel computing make it easier to learn.

PARALLEL
Built from the ground up to implement parallel algorithms at your desired level of abstraction. No need to trade low-level control for convenience.

FAST
Chapel is a compiled language, generating efficient machine code that meets or beats the performance of other languages.

SCALABLE
Chapel enables application performance at any scale, from laptops to clusters, the cloud, and the largest supercomputers in the world.

GPU-ENABLED
Chapel supports vendor-neutral GPU programming with the same language features used for distributed execution. No boilerplate. No cryptic APIs.

OPEN
Entirely open-source using the Apache 2.0 license. Built by a great community of developers. Join us!

PLAN

- **Announcements**

- ICA12 today

- **Last time**

- Please go complete the survey at <https://scsonline.ucatt.arizona.edu>
- Group work time

- **Today**

- Outro
- ICA12: Final exam quiz
- Parallelizing histogram in Chapel
- Other parallel constructs